

# DS4300 - Large-Scale Storage and Retrieval

---

## NoSQL Overview

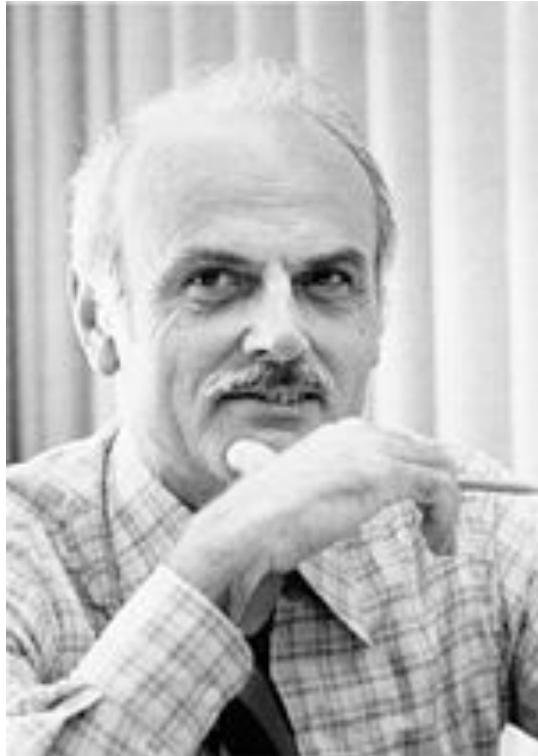


Northeastern University

John Rachlin, PhD  
Northeastern University

# Codd, 1970. The origins of the relational model

---



Edgar Frank Codd  
1923-2003

## *Information Retrieval*

---

### A Relational Model of Data for Large Shared Data Banks

E. F. CODD

*IBM Research Laboratory, San Jose, California*

Future users of large data banks must be protected from having to know how the data is organized in the machine (the internal representation). A prompting service which supplies such information is not a satisfactory solution. Activities of users at terminals and most application programs should remain unaffected when the internal representation of data is changed and even when some aspects of the external representation are changed. Changes in data representation will often be needed as a result of changes in query, update, and report traffic and natural growth in the types of stored information.

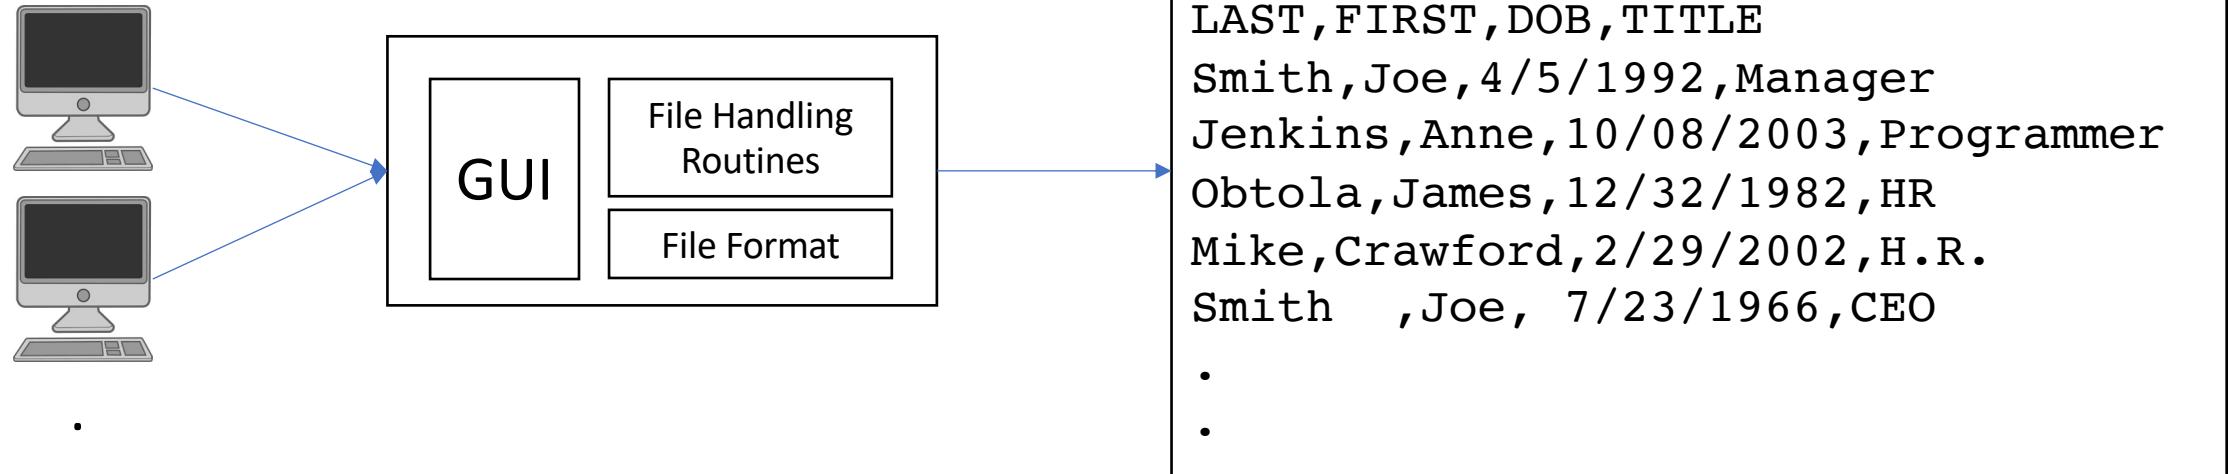


# Before Databases: File-Based Systems

**Problem:** You need to build a human-resources (HR) application to manage information about your employees.

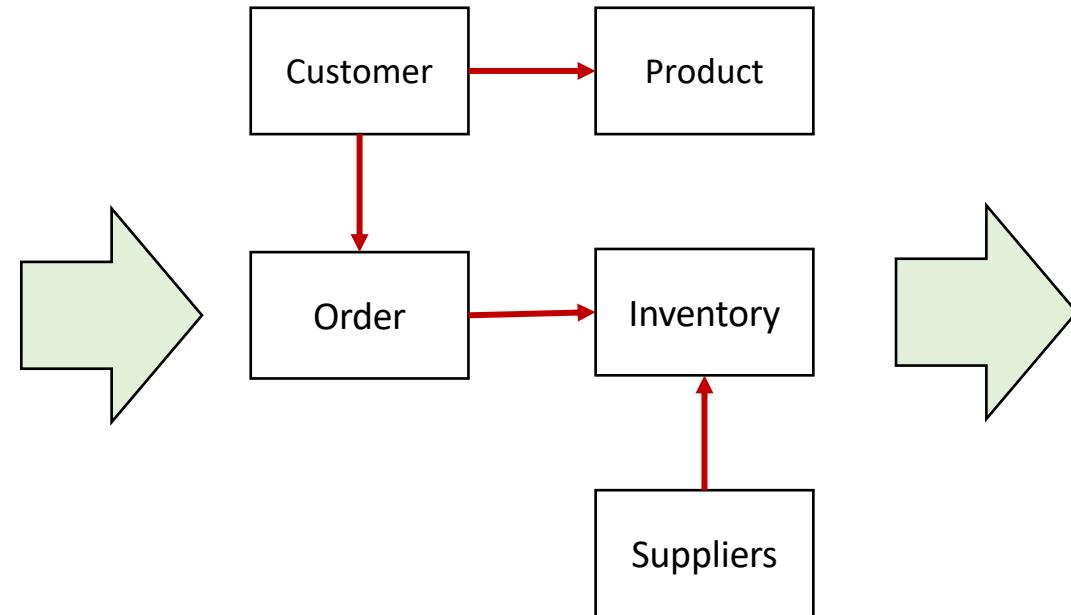
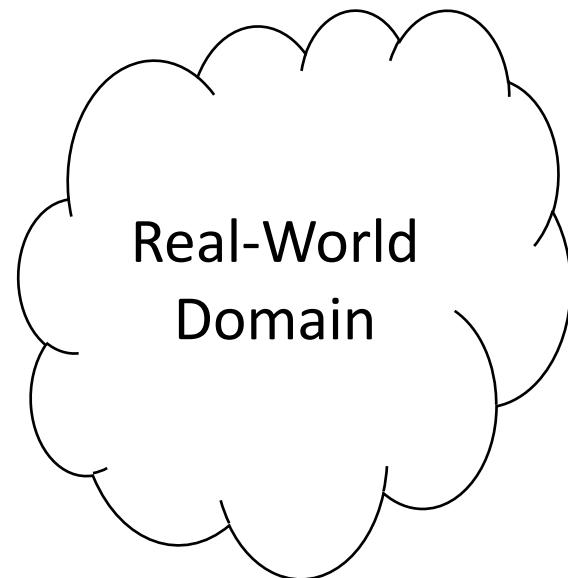
**User requirements:**

- Search ("Show me information about Joe Smith")
- Add / Delete ("Enter data about Anne Jenkins, our new hire. Delete Saniya Lee, she retired")
- Update ("James Obatola was promoted, update his title.")



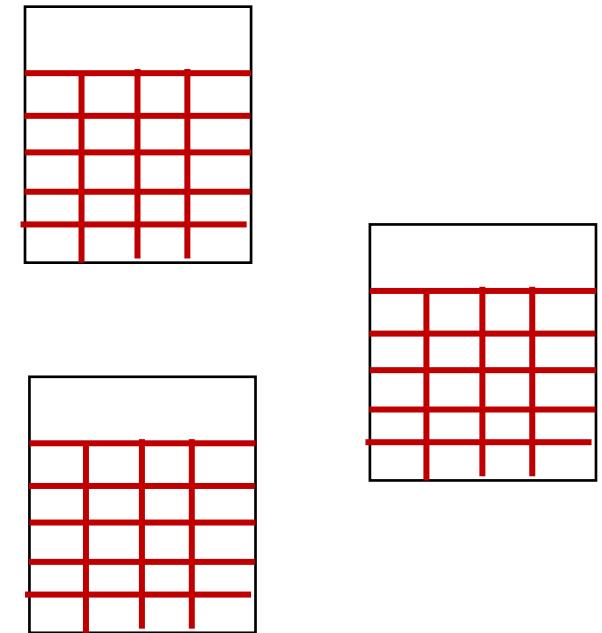
# What is a data model?

A **data model** is a description of the world, a domain, or an enterprise. The model determines how we interact with the data in a database both in terms of the operations we can perform, and the questions we can ask.



- **Entities** (things, events, concepts)
- **Attributes** (annotations on entities)
- **Relationships** (connections or associations between entities)

Entity-Relationship Diagram  
(A *conceptual* data model)



# Relational Data Model

Entities

Branch

branchNo	street	city	postCode
B005	22 Deer Rd	London	SW1 4EH
B007	16 Argyll St	Aberdeen	AB2 3SU
B003	163 Main St	Glasgow	G11 9QX
B004	32 Manse Rd	Bristol	BS99 1NZ
B002	56 Clover Dr	London	NW10 6EU

Staff

staffNo	fName	lName	position	sex	DOB	salary	branchNo
SL21	John	White	Manager	M	1-Oct-45	30000	B005
SG37	Ann	Beech	Assistant	F	10-Nov-60	12000	B003
SG14	David	Ford	Supervisor	M	24-Mar-58	18000	B003
SA9	Mary	Howe	Assistant	F	19-Feb-70	9000	B007
SG5	Susan	Brand	Manager	F	3-Jun-40	24000	B003
SL41	Julie	Lee	Assistant	F	13-Jun-65	9000	B005

A column or **attribute** or **field** providing more details about the particular entity.

- Name
- Datatype
- Other attributes

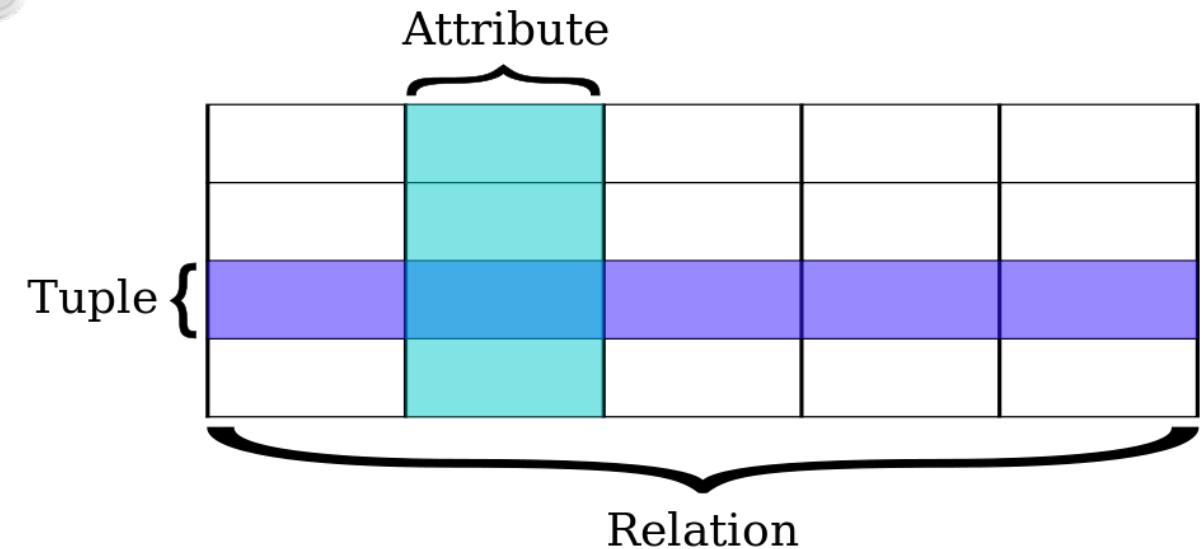
A row or **record** describing one instance of the entity

# Three pillars of Relational Databases

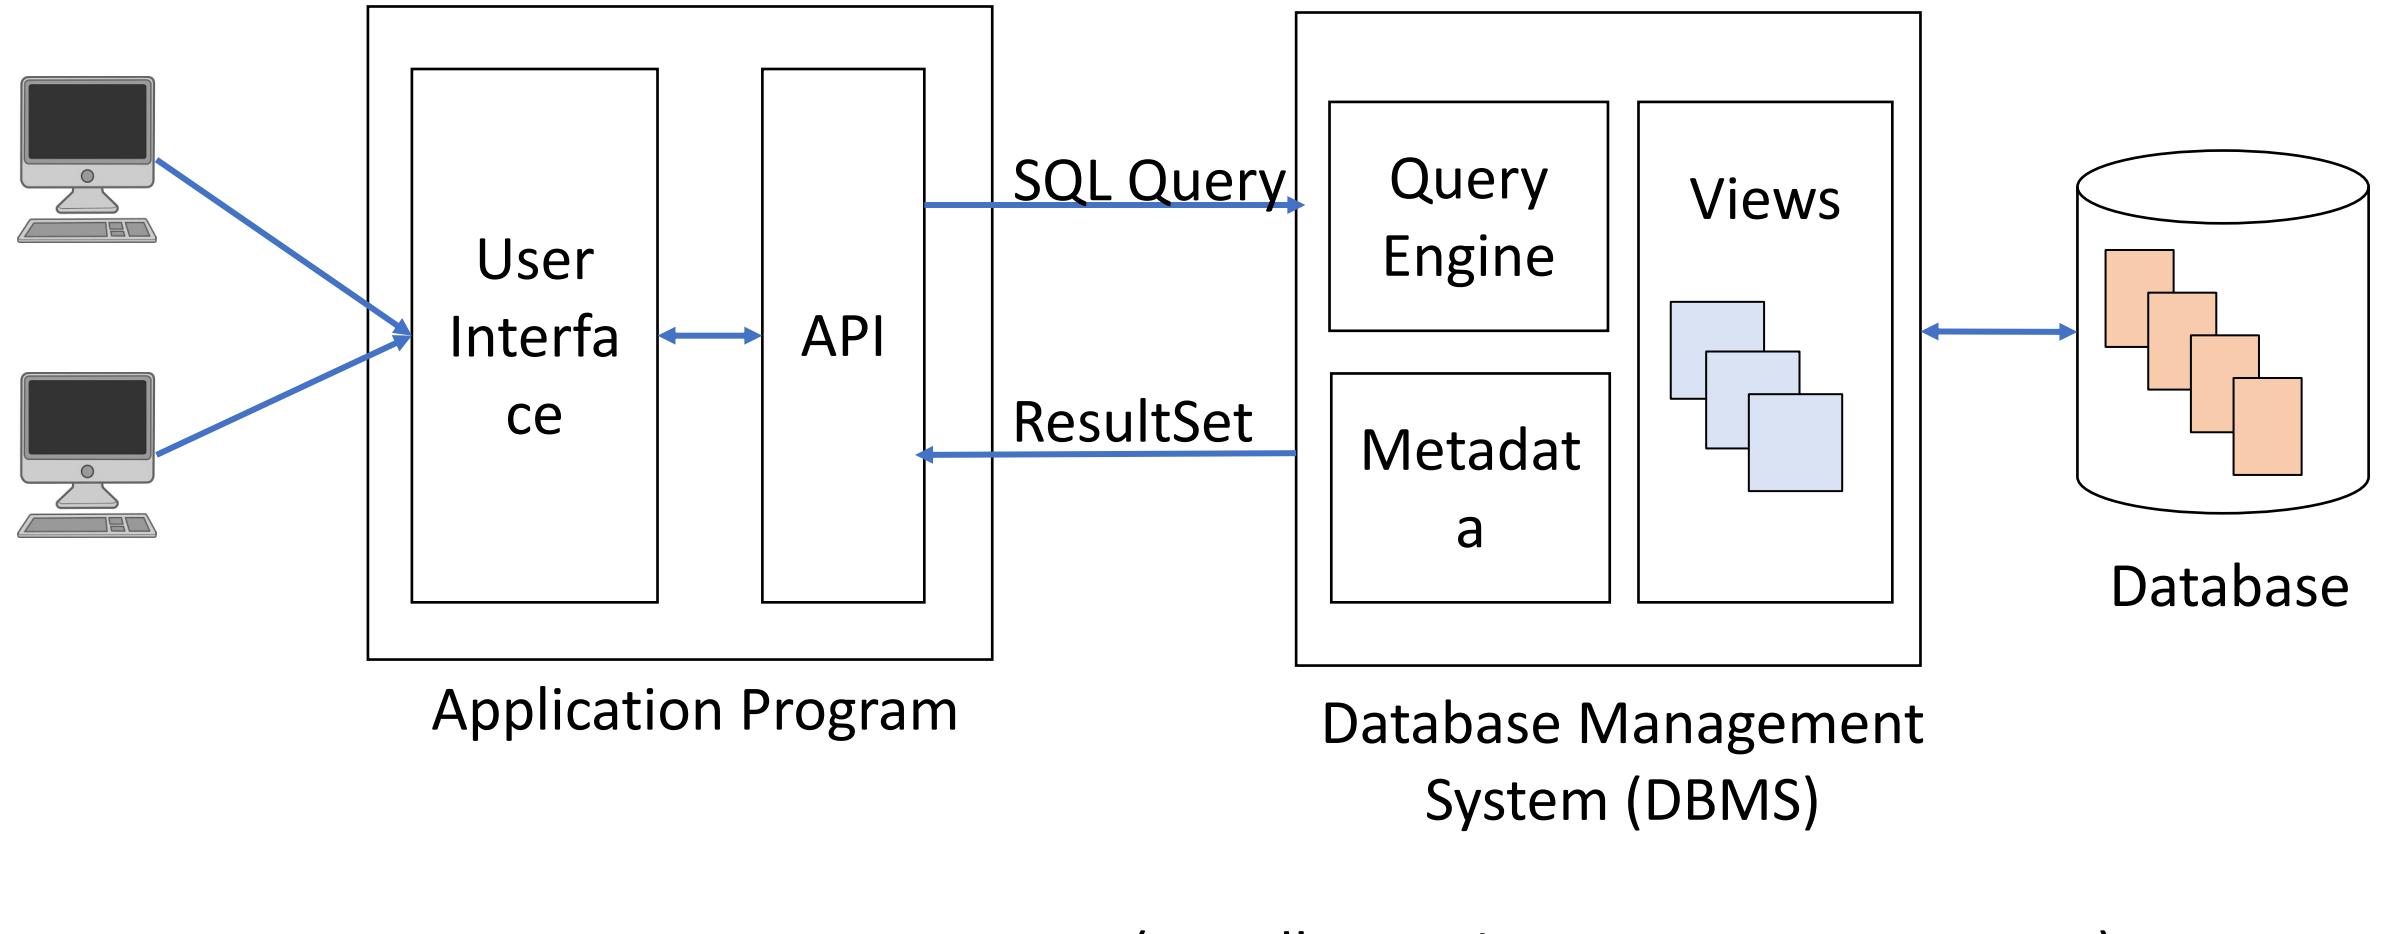
# relational model

# SQL

# ACID



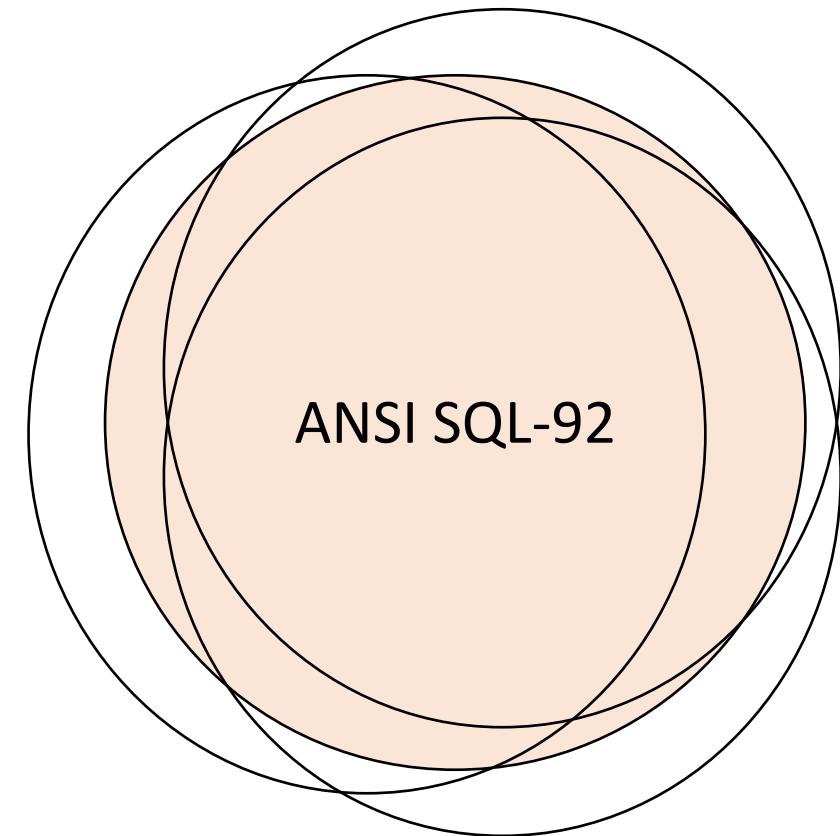
# Database Application



# SQL: The language of databases

---

- SQL = Structured Query Language
- Pronounced *either* S.Q.L. or *See*-quel (both are perfectly acceptable.)
- The language was *standardized* over time (SQL-86, SQL-89, SQL-92) but most vendors do not implement the standard completely or they introduce vendor-specific features / dialects.



One standard, many dialects



# The dominance of the relational model today

---

- Overcoming the limitations of file-based systems was a major technical achievement!
- Key features:
  - ✓ Data persistence
  - ✓ Concurrency support: multiple users simultaneously accessing data
  - ✓ Transaction support
  - ✓ Application integration
- A standard query language around which engines could be developed and optimized (though there are many dialects).
- Relational databases transformed how data is modeled, accessed, and transformed and have been **the dominant database model** since their development in the 1970's



# Why learn about databases?

---

Answer: Because databases are *ubiquitous*!

u·biq·ui·tous

/yōō'bikwədəs/ 🔍

*adjective*

*adjective: ubiquitous*

present, appearing, or found everywhere.

"his ubiquitous influence was felt by all the family"

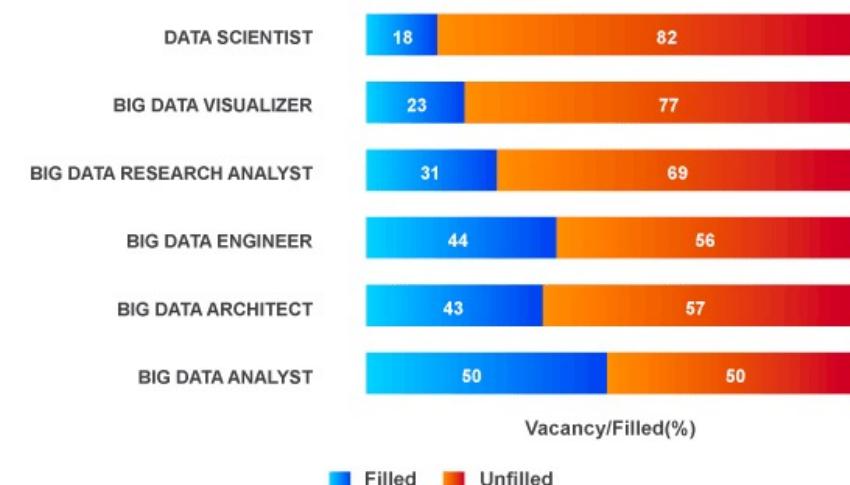
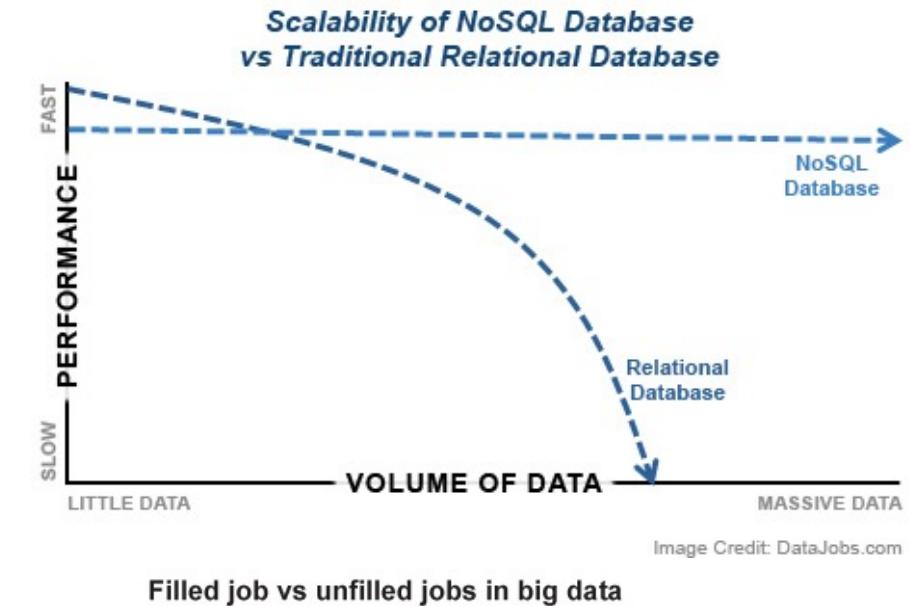
*synonyms:* omnipresent, ever-present, everywhere, all over the place, pervasive, universal, worldwide, global; More

*antonyms:* rare



# Why learn about NoSQL / Non-relational models?

- Traditional relational databases don't always scale well
- Stringent relational schema not ideal for unstructured data
- Expanding the breadth of your knowledge in data management
- \$\$\$



# Data Science Careers



Among the various industries that employ NoSQL specialists, the Healthcare Services industry pays the best, with a whopping average salary of \$135,948, followed by the Software Development industry, offering an average salary of \$120,043.

Source: <https://www.simplilearn.com/why-nosql-skills-are-crucial-for-big-data-career-article> (Dec, 2016)

## MODERN DATA SCIENTIST

Data Scientist, the sexiest job of the 21st century, requires a mixture of multidisciplinary skills ranging from an intersection of mathematics, statistics, computer science, communication and business. Finding a data scientist is hard. Finding people who understand who a data scientist is, is equally hard. So here is a little cheat sheet on who the modern data scientist really is.

### MATH & STATISTICS

- ★ Machine learning
- ★ Statistical modeling
- ★ Experiment design
- ★ Bayesian inference
- ★ Supervised learning: decision trees, random forests, logistic regression
- ★ Unsupervised learning: clustering, dimensionality reduction
- ★ Optimization: gradient descent and variants



### PROGRAMMING & DATABASE

- ★ Computer science fundamentals
- ★ Scripting language e.g. Python
- ★ Statistical computing packages, e.g., R
- ★ Databases: SQL and NoSQL
- ★ Relational algebra
- ★ Parallel databases and parallel query processing
- ★ MapReduce concepts
- ★ Hadoop and Hive/Pig
- ★ Custom reducers
- ★ Experience with xaaS like AWS

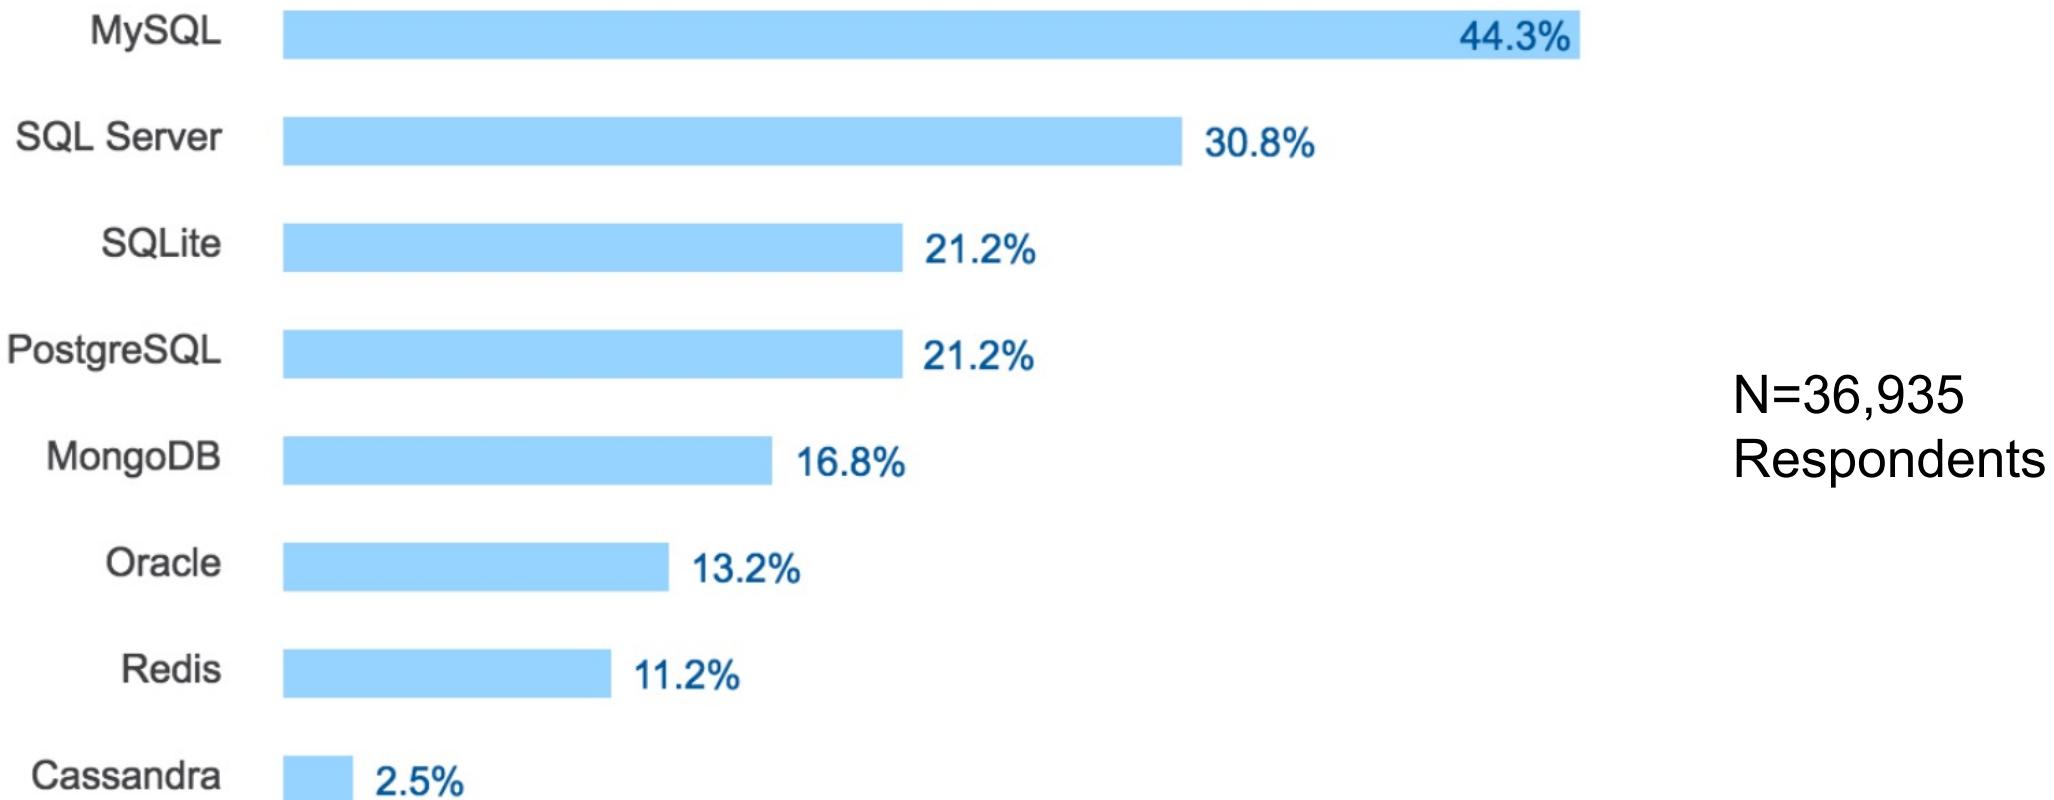
### COMMUNICATION & VISUALIZATION

- ★ Able to engage with senior management
- ★ Story telling skills
- ★ Translate data-driven insights into decisions and actions
- ★ Visual art design
- ★ R packages like ggplot or lattice
- ★ Knowledge of any of visualization tools e.g. Flare, D3.js, Tableau



# StackOverflow survey of most popular databases

---

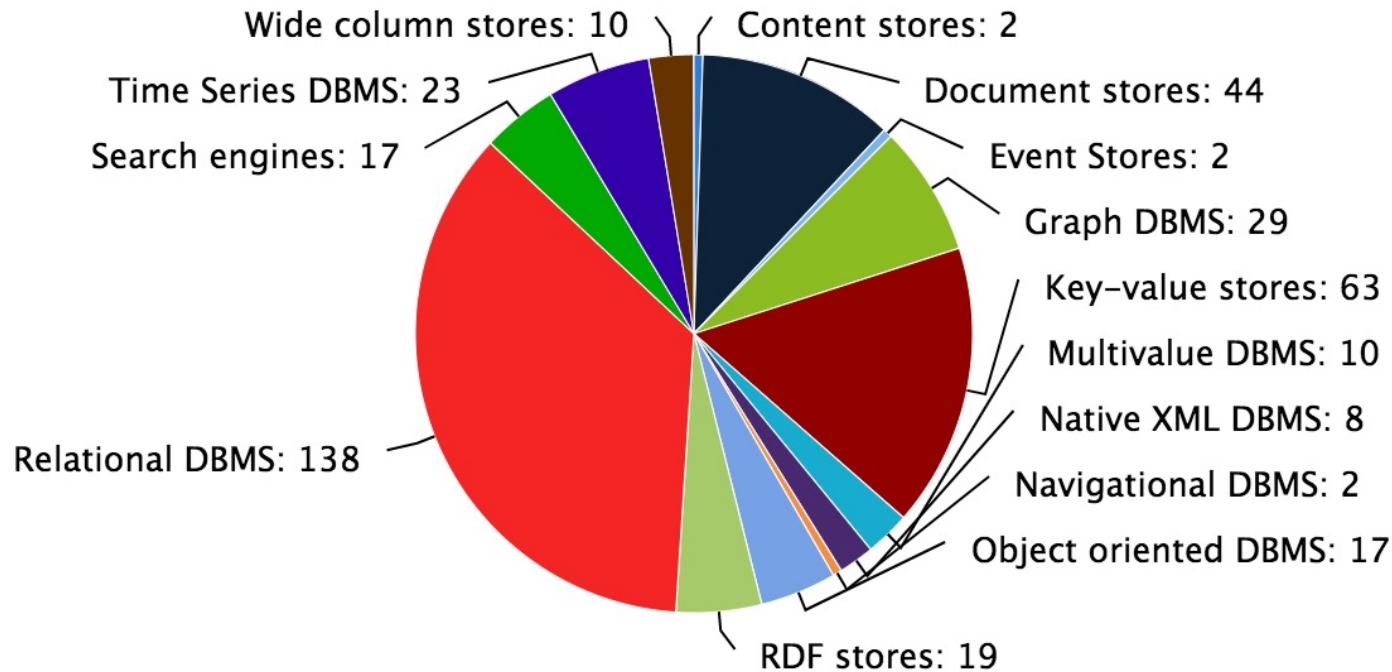


[https://www.eversql.com/most-popular-databases-in-2017 according-to-stackoverflow-survey/](https://www.eversql.com/most-popular-databases-in-2017-according-to-stackoverflow-survey/)



# DBMS Popularity

## Number of systems per category, January 2018



DB-Engines lists 341 different database management systems, which are classified according to their database model (e.g. relational DBMS, key-value stores etc.). This pie-chart shows the number of systems in each category. Some of the systems belong to more than one category.

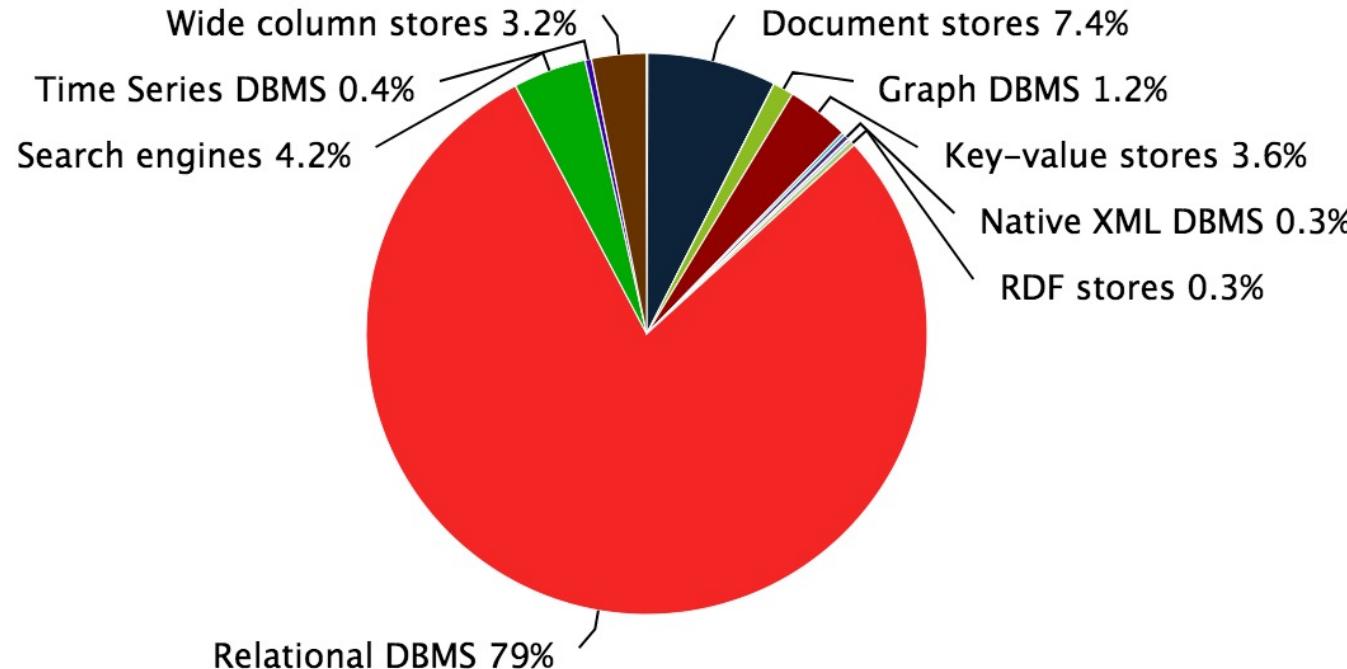
© 2018, DB-Engines.com



Northeastern University

# DBMS Popularity

## Ranking scores per category in percent, January 2018



© 2018, DB-Engines.com

This chart shows the popularity of each category. It is calculated with the popularity (i.e. the ranking scores) of all individual systems per category. The sum of all ranking scores is 100%.



# DB-Engines Ranking

337 systems in ranking, November 2017

Rank			DBMS	Database Model	Score		
Nov 2017	Oct 2017	Nov 2016			Nov 2017	Oct 2017	Nov 2016
1.	1.	1.	Oracle	Relational DBMS	1360.05	+11.25	-52.96
2.	2.	2.	MySQL	Relational DBMS	1322.03	+23.20	-51.53
3.	3.	3.	Microsoft SQL Server	Relational DBMS	1215.08	+4.76	+1.27
4.	4.	4.	PostgreSQL	Relational DBMS	379.92	+6.64	+54.10
5.	5.	5.	MongoDB	Document store	330.47	+1.07	+5.00
6.	6.	6.	DB2	Relational DBMS	194.06	-0.53	+12.61
7.	7.	↑ 8.	Microsoft Access	Relational DBMS	133.31	+3.86	+7.34
8.	8.	↓ 7.	Cassandra	Wide column store	124.21	-0.58	-9.76
9.	9.	9.	Redis	Key-value store	121.18	-0.87	+5.64
10.	10.	↑ 11.	Elasticsearch	Search engine	119.41	-0.82	+16.84

<https://db-engines.com/en/ranking>



# DB-Engines Ranking

371 systems in ranking, May 2021

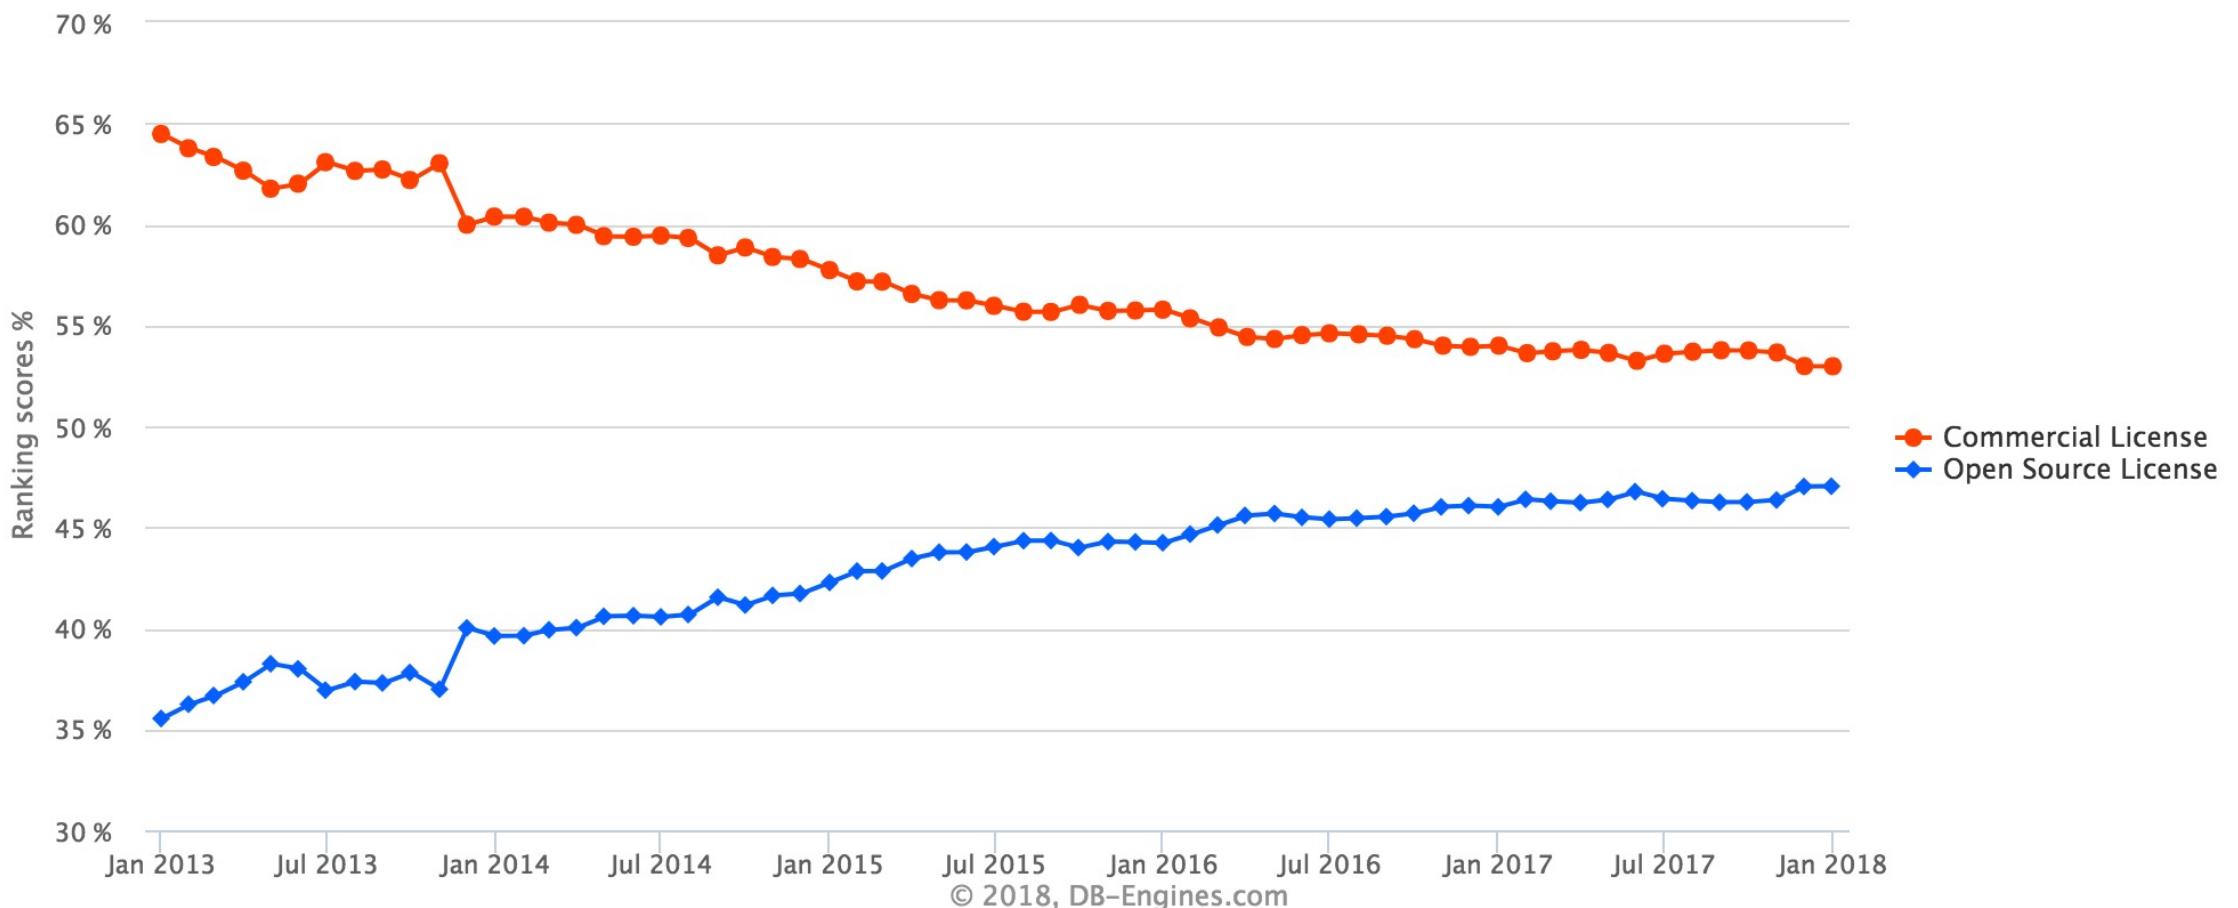
Rank			DBMS	Database Model	Score		
May 2021	Apr 2021	May 2020			May 2021	Apr 2021	May 2020
1.	1.	1.	Oracle	Relational, Multi-model	1269.94	-4.98	-75.50
2.	2.	2.	MySQL	Relational, Multi-model	1236.38	+15.69	-46.26
3.	3.	3.	Microsoft SQL Server	Relational, Multi-model	992.66	-15.30	-85.64
4.	4.	4.	PostgreSQL	Relational, Multi-model	559.25	+5.73	+44.45
5.	5.	5.	MongoDB	Document, Multi-model	481.01	+11.04	+42.02
6.	6.	6.	IBM Db2	Relational, Multi-model	166.66	+8.88	+4.02
7.	7.	↑ 8.	Redis	Key-value, Multi-model	162.17	+6.28	+18.69
8.	8.	↓ 7.	Elasticsearch	Search engine, Multi-model	155.35	+3.18	+6.23
9.	9.	9.	SQLite	Relational	126.69	+1.64	+3.66
10.	10.	10.	Microsoft Access	Relational	115.40	-1.33	-4.50

<https://db-engines.com/en/ranking>



# The growing importance of open-source databases

## Popularity trend



The above chart shows the historical trend of the popularity of open source and commercial database management systems.

# Most popular databases, January 2018

## Commercial

Rank	System	Score	Overall Rank
1.	Oracle	1342	1.
2.	Microsoft SQL Server	1148	3.
3.	DB2	190	6.
4.	Microsoft Access	127	7.
5.	Teradata	73	12.

## Open-Source

Rank	System	Score	Overall Rank
1.	MySQL	1300	2.
2.	PostgreSQL	386	4.
3.	MongoDB	331	5.
4.	Cassandra	124	8.
5.	Redis	123	9.



# The NoSQL Landscape: Maturity vs. Features

Figure 1. Magic Quadrant for Operational Database Management Systems



# What is Big Data?

- **Big Data** is generally concerned with the acquisition, storage, and analysis of extremely large volumes of data, typically many Terabytes or more.
- What is consider **big data** today may not be tomorrow. For example, in the 1980's, 10-100 GB might have been considered "big data" with disk drives typically in the 20-100 MB range.
- Circa 2018: 100-1000 TB or more would likely be considered big data by most companies (but not all!)

Prefixes for multiples of bits (bit) or bytes (B)					
Decimal			Binary		
Value	SI		Value	IEC	JEDEC
1000	k kilo		1024	Ki kibi	K kilo
$1000^2$	M mega		$1024^2$	Mi mebi	M mega
$1000^3$	G giga		$1024^3$	Gi gibi	G giga
$1000^4$	T tera		$1024^4$	Ti tebi	–
$1000^5$	P peta		$1024^5$	Pi pebi	–
$1000^6$	E exa		$1024^6$	Ei exbi	–
$1000^7$	Z zetta		$1024^7$	Zi zebi	–
$1000^8$	Y yotta		$1024^8$	Yi yobi	–

V • T • E



# Time is important too!

	N	10^N Sec	Min	Hrs	Days	Years	Centuries
Thousand / Kilo	-	1	0.0	0.0	0.0	0.0	0.0
	1	10	0.2	0.0	0.0	0.0	0.0
	2	100	1.7	0.0	0.0	0.0	0.0
	3	1,000	16.7	0.3	0.0	0.0	0.0
	4	10,000	166.7	2.8	0.1	0.0	0.0
	5	100,000	1,667	27.8	1.2	0.0	0.0
Million / Mega	6	1,000,000	16,667	278	11.6	0.0	0.0
	7	10,000,000	166,667	2,778	116	0.3	0.0
	8	100,000,000	1,666,667	27,778	1,157	3.2	0.0
Billion / Giga	9	1,000,000,000	16,666,667	277,778	11,574	32	0.3
	10	10,000,000,000	166,666,667	2,777,778	115,741	317	3.2
	11	100,000,000,000	1,666,666,667	27,777,778	1,157,407	3,169	31.7
Trillion / Tera	12	1,000,000,000,000	16,666,666,667	277,777,778	11,574,074	31,688	317
	13	1.E+13	2.E+11	3.E+09	1.E+08	3.E+05	3.E+03
	14	1.E+14	2.E+12	3.E+10	1.E+09	3.E+06	3.E+04
Quadrillion / Peta	15	1.E+15	2.E+13	3.E+11	1.E+10	3.E+07	3.E+05
	16	1.E+16	2.E+14	3.E+12	1.E+11	3.E+08	3.E+06
	17	1.E+17	2.E+15	3.E+13	1.E+12	3.E+09	3.E+07
Quintillian / Exa	18	1.E+18	2.E+16	3.E+14	1.E+13	3.E+10	3.E+08
	19	1.E+19	2.E+17	3.E+15	1.E+14	3.E+11	3.E+09
	20	1.E+20	2.E+18	3.E+16	1.E+15	3.E+12	3.E+10



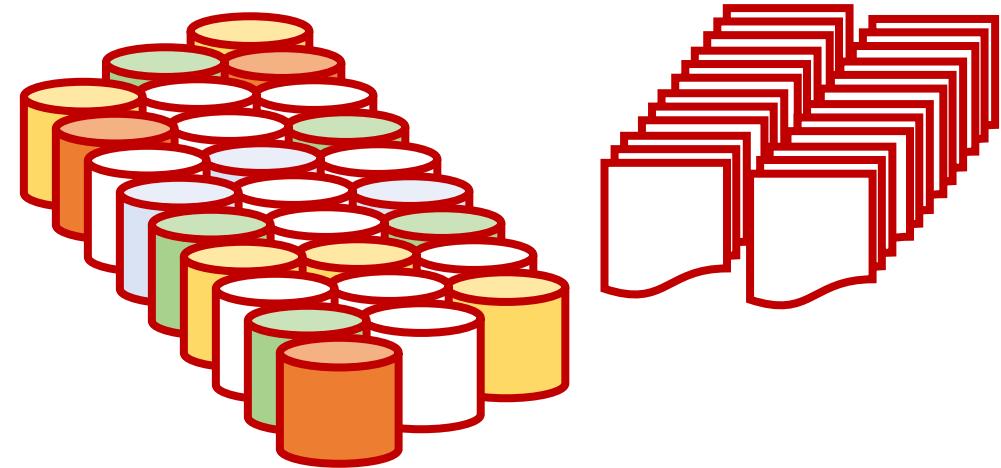
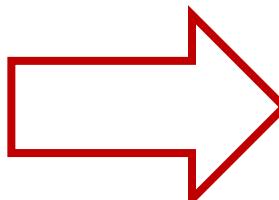
# Volume, Velocity, Variety

- **Volume:** How much data do you have?
- **Velocity:** How quickly is your data growing?
- **Variety:** How diverse or complex is your data?
- **Veracity:** How trustworthy is your data?



- Manageable volumes
- Limited growth
- Fixed schema, highly structured

**Traditional file systems  
and relational databases**



- Massive volumes
- Rapid growth with *streaming* data
- Both structured and un-structured or semi-structured datasets (tables, documents log files, images, videos)

**Big Data Solutions**



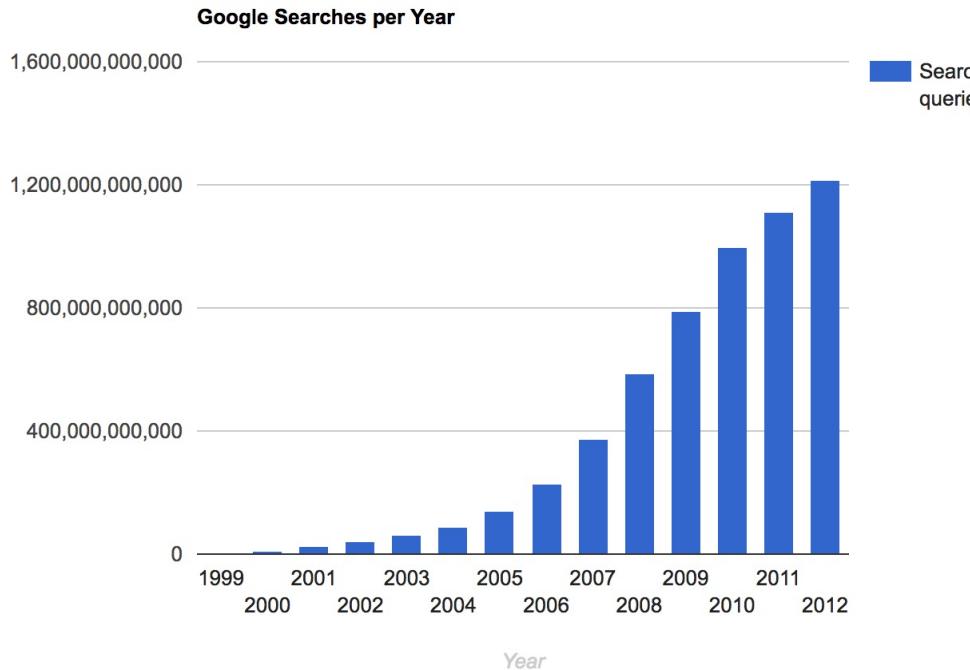
# Sources of Big Data

- Activity data (online transactions, credit card data, stock trading)
- Sensor networks
- On-line gaming systems supporting millions of users
- Biological databases, including genomic data
- Scientific instruments (LHC, LSST)
- Social Networking (twitter, Facebook, LinkedIn, text messaging)

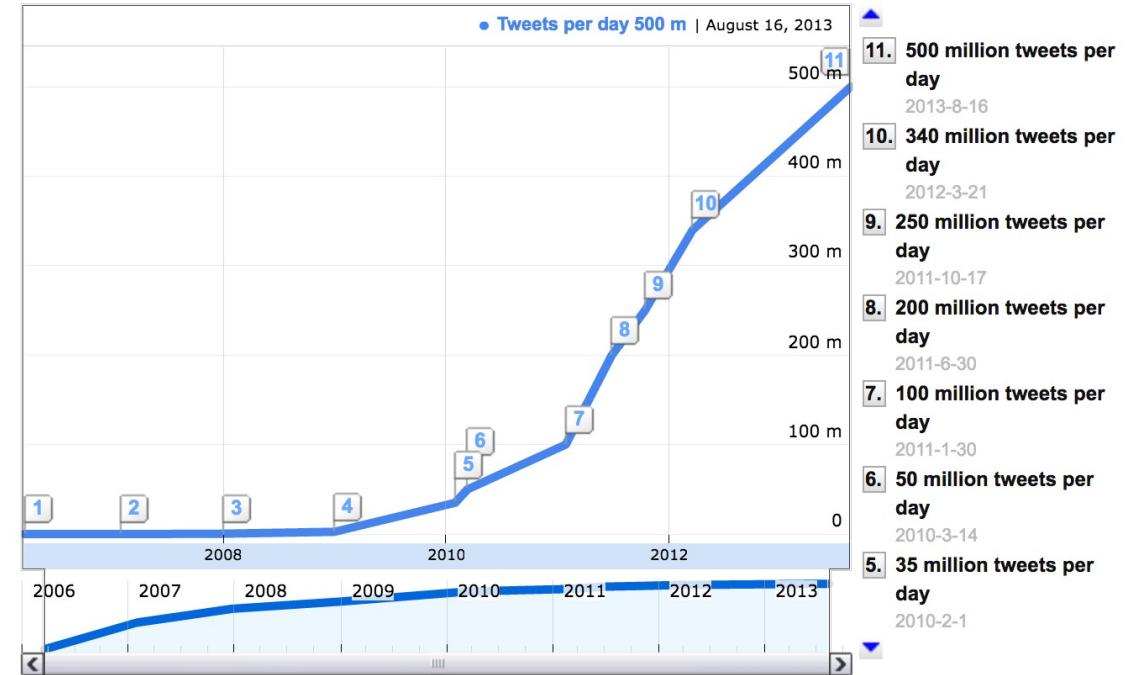


# Searches and Tweets

Google now processes over 40,000 search queries every second on average ([visualize them here](#)), which translates to over **3.5 billion searches per day** and **1.2 trillion searches per year** worldwide. The chart below shows the number of searches per year throughout Google's history:



Every second, on average, around 6,000 tweets are tweeted on Twitter ([visualize them here](#)), which corresponds to over **350,000 tweets sent per minute**, **500 million tweets per day** and around 200 billion tweets per year. The chart below shows the number of tweets per day throughout Twitter's history:

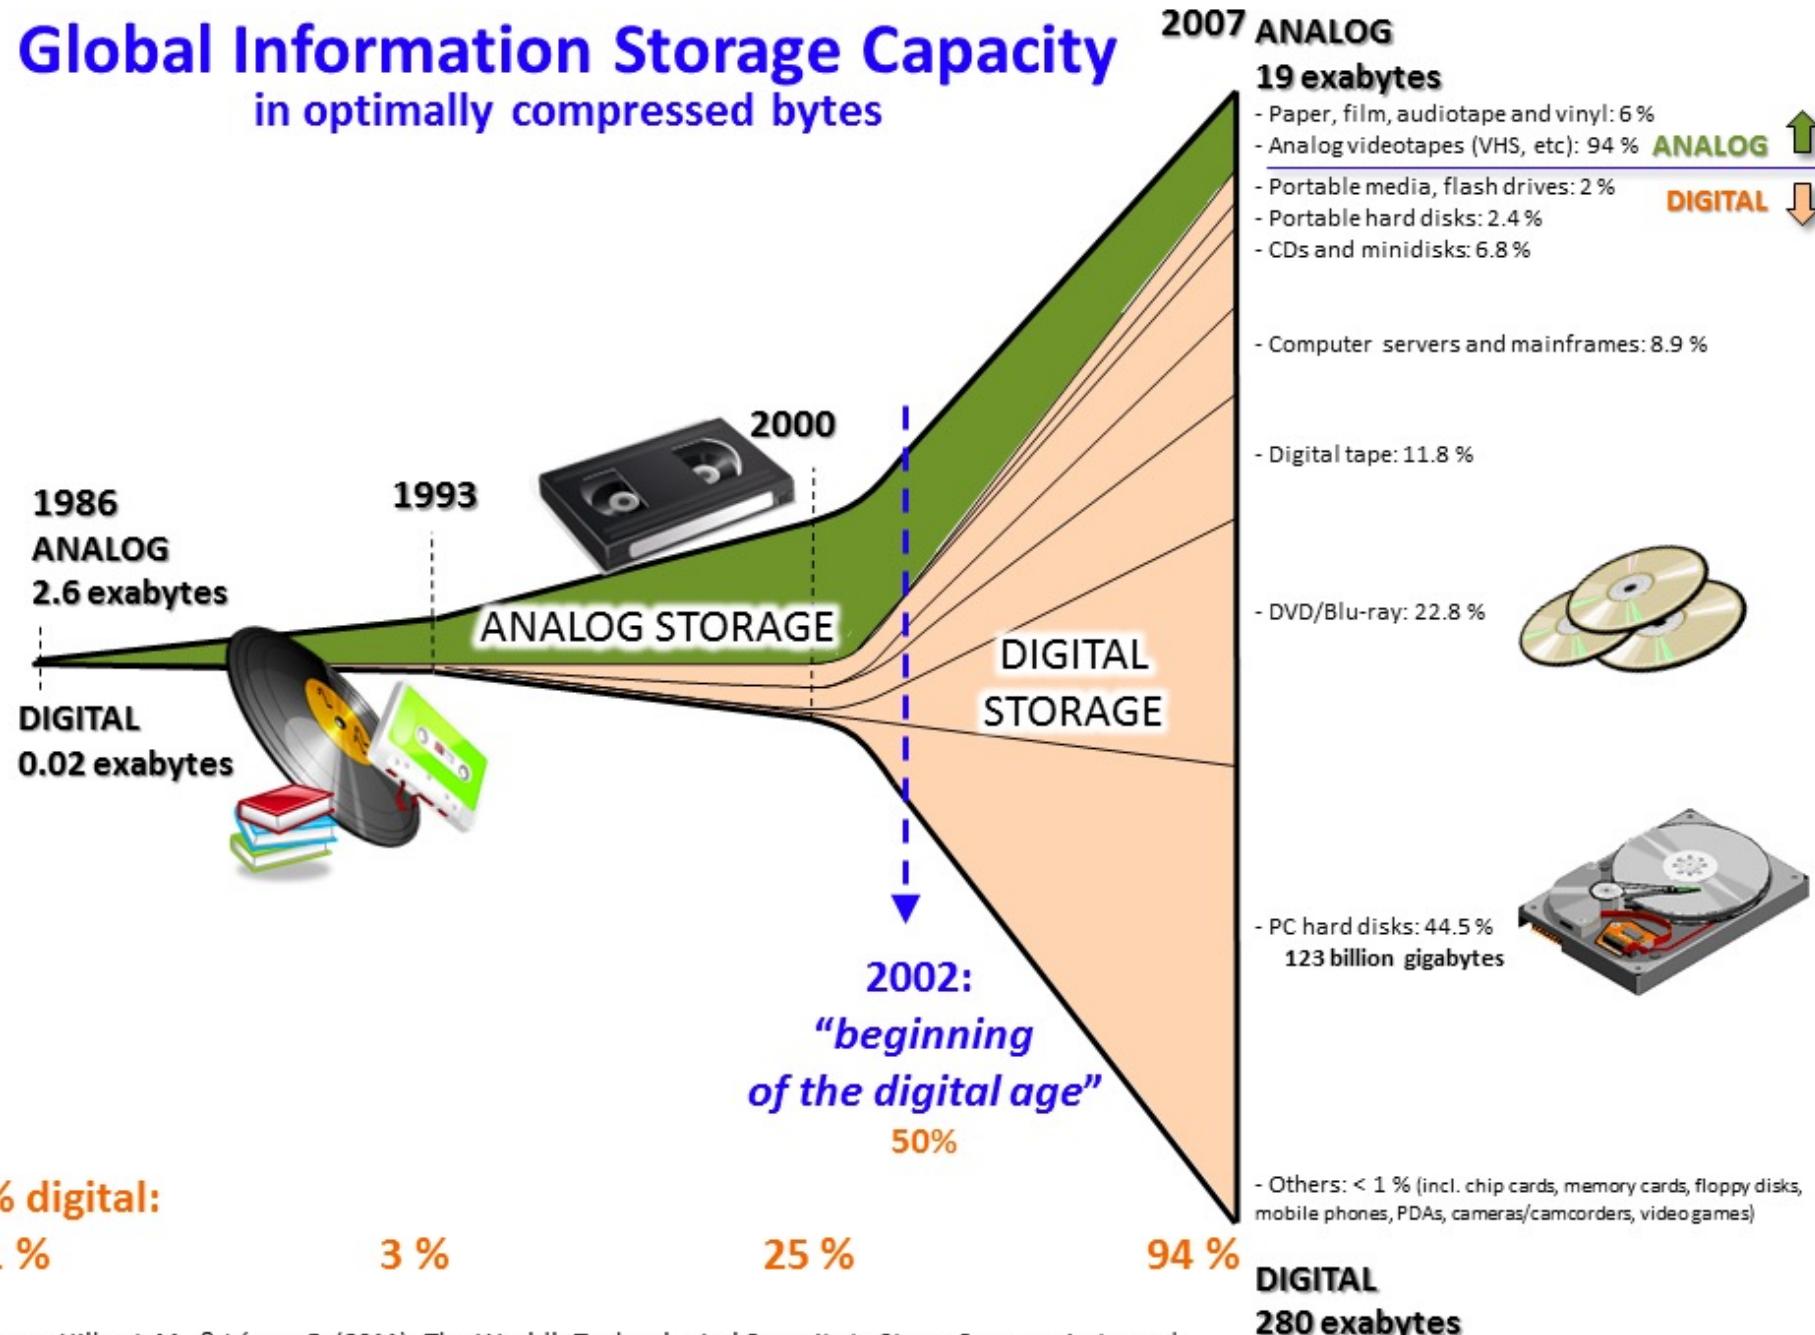


Source: <http://www.internetlivestats.com/>



# Global Information Storage Capacity

in optimally compressed bytes

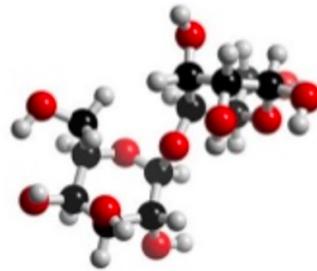


# Bioinformatics Databases

## Sequences



## Structure

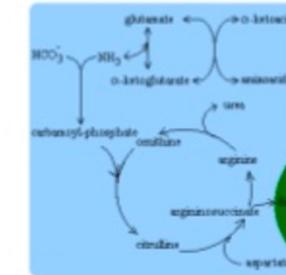


## Genome

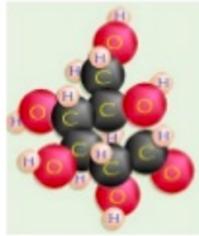
Complete Genome of  
*Arabidopsis thaliana*



## Pathways



## Lipids, Carbohydrates



## Literature



**Primary Database:** Direct experimental results

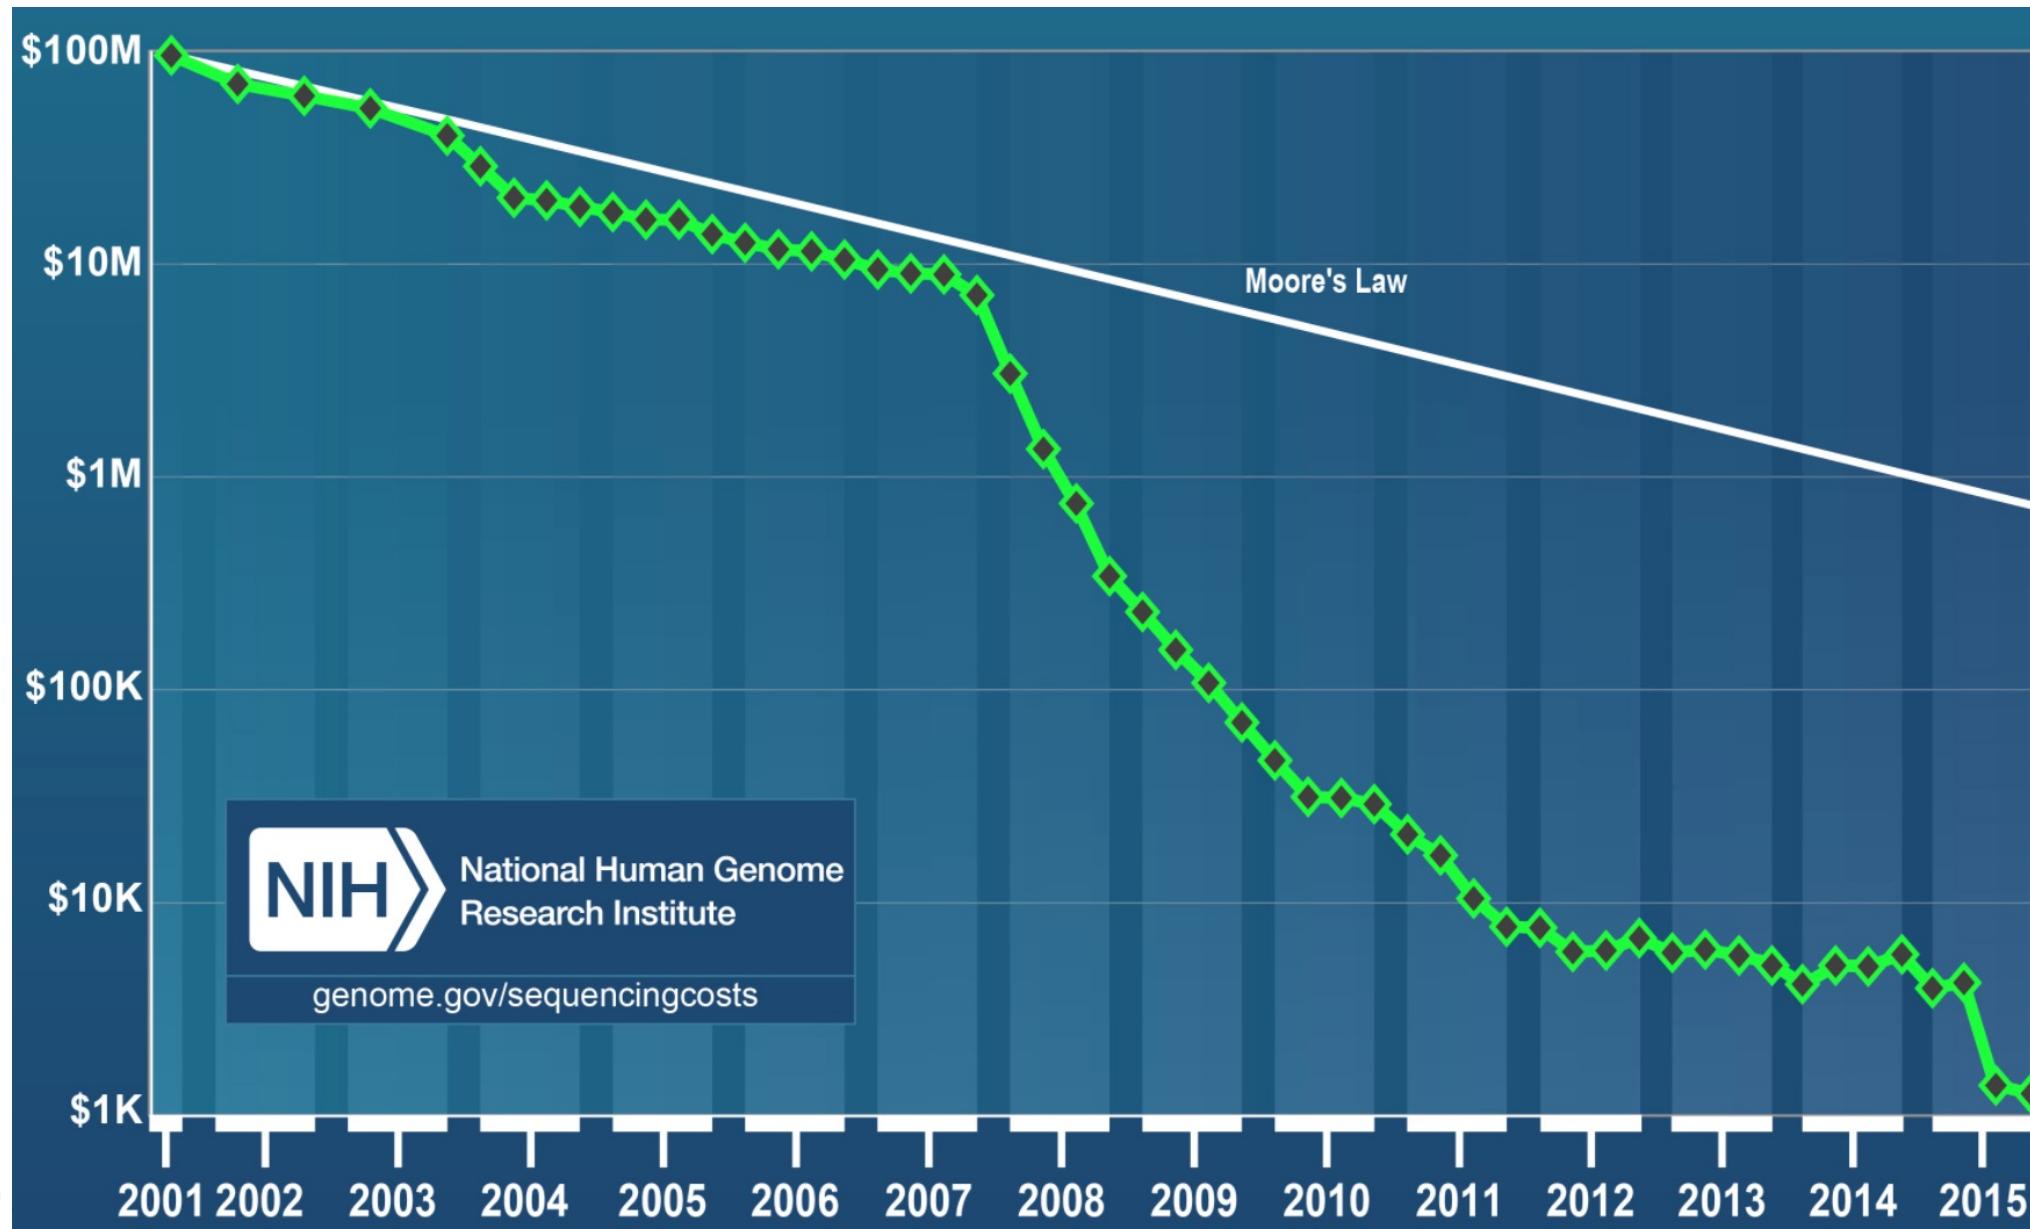
**Secondary Database:** Derived databases from transformation & analysis

Source: <https://www.slideshare.net/shwetakagliwal/biological-databases-11267007>

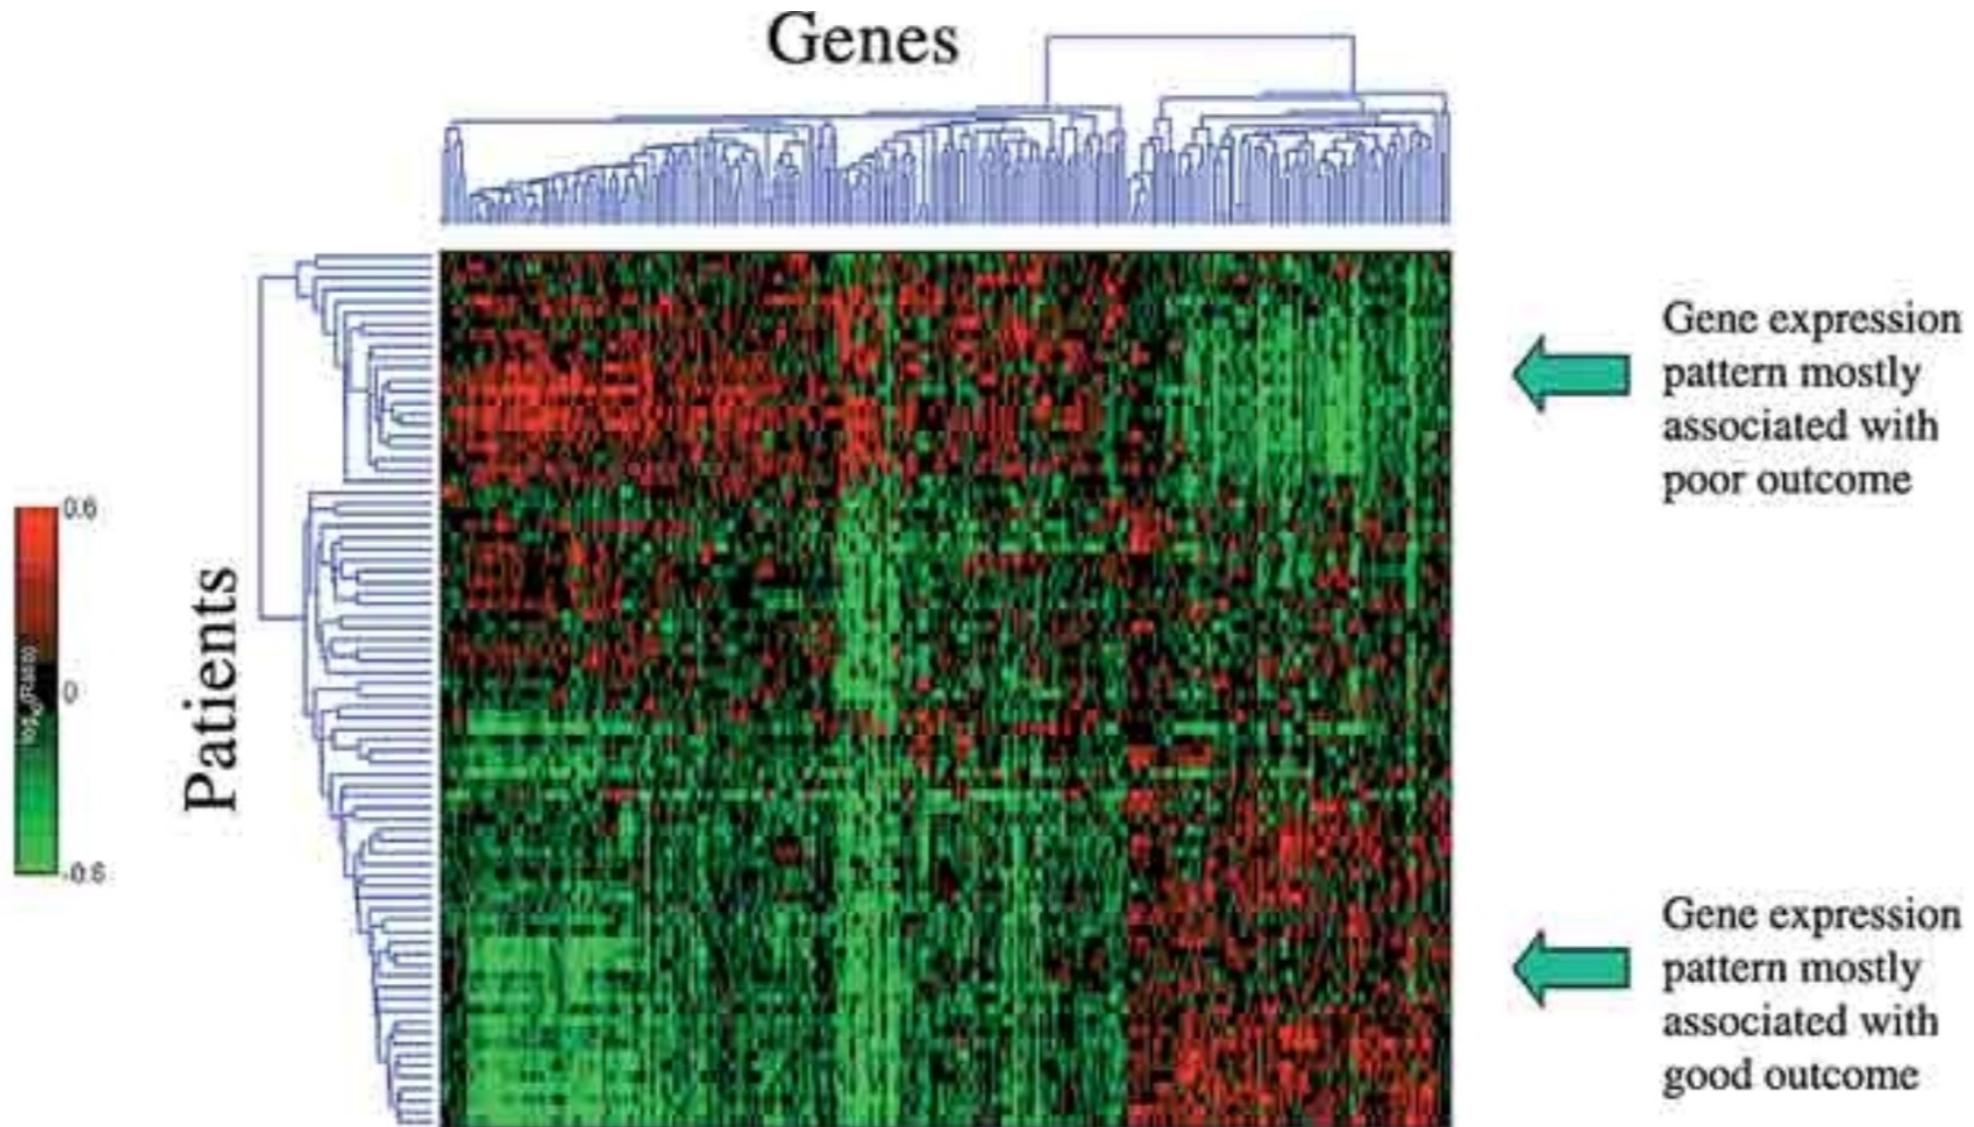


Northeastern University

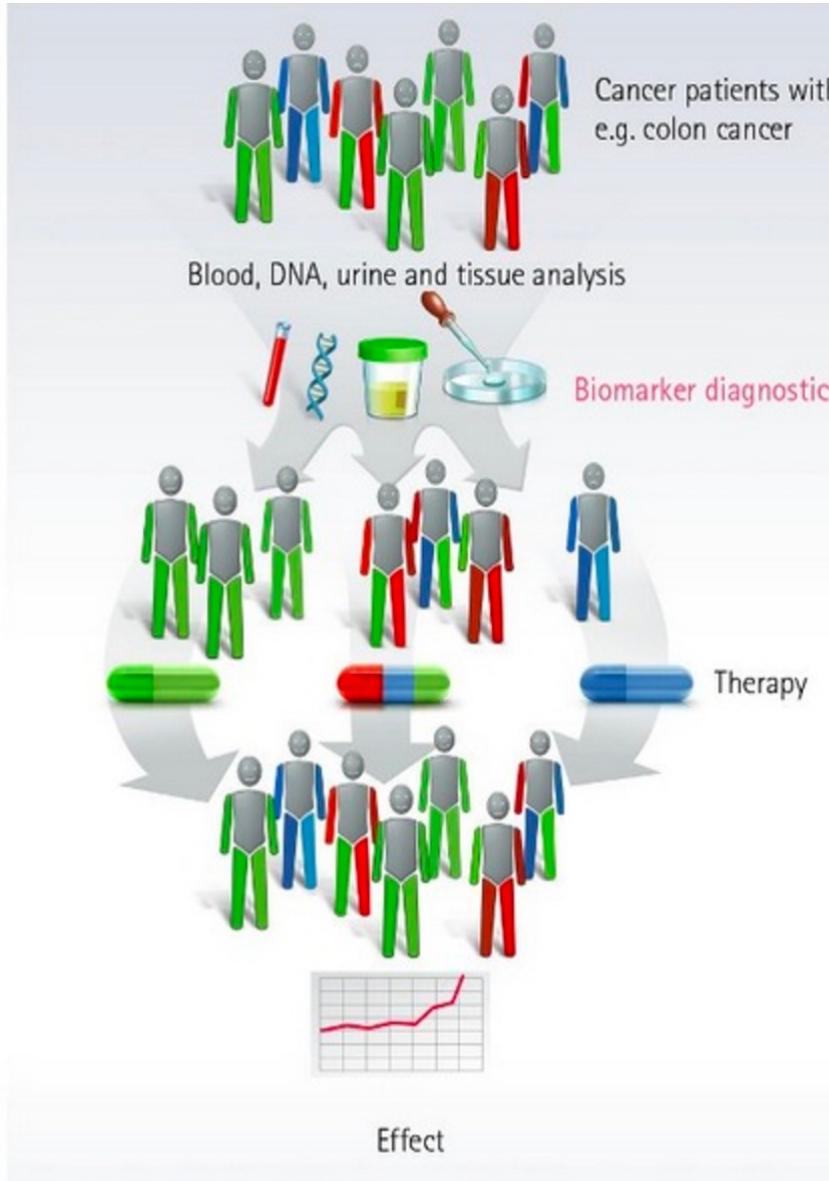
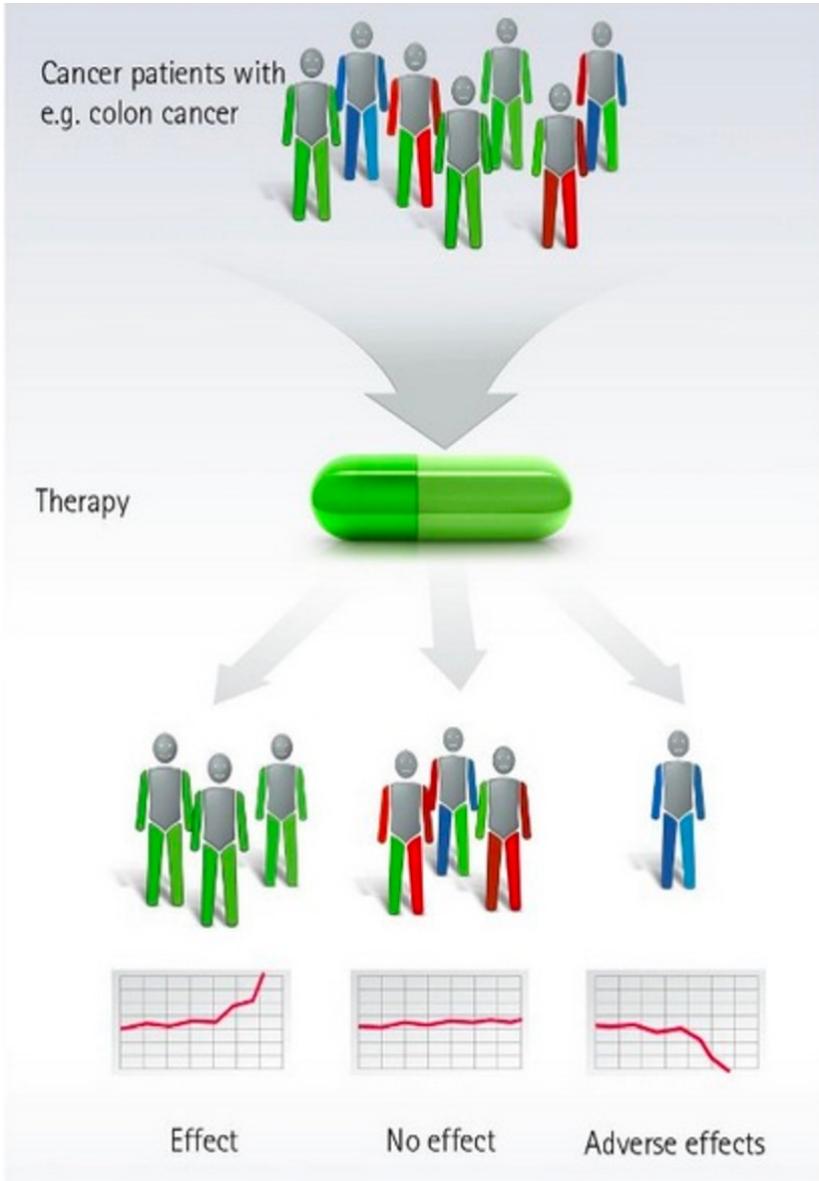
# Cost to sequence a single human genome.



# DNA Microarrays



# The coming age of personalized medicine



The promise of personalized medicine:

- Improved efficacy
- Reduced adverse side-effects
- Reduced “trial-and-error” delays in the treatment of time-critical diseases such as cancer.

# EMR: Electronic Medical Records

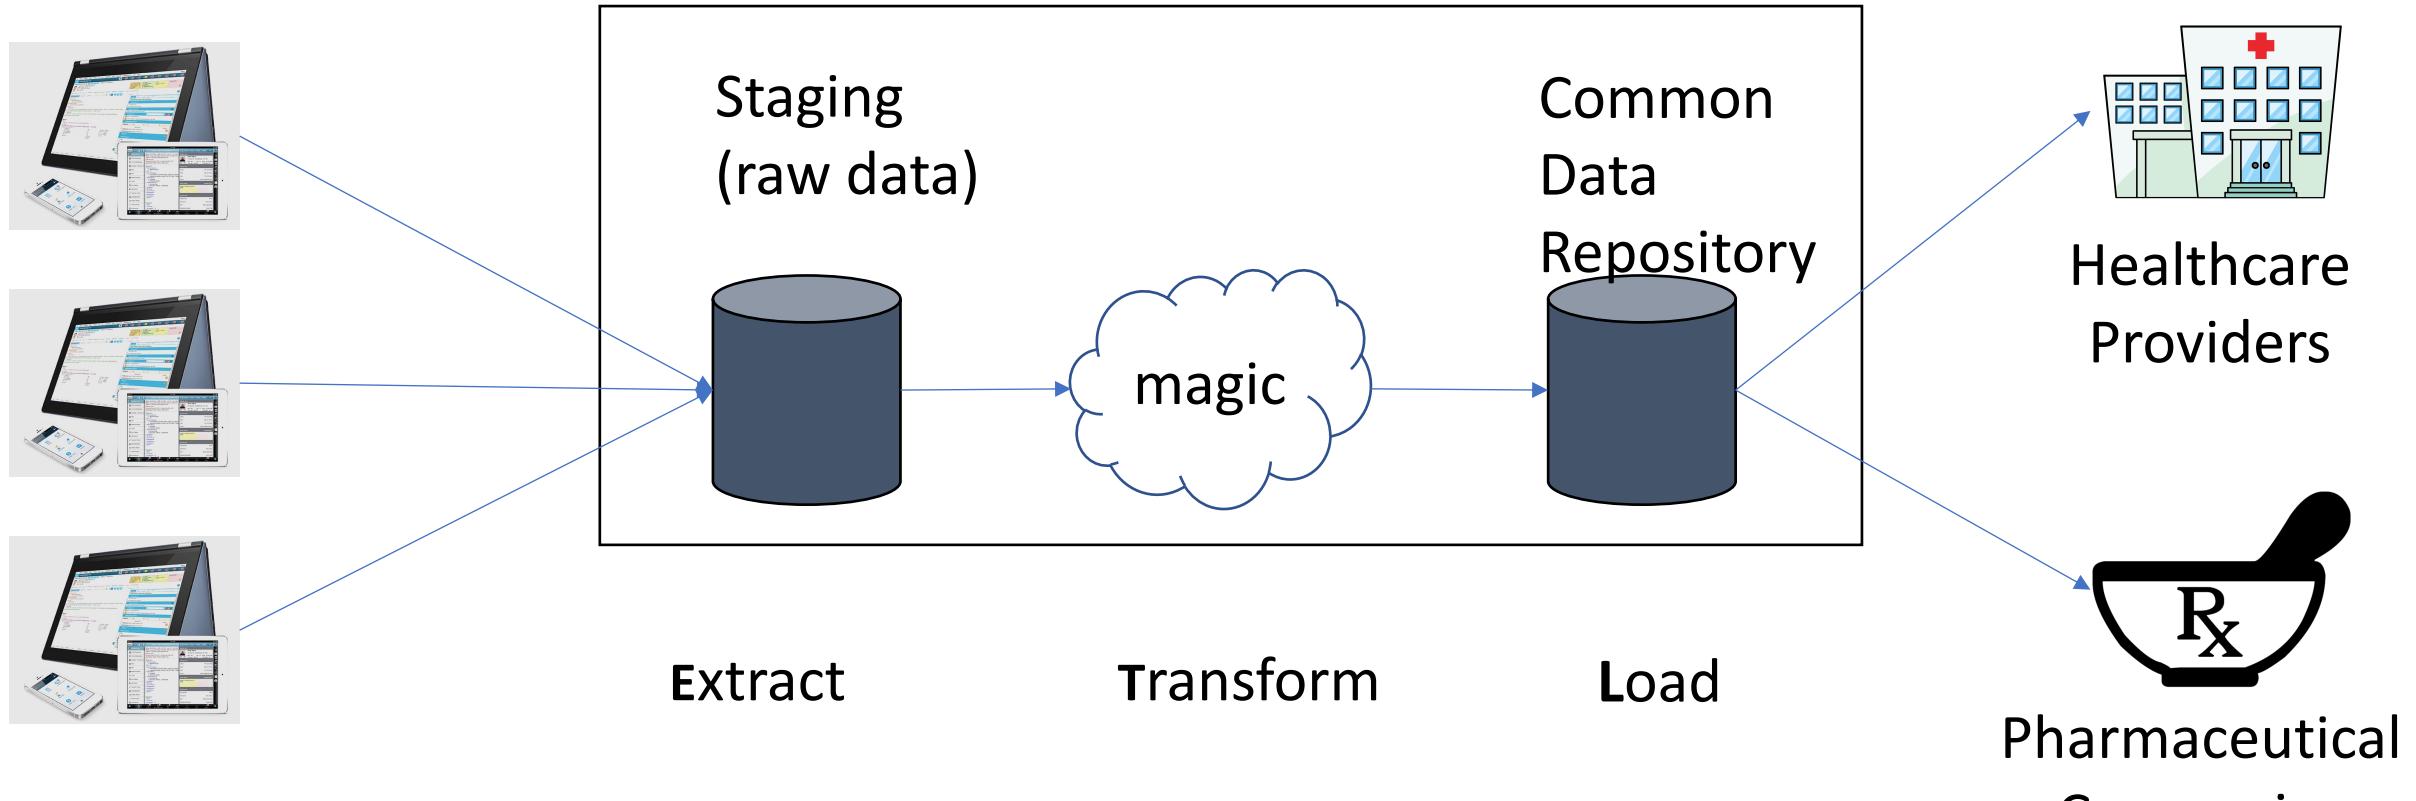


- Patient Demographics
- Doctors
- Insurance
- Medical history
- Allergies
- Procedures / Measurements (BP, Temp, O<sub>2</sub>)
- Order Lab Tests → Lab Results
- Dr. Notes

Next time you get a physical, notice how much time your doctor spends at a computer terminal!



# Optum Analytics: Aggregator of 90+ million patient records



• ***What questions would you ask if you had the medical records for 90+ million patients?***



# Personalized Health / Wearable Sensors



<https://finance.yahoo.com/news/exclusive-fitbits-6-billion-nights-sleep-data-reveals-us-110058417.html>



# The dawn of big data presented new challenges

*Data assets are [a] major component of the balance sheet, replacing traditional physical assets of the 20th century.*

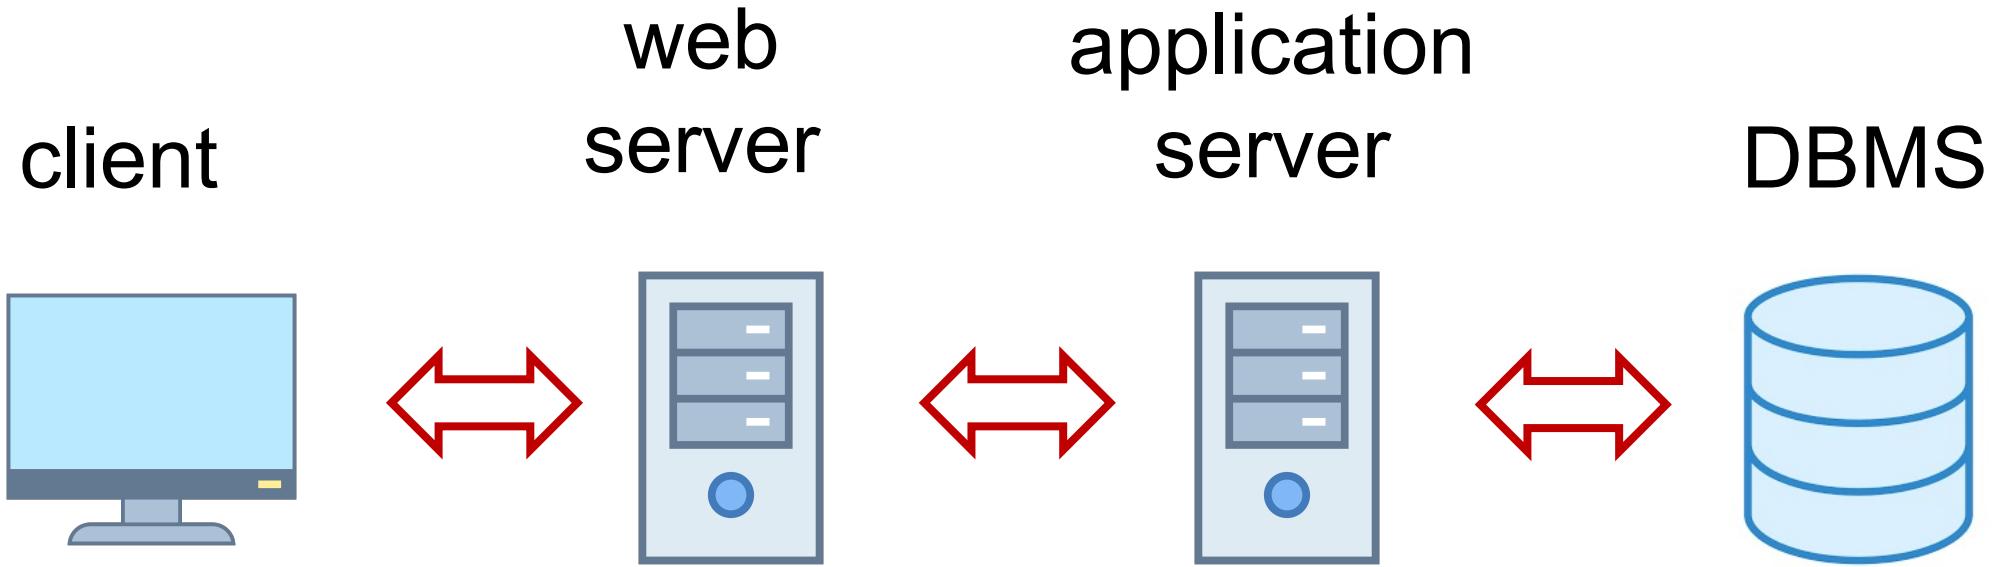
-- Dr. Ralph Kimball: *Rethinking EDW in the era of expansive data management*

- In the Internet age, we expect *instantaneous* answers to *any* question: medical, financial, scientific, products & services, and turkey recipes!
- It is no longer ok to prune data, say, for the last N days. (We can build better machine learning models, recommendation engines, targeted ads, and predictions if we have *all* of the historical data.)



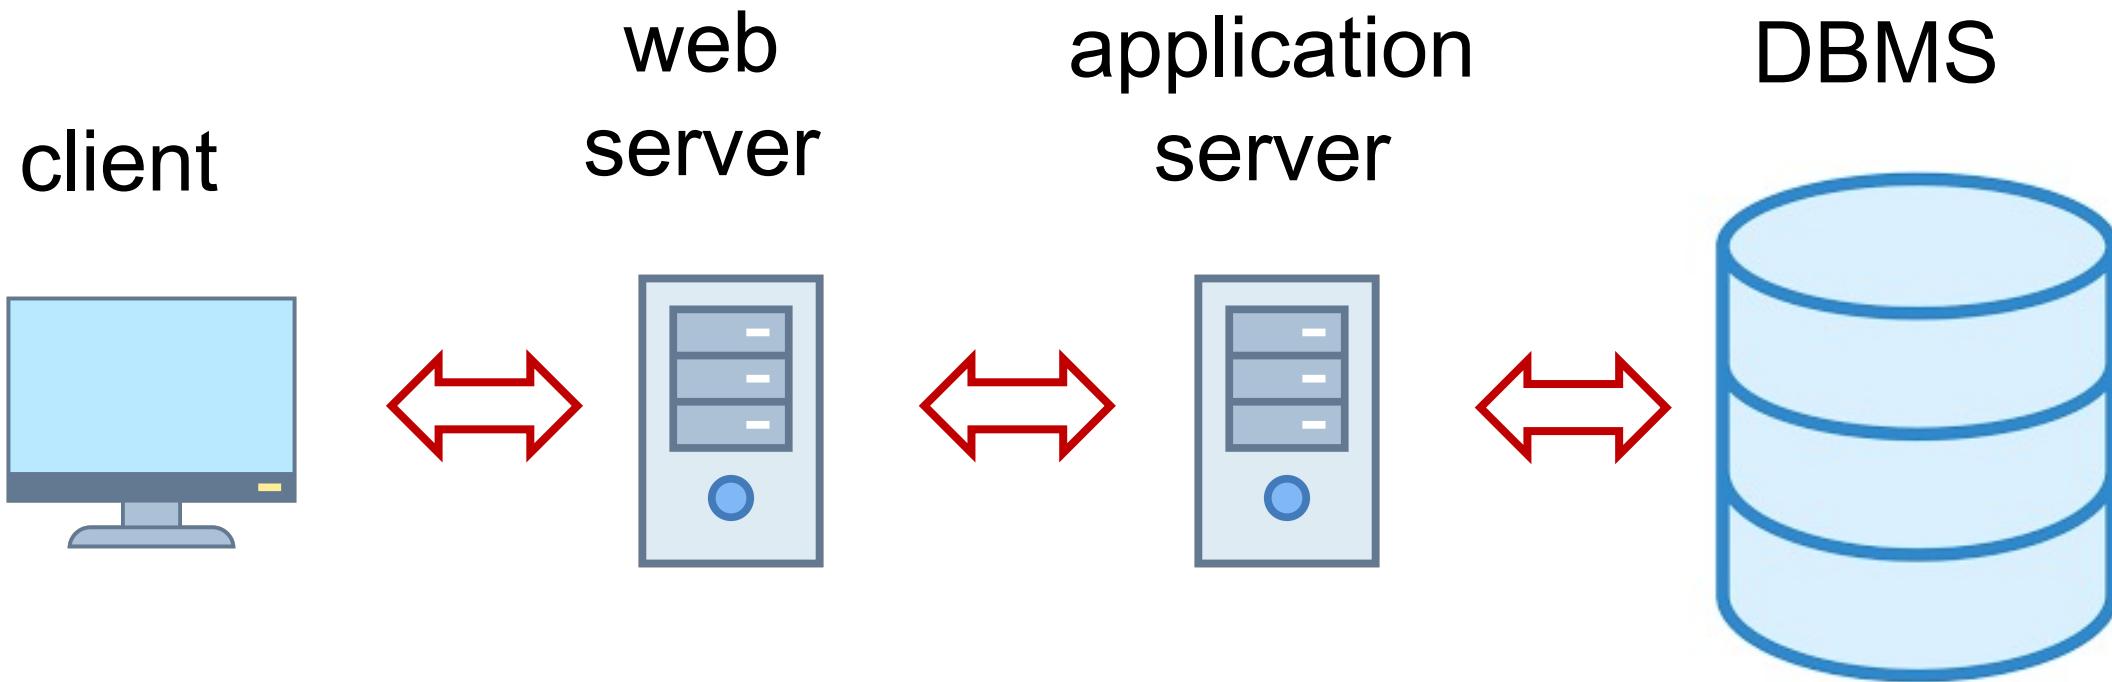
# Early stages of a company. N-tier client-server is adequate

---



# Early stage growth. Vertical scaling

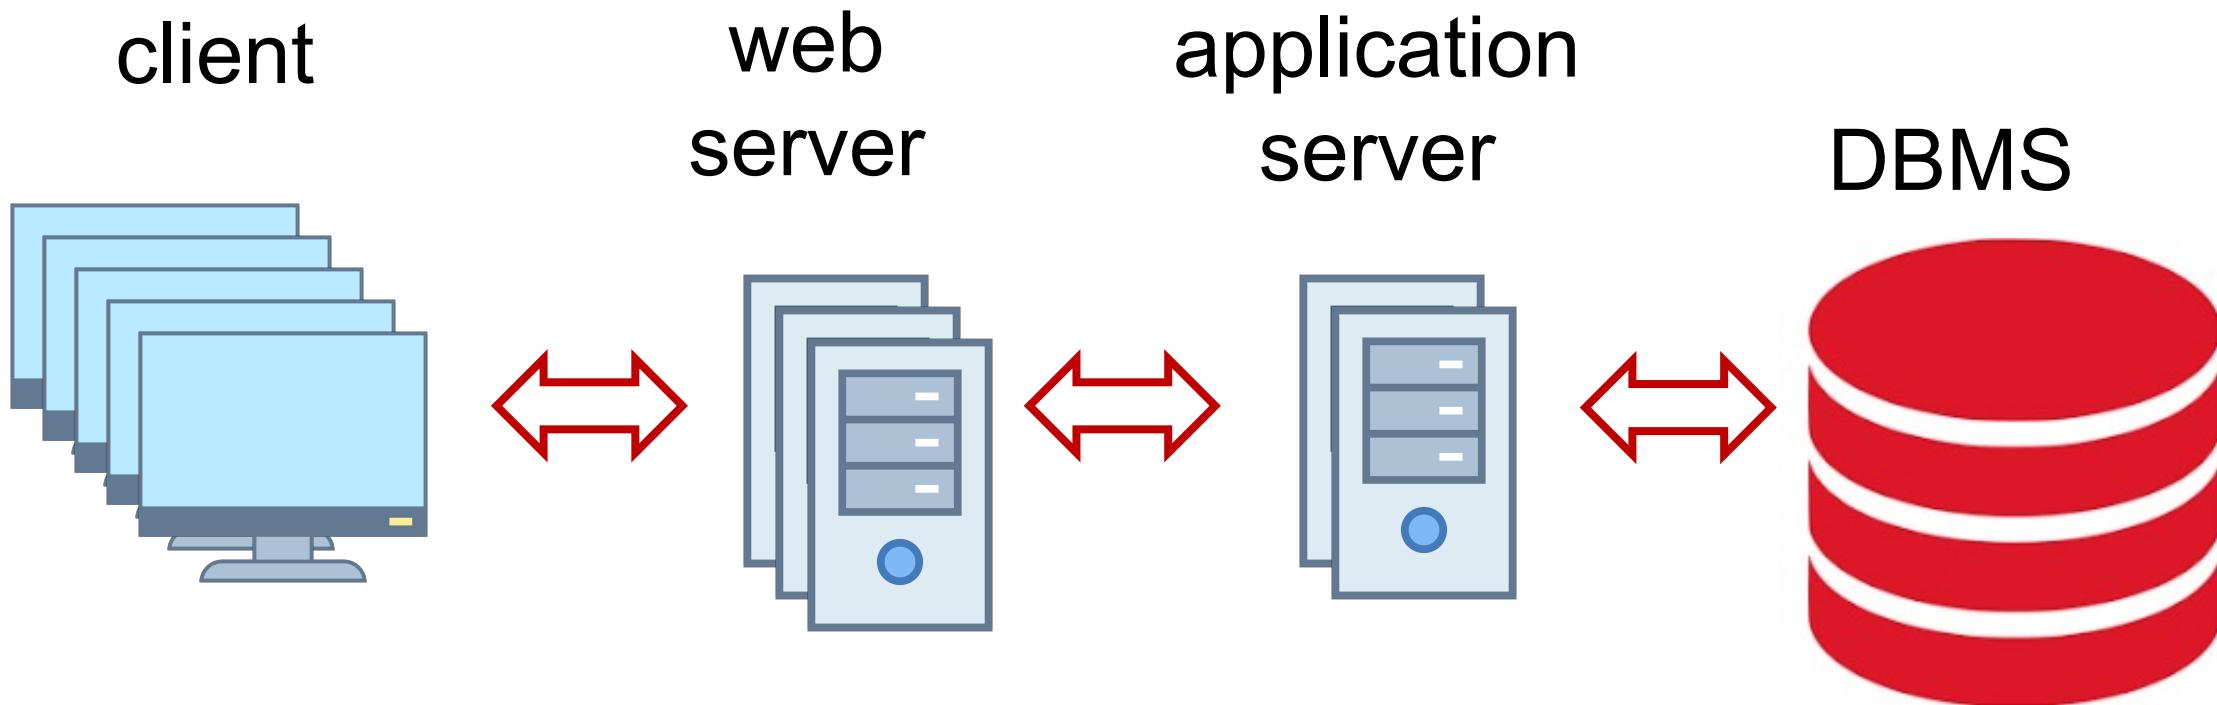
---



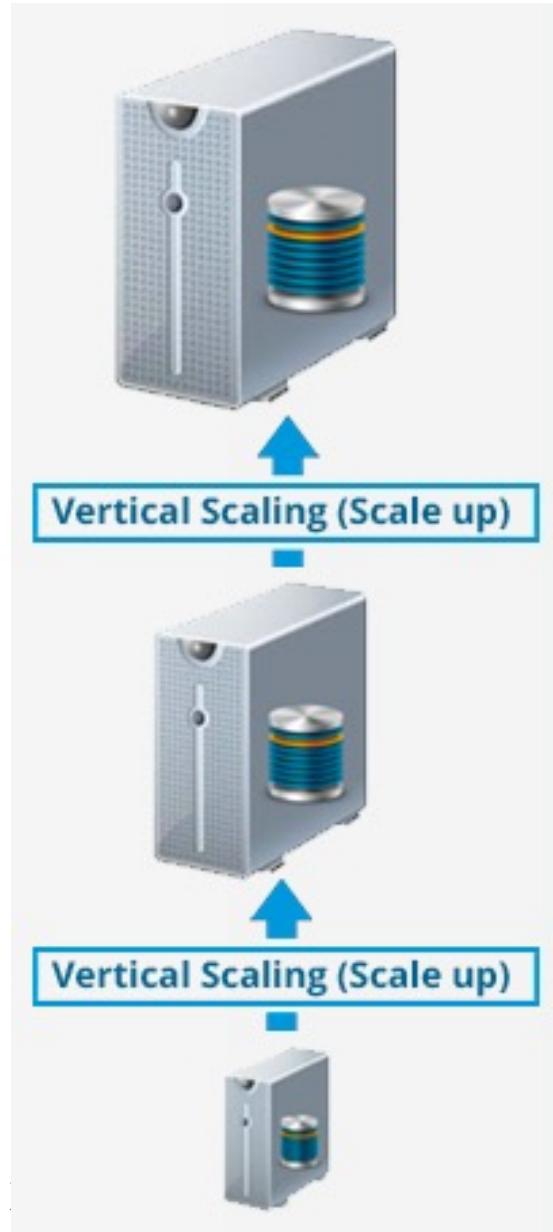
Beef up your DBMS: More RAM, more cores, faster disks, ***more \$\$\$***



# Growing user base supported via load balancing



# Scaling Up: Advantages and Disadvantages



## ADVANTAGES

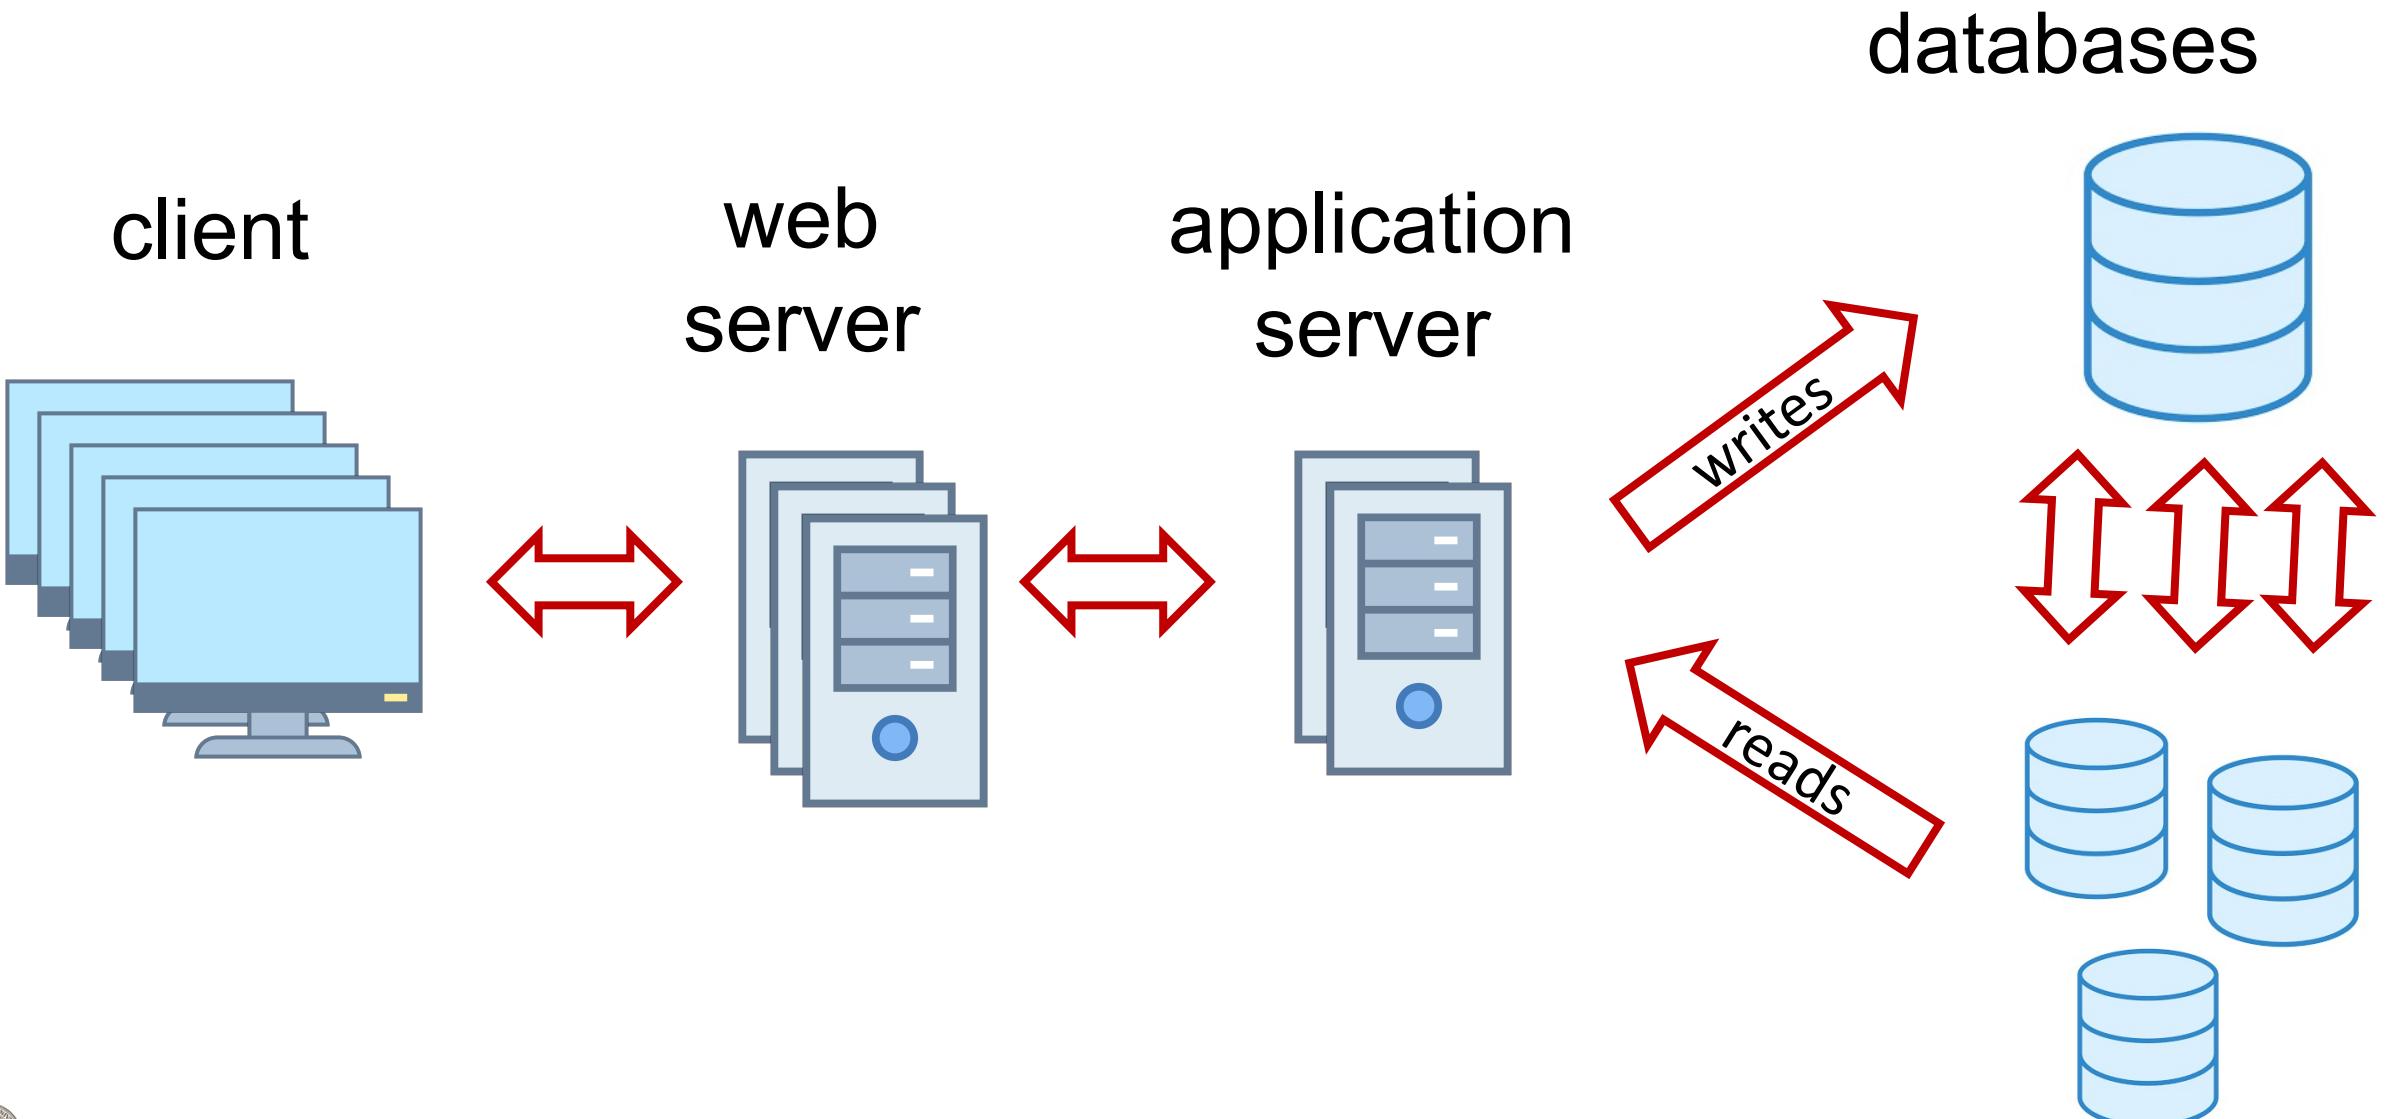
- Leveraging a platform you already know
- Software migration is easier
- Administration is easier
- Relatively easy to manage: add more RAM, CPU or buy a new server if necessary
- Cooling overhead is lower
- Networking hardware overhead is lower

## DISADVANTAGES

- Exponential server cost and licensing costs (~10's of millions \$)
- Vendor lock-in
- Single point of failure
- Limits on available hardware expansion
- Reliance on specialty human resources

*Once you are in the realm of “big data” you need another solution.*

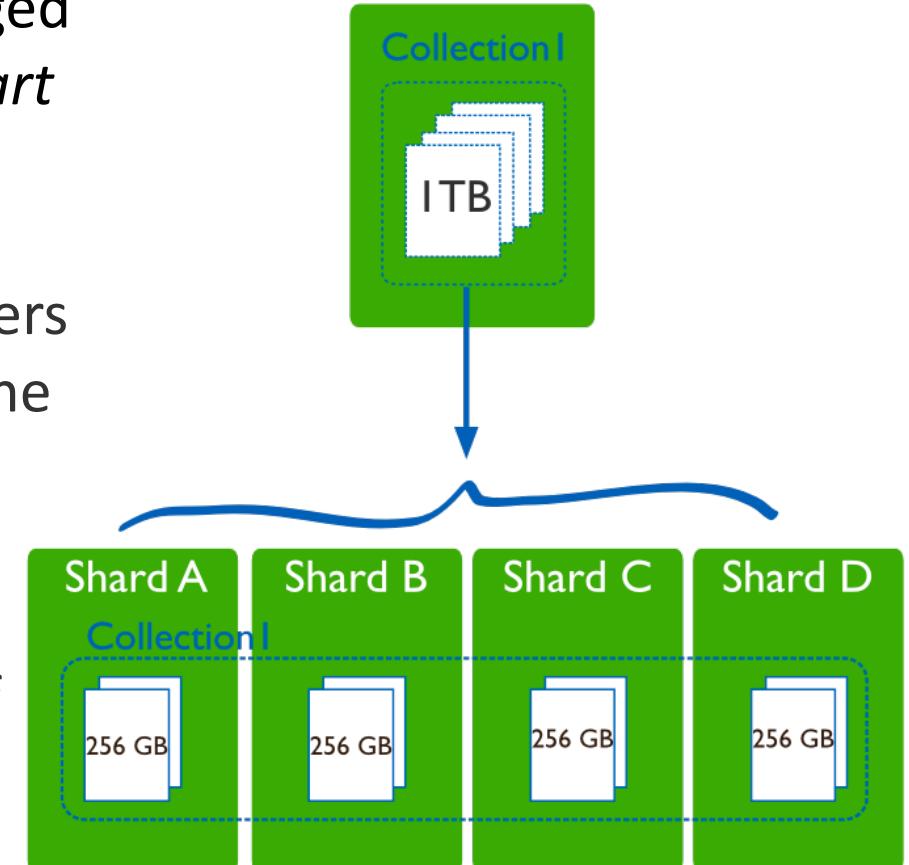
# Replication: Adding secondary DB servers to keep up with demand



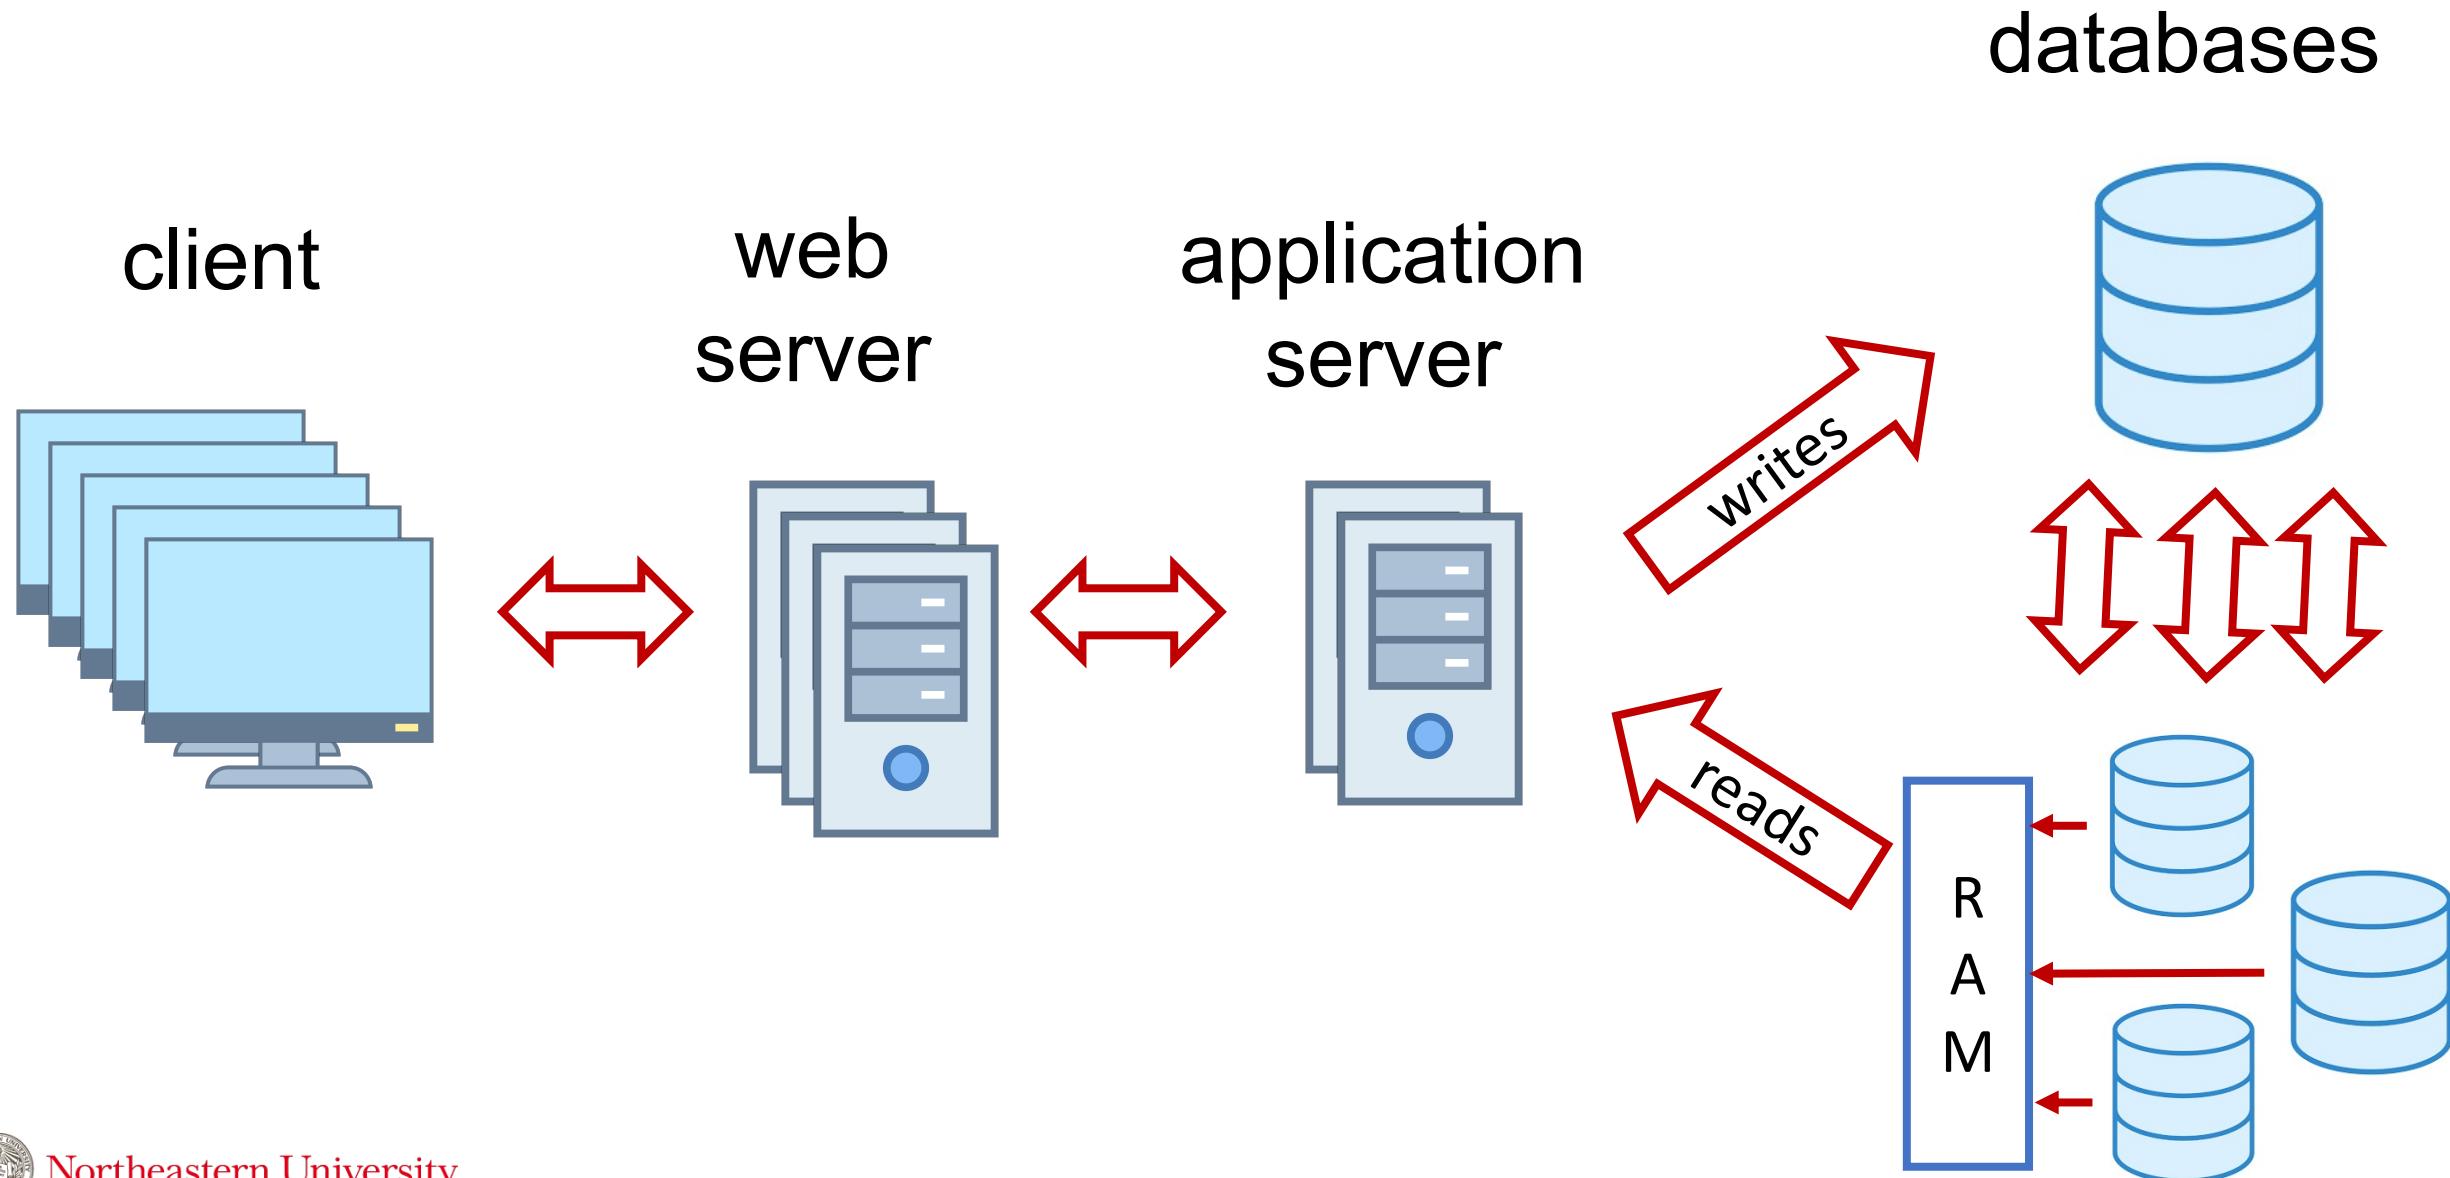
# Sharding: Trading application simplicity for speed

**Sharding** is a type of database partitioning that separates very large databases into smaller, faster, more easily managed parts called data **shards**. The word shard means *a small part of a whole*.

- Relational databases could also be run as separate servers for different shards (sets of data), effectively sharding the database.
- While this separates the load, all the sharding has to be controlled by the application which has to keep track of which database server to talk to for each bit of data.
- *Also, we lose any querying, referential integrity, transactions, or consistency controls that cross shards.*

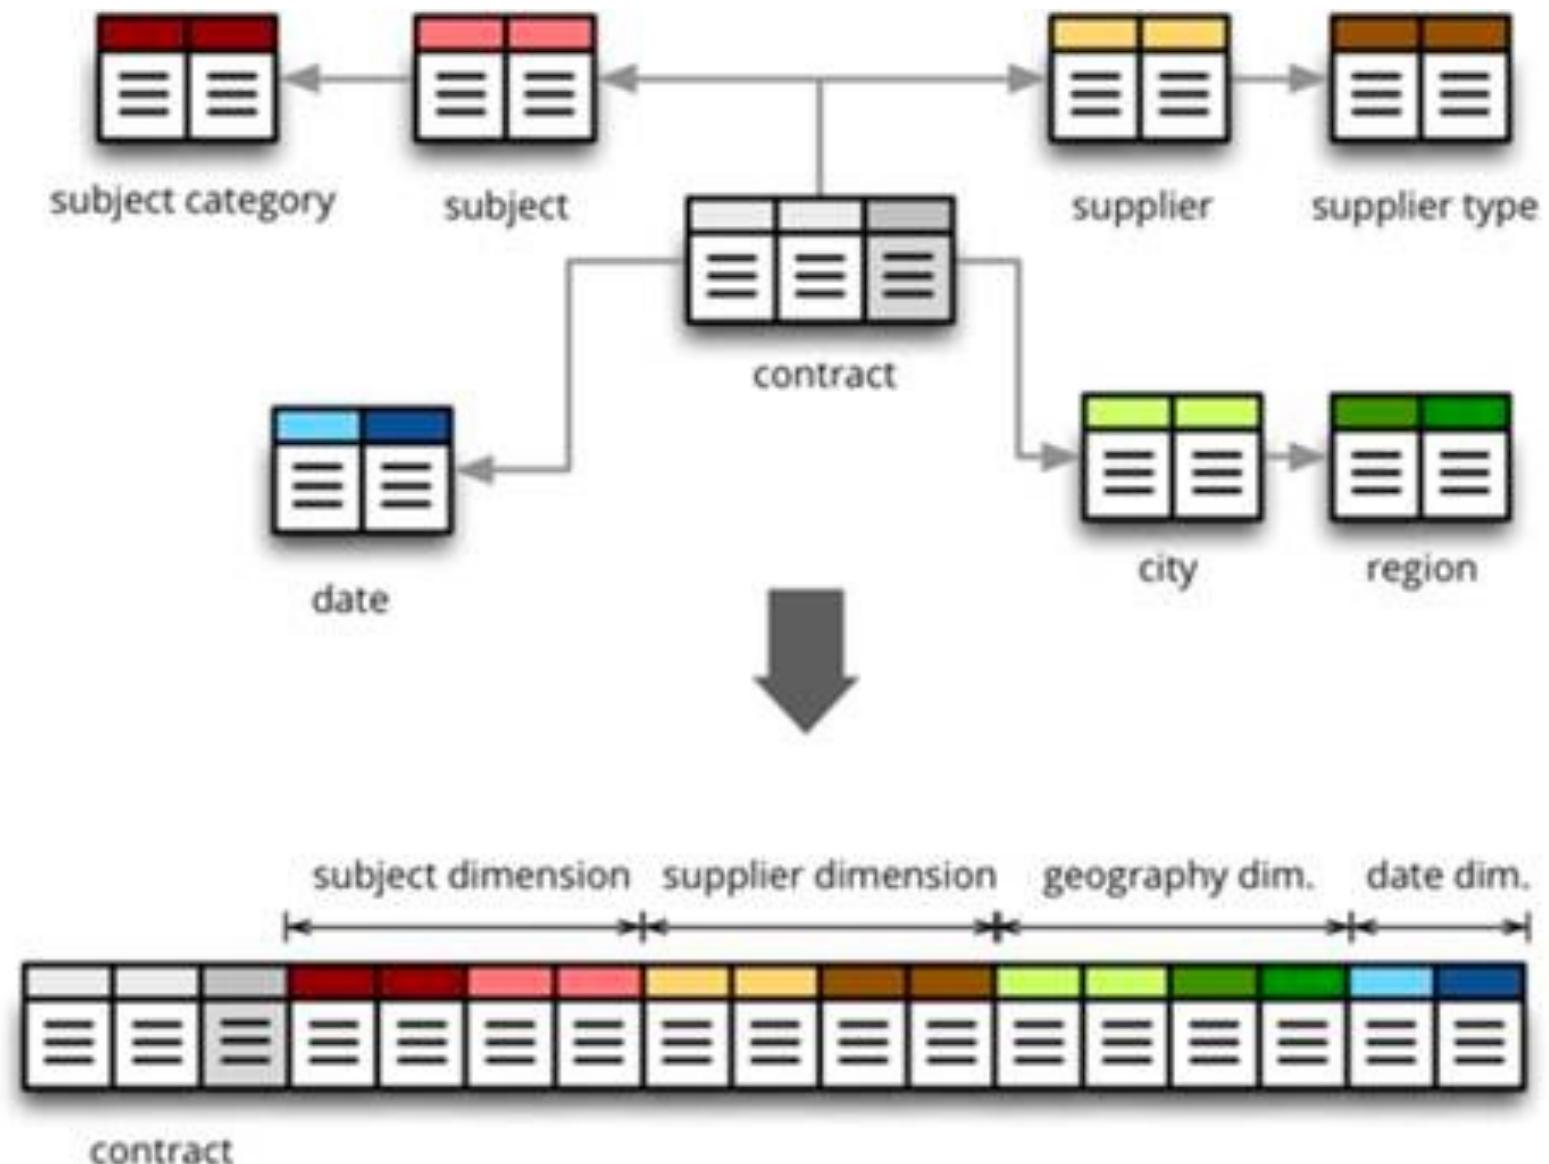


# Memcached: trading consistency for speed.



# Denormalization: Trading write performance for speed

- We also incur storage overhead
- Complex JOINs are simplified or eliminated making queries much faster



# More desperate measures

---

- Eliminating stored procedures in favor of off-line processing
- Eliminating secondary indexes which slow down inserts and updates, limiting access to data via primary keys only
- **Materialized views:** To support efficient querying, a common solution is to generate, in advance, a view that materializes the data in a format suited to the required results set.



# The Benefits of Relational Databases

- Data persistence via a well-defined schema (without having to know the details of how the data is organized on disk)
- SQL – The query standard (with many dialects). It is both expressive and flexible.
- Data integration: Data can be combined via JOINS for deeper insights
- Reporting / Aggregate analysis
- ACID Transactions

Branch

branchNo	street	city	postCode
B005	22 Deer Rd	London	SW1 4EH
B007	16 Argyll St	Aberdeen	AB2 3SU
B003	163 Main St	Glasgow	G11 9QX
B004	32 Manse Rd	Bristol	BS99 1NZ
B002	56 Clover Dr	London	NW10 6EU

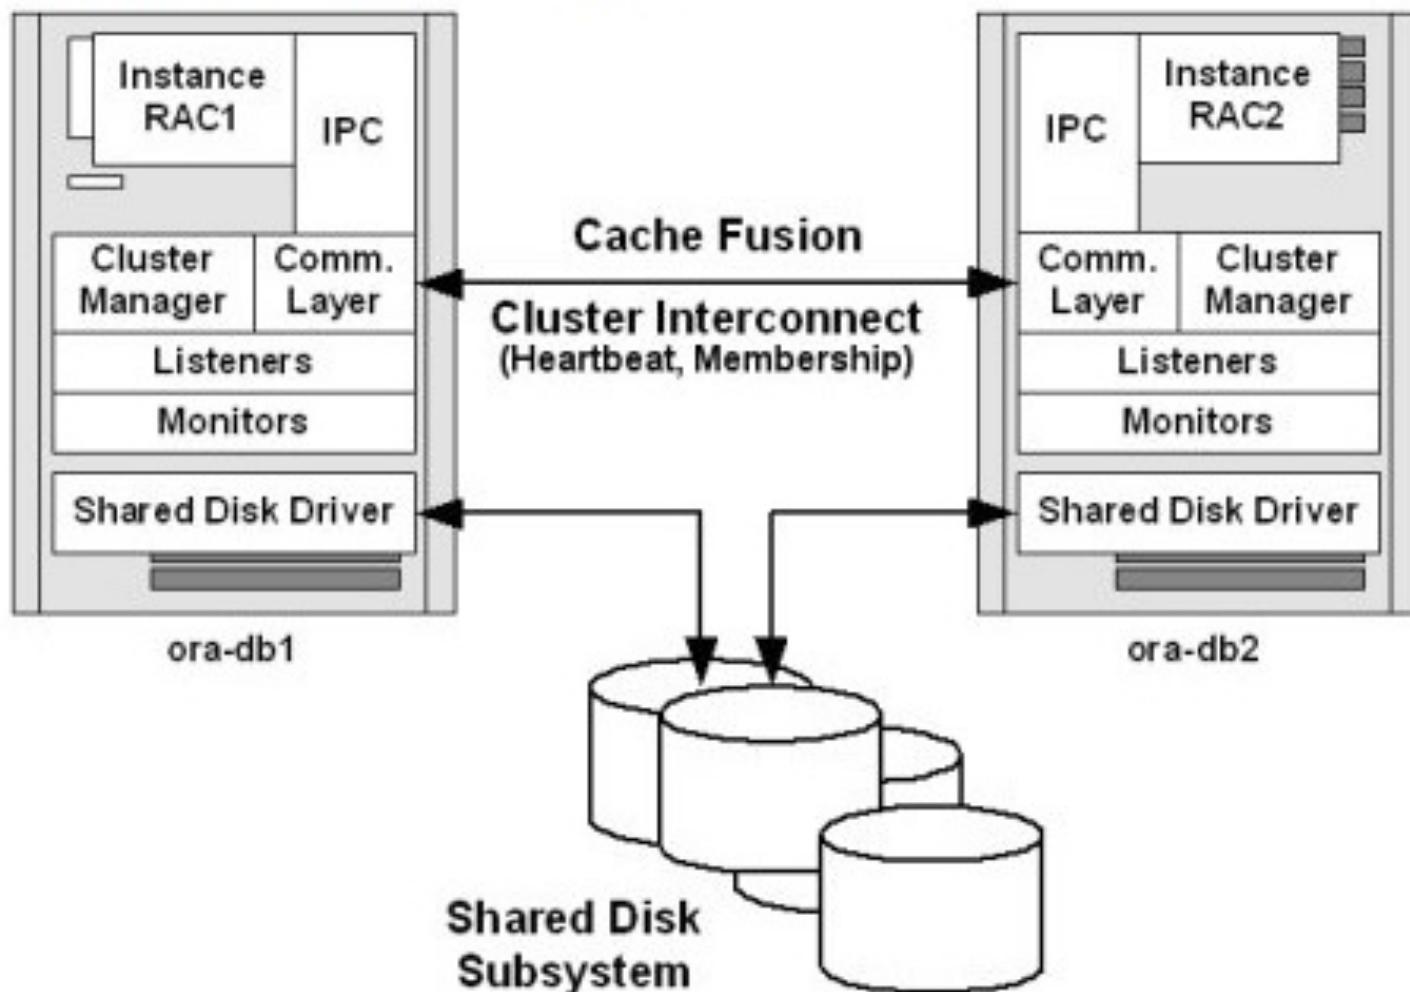
Staff

staffNo	fName	IName	position	sex	DOB	salary	branchNo
SL21	John	White	Manager	M	1-Oct-45	30000	B005
SG37	Ann	Beech	Assistant	F	10-Nov-60	12000	B003
SG14	David	Ford	Supervisor	M	24-Mar-58	18000	B003
SA9	Mary	Howe	Assistant	F	19-Feb-70	9000	B007
SG5	Susan	Brand	Manager	F	3-Jun-40	24000	B003
SL41	Julie	Lee	Assistant	F	13-Jun-65	9000	B005



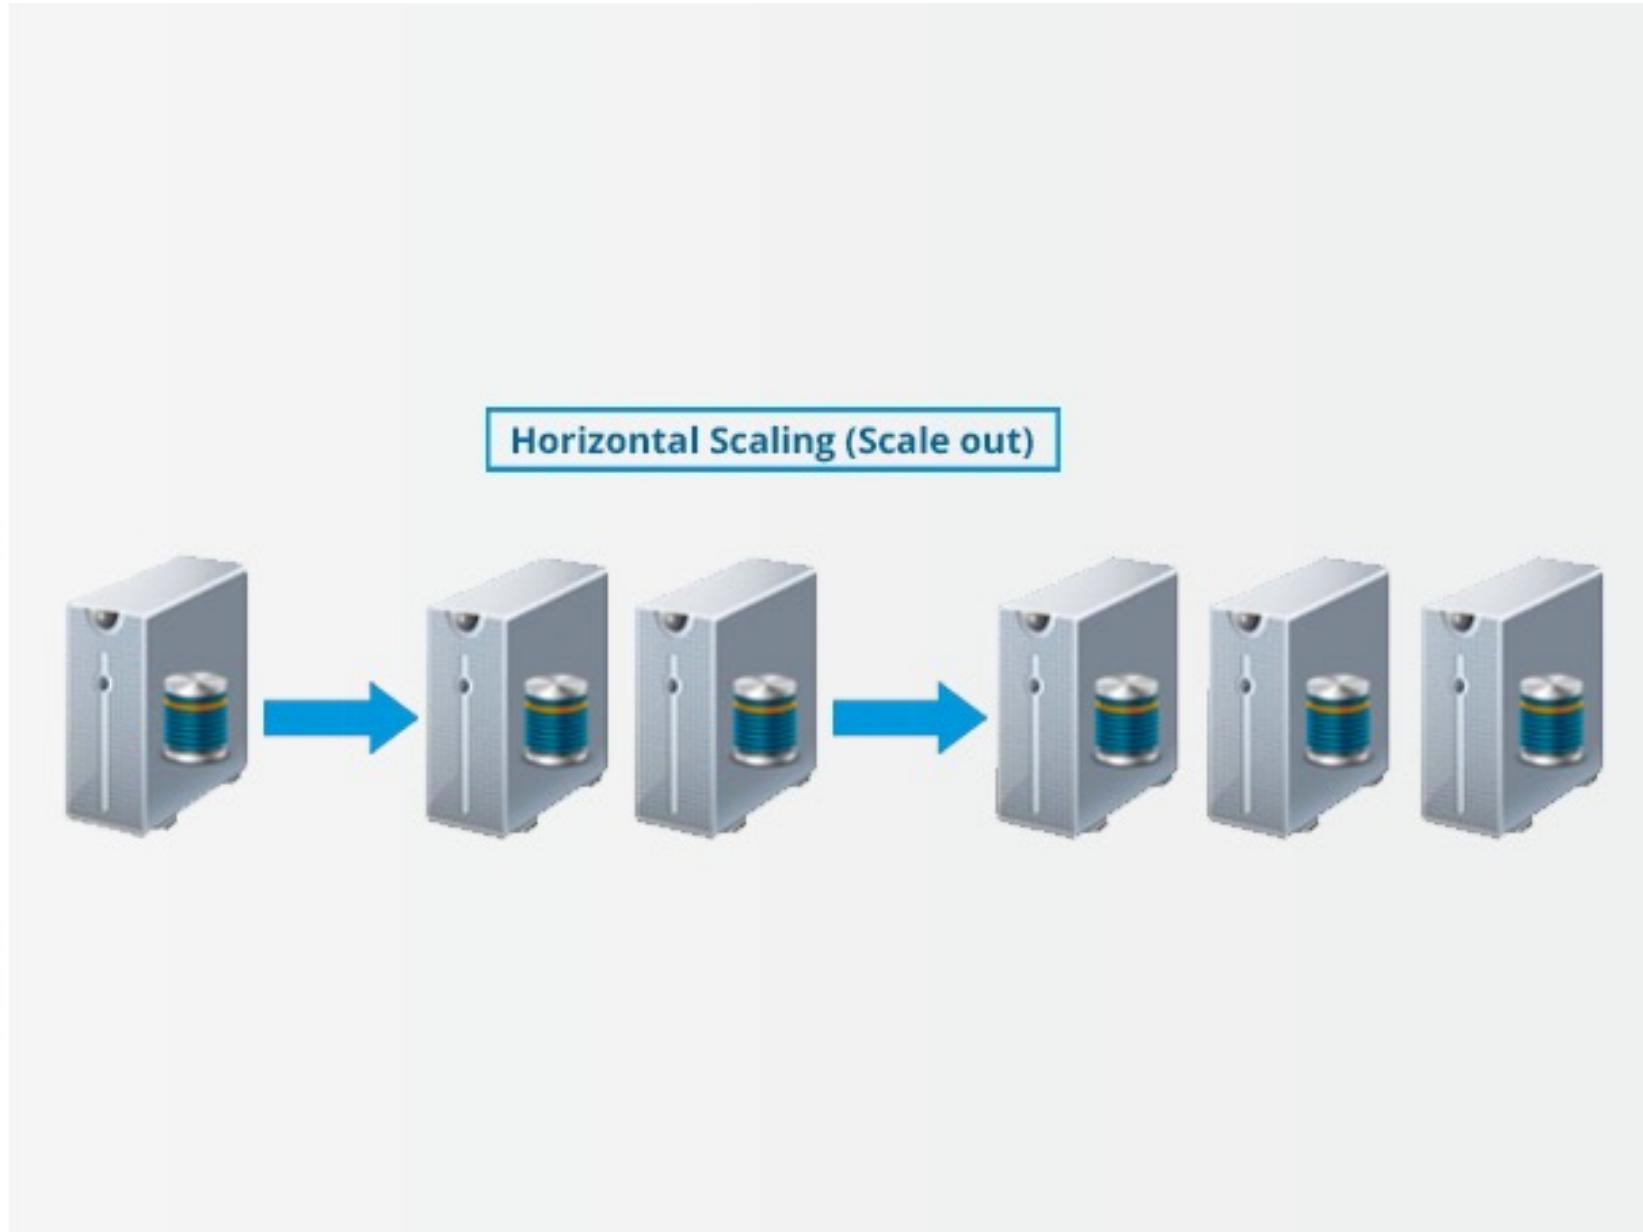
# Scale up: Oracle Real Application Cluster (RAC)

The following diagram, adapted from *Oracle Real Application Clusters* by Murali Vallath [Elsevier Digital Press, 2004] illustrates some of the basic components of an Oracle RAC cluster under Oracle 9i.



Oracle Real Application Cluster (RAC) is a single database that is shared by multiple Oracle instances. It provides opportunities for improved database availability, performance and scalability. Oracle RAC lowers the TCO of a high-end database and IT assets, since a complex Oracle database can be set-up on low-cost, standard, modular parts. Oracle RAC has matured into a true enterprise, scalable solution, but the adoption curve and skill-set required to maintain RAC is relatively slow, as most organizations who have installed RAC only use it as a two-node cluster.

# Scaling Up vs. Scaling Out



# Advantages and disadvantages of *Scaling out*

---

## Advantages

- Horizontal scaling using commodity hardware
- Greater reliability

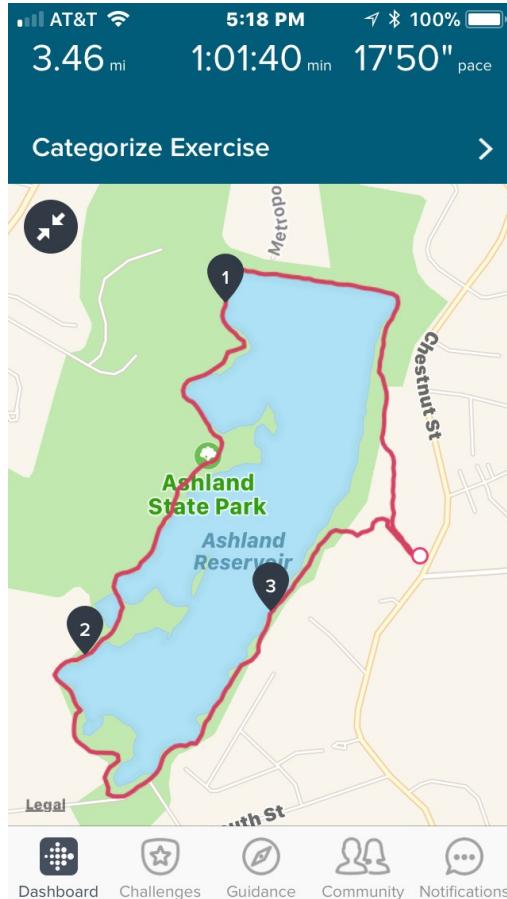
## Disadvantages

- Incompatible with existing relational storage models
- Relational DBs that introduce cluster-based mechanisms have high licensing costs
- The computing models are more complicated
- Security issues may be more complicated or simply less mature
- The technologies are more cutting-edge and less mature making it more challenging to find answers to problems you may encounter.



# Finding GPS Anomalies

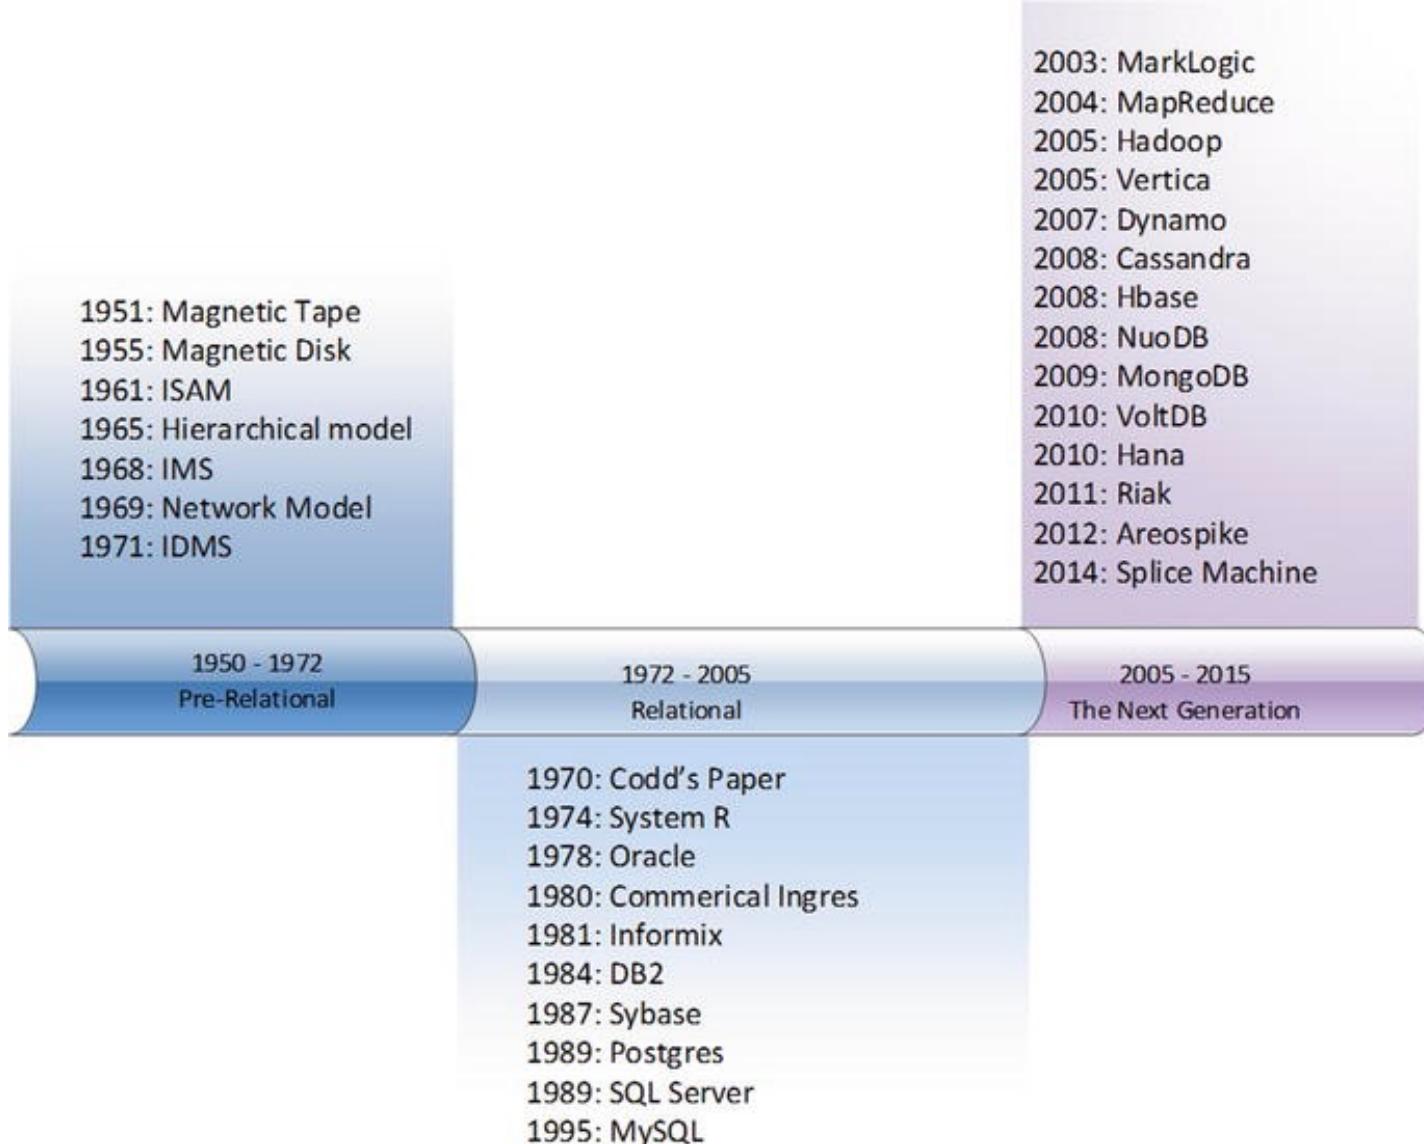
How would you parallelize this?



anomalies = 0  
for each user  
for each exercise  
for each gps measurement  
 $\text{distance} = \text{gps.current} - \text{gps.last}$   
if ( $\text{distance} > 10\text{m}$ ) anomalies++  
  
return anomalies

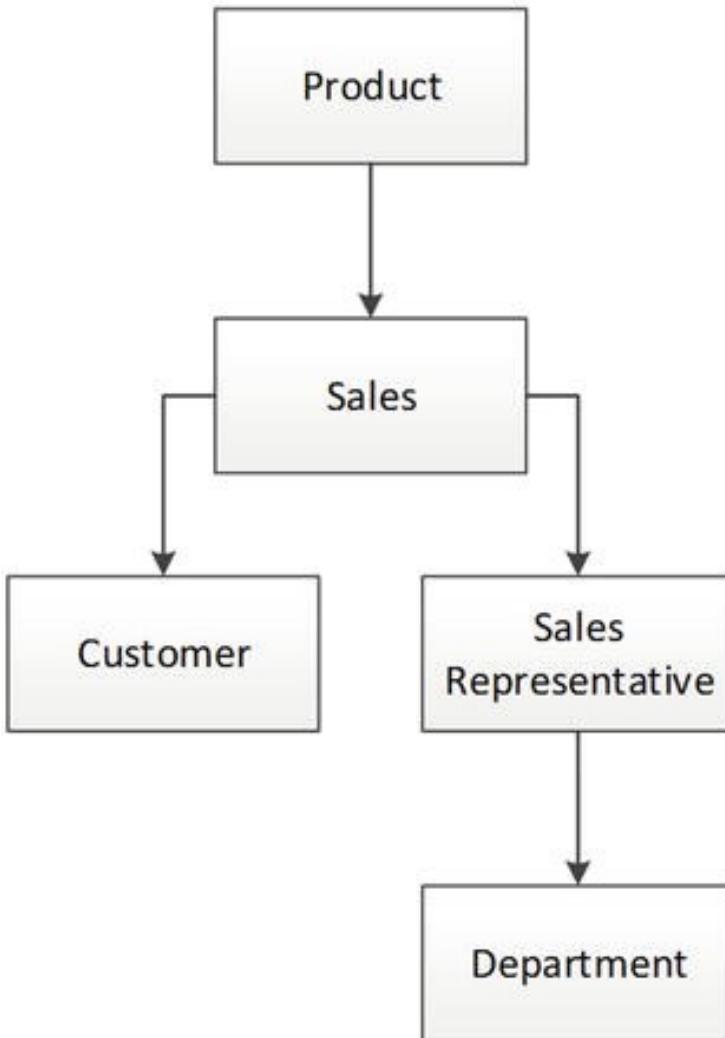


# The three database revolutions

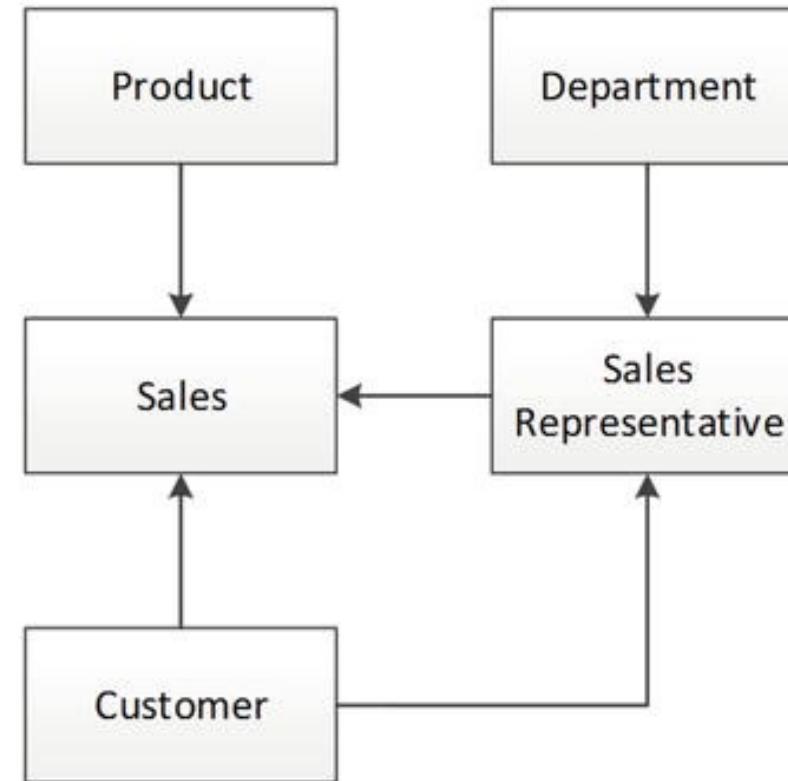


# Pre-Relational

Hierarchical Model



Network Model



# Problems with pre-relational databases

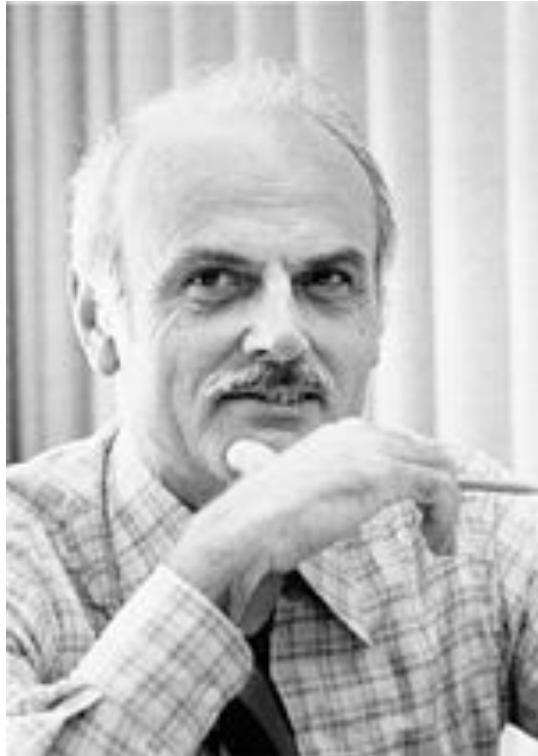
---

- **Hard to use:** required programming expertise
- **Arbitrary representations** – no logical foundation
- **Logical/Physical Dependence** – the representation of the data is closely tied to how it is stored on disk.



# Codd, 1970. The origins of the relational model

---



Edgar Frank Codd  
1923-2003

## *Information Retrieval*

---

### A Relational Model of Data for Large Shared Data Banks

E. F. CODD

*IBM Research Laboratory, San Jose, California*

Future users of large data banks must be protected from having to know how the data is organized in the machine (the internal representation). A prompting service which supplies such information is not a satisfactory solution. Activities of users at terminals and most application programs should remain unaffected when the internal representation of data is changed and even when some aspects of the external representation are changed. Changes in data representation will often be needed as a result of changes in query, update, and report traffic and natural growth in the types of stored information.



# Relational Data Model

Entities

Branch			
branchNo	street	city	postCode
B005	22 Deer Rd	London	SW1 4EH
B007	16 Argyll St	Aberdeen	AB2 3SU
B003	163 Main St	Glasgow	G11 9QX
B004	32 Manse Rd	Bristol	BS99 1NZ
B002	56 Clover Dr	London	NW10 6EU

Staff

staffNo	fName	IName	position	sex	DOB	salary	branchNo
SL21	John	White	Manager	M	1-Oct-45	30000	B005
SG37	Ann	Beech	Assistant	F	10-Nov-60	12000	B003
SG14	David	Ford	Supervisor	M	24-Mar-58	18000	B003
SA9	Mary	Howe	Assistant	F	19-Feb-70	9000	B007
SG5	Susan	Brand	Manager	F	3-Jun-40	24000	B003
SL41	Julie	Lee	Assistant	F	13-Jun-65	9000	B005

A column or **attribute** or **field** providing more details about the particular entity.

- Name
- Datatype
- Other attributes

A row or **record** describing one instance of the entity



# Key concepts of the relational model

---

- **Tuples**, an unordered set of **attribute** values. In an actual database system, a tuple corresponds to a row, and an attribute to a column value.
- A **relation**, which is a collection of distinct tuples and corresponds to a table in relational database implementations.
- **Constraints**, which enforce consistency of the database. Key constraints are used to identify tuples and relationships between tuples.
- **Operations** on relations such as joins, projections, unions, and so on. These operations always return relations. In practice, this means that a query on a table returns data in a tabular format.

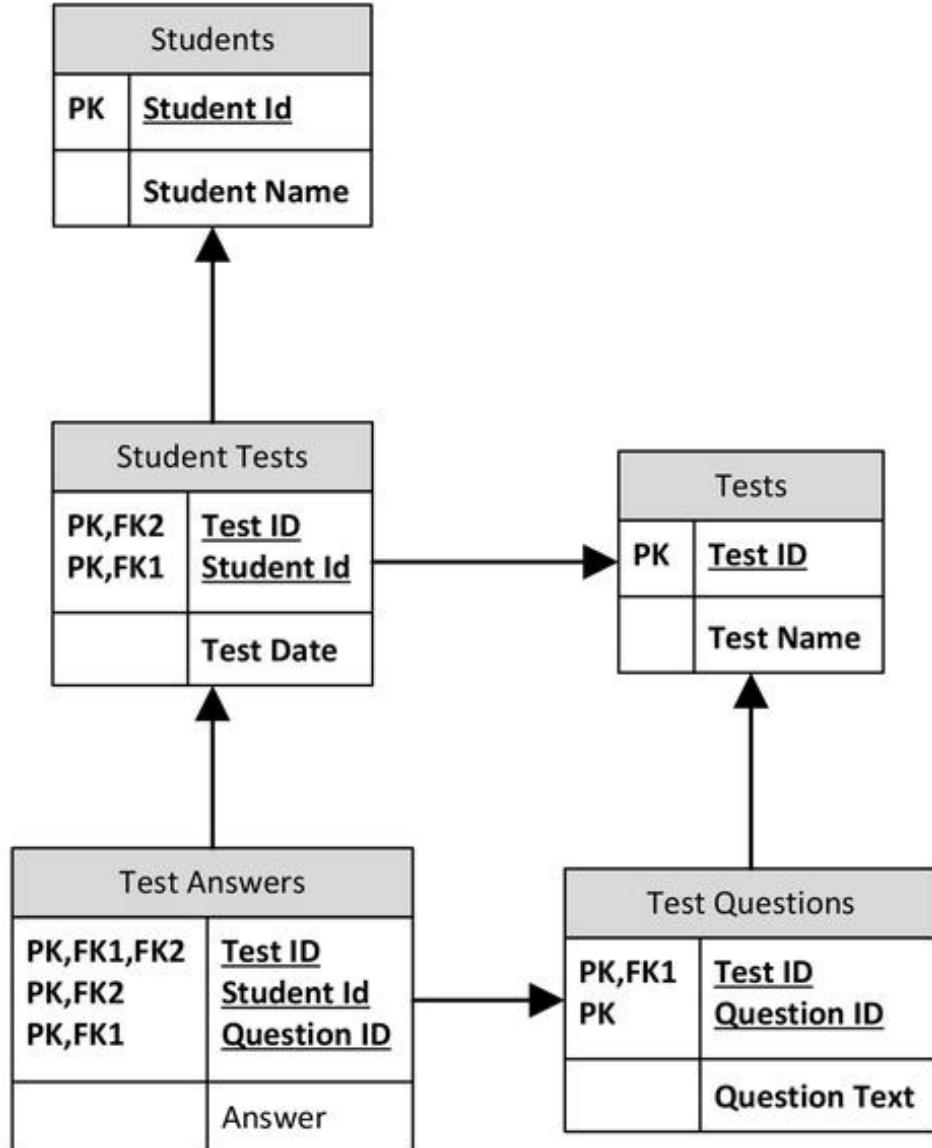


# Normalization

Un-normalized data

Test scores
Student Name
Test Name
Test Date
Answer 1
Answer 2
Answer 3
Answer 4
Answer 5
Answer 6
Answer N

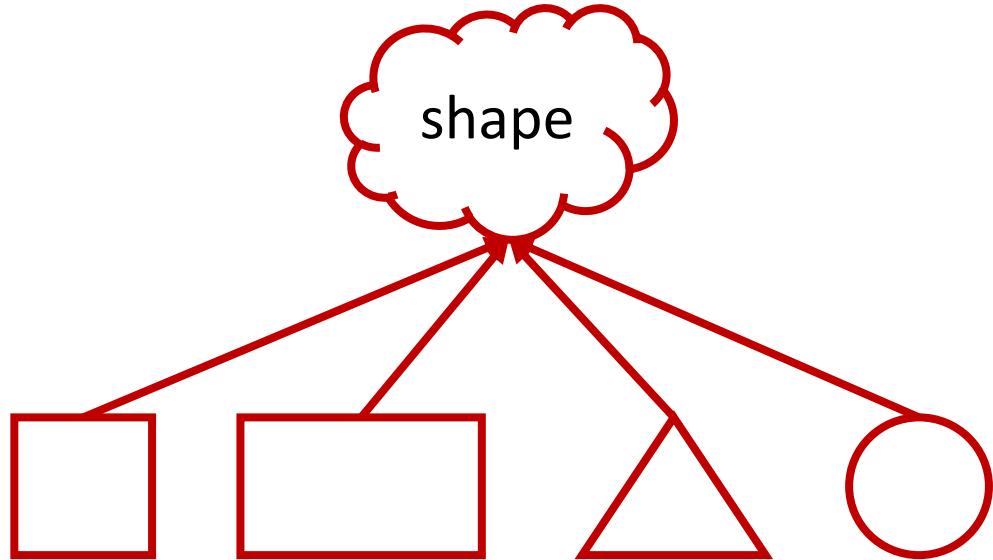
Normalized data



*“All non-key attributes depend on the key, the whole key, and nothing but the key -- so help me Codd!”*

# Object-Oriented Design

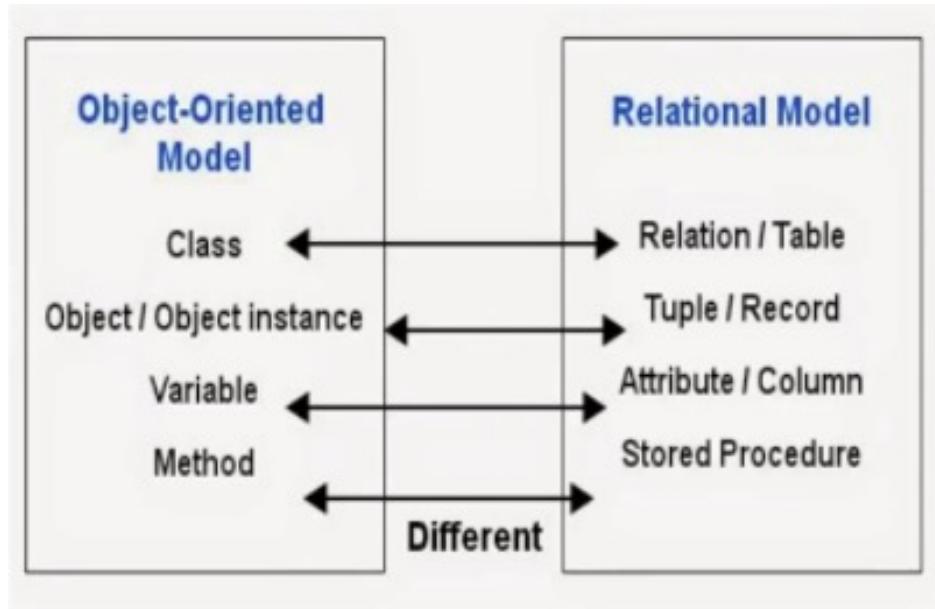
- **Encapsulation:** A *class* encapsulates both data and methods.
- **Inheritance:** Classes can inherit properties & methods of parent classes.



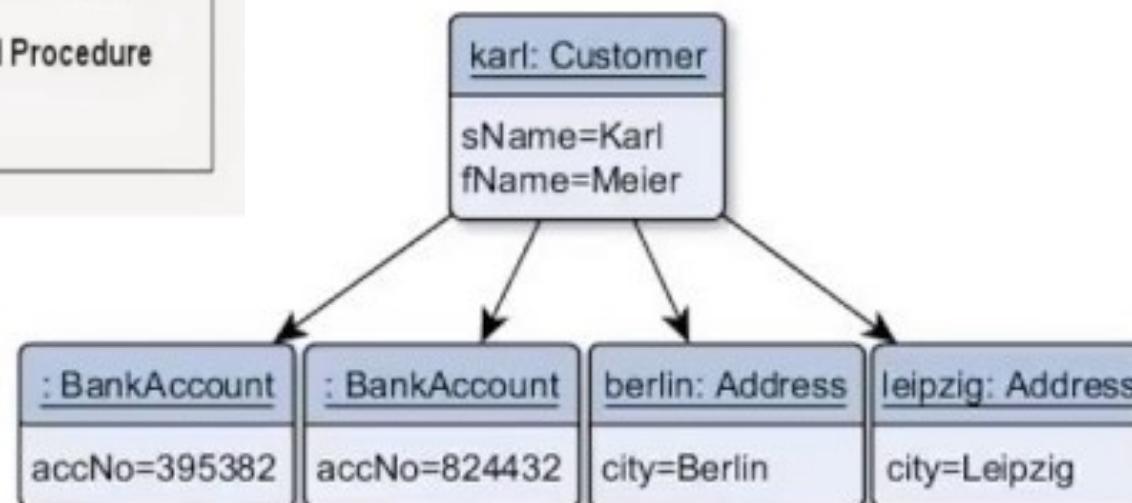
```
Class Circle extends Shape  
{  
    private double radius;  
  
    double area() { return pi * r * r; }  
}
```



# Relational vs. Object-Oriented Databases (1990's)



Object-Oriented Data Model



Relational Data Model

Customers		
ID	SName	FName
1	Karl	Meier

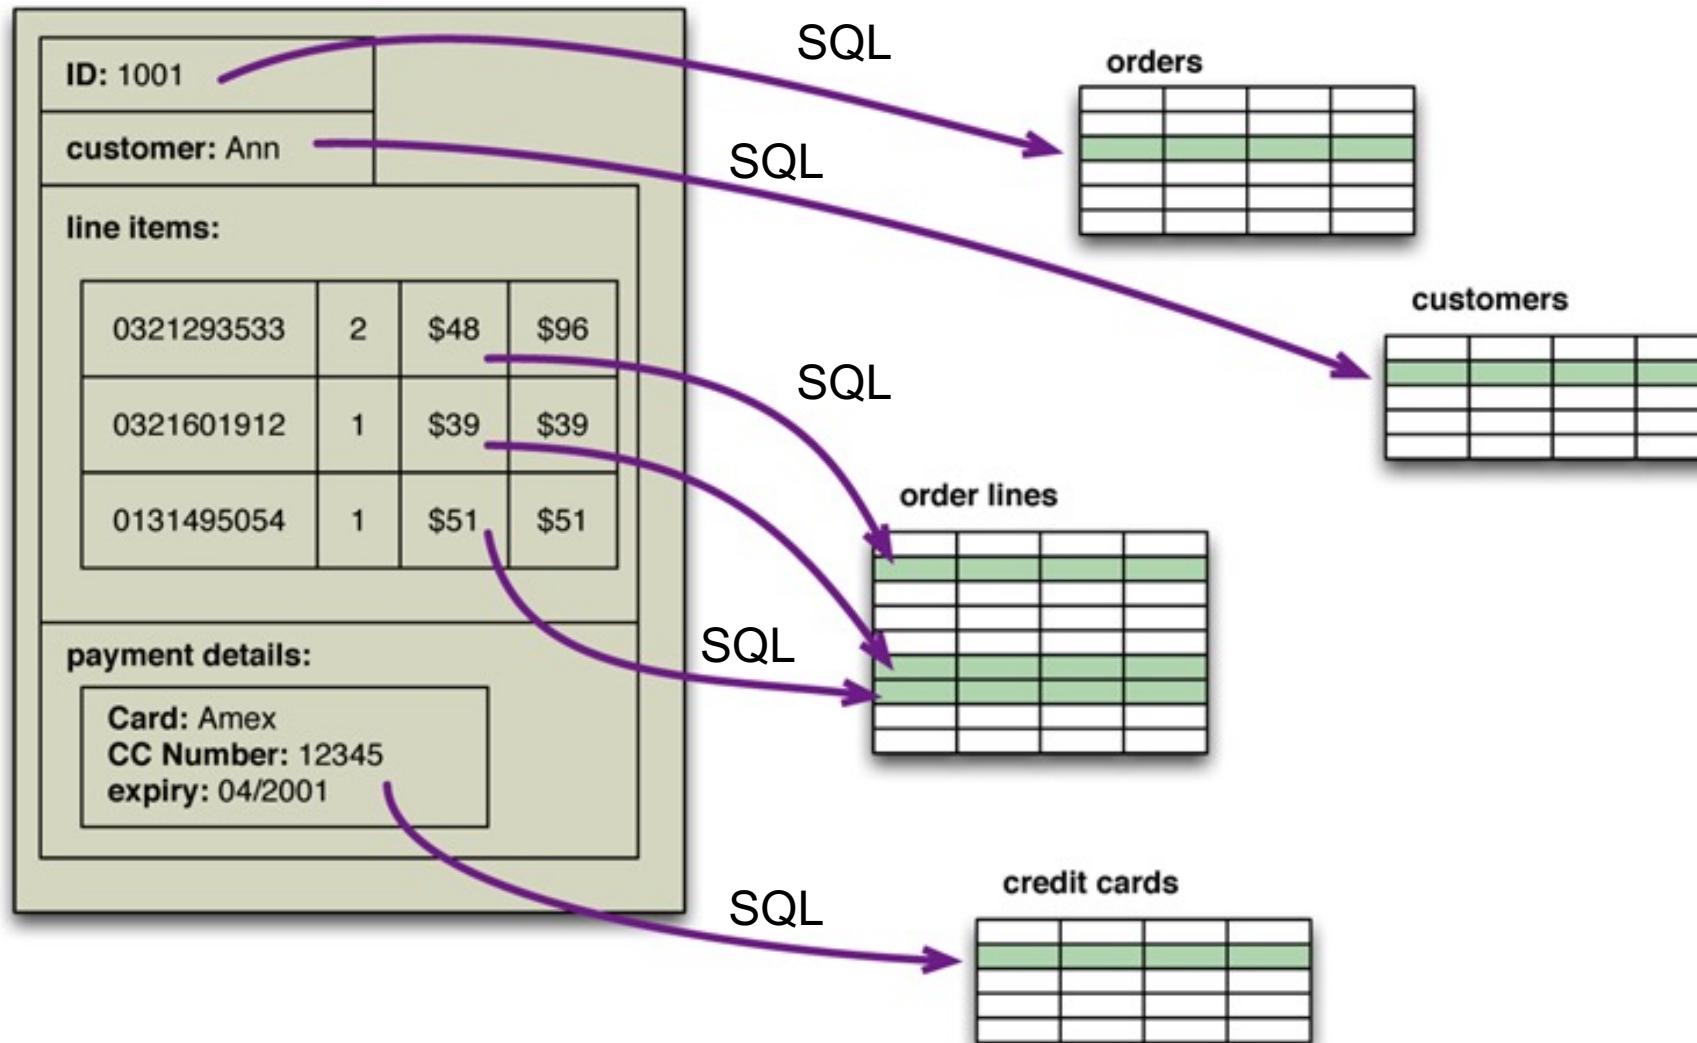
Addresses		
ID	City	CID
1	Berlin	1
2	Leipzig	1

BankAccounts		
ID	AccNo	CID
1	395382	1
2	824432	1



# Why NoSQL?: Impedance mismatch problem



# Beyond Relational Databases

---

*A relational database is like a garage that forces you to take your car apart and store the pieces in little drawers.*

-- Object-oriented data community, mid-1990s

*An object database is like a closet which requires that you hang up your suit, with tie, underwear, belt, socks and shoes all attached.*

-- David Ensor, same period.



# What happened to Object-Oriented Databases?

---

- Object-oriented databases were an attempt in the 1980's and 1990's to store objects in a DB and thus avoid the impedance mismatch problem.
- *The idea never took off*, although object-oriented programming is now universal. Why?
  - a. **SQL became an industry standard** as a foundation for supporting data integration efforts
  - b. The problem was abated to some extent by the emergence of object-relational mapping (ORM) frameworks such as **Hibernate**
  - c. **Professional divide** between application developers and database administrators
  - d. Big data wasn't a technological driver (yet!)



# Transaction models: ACID

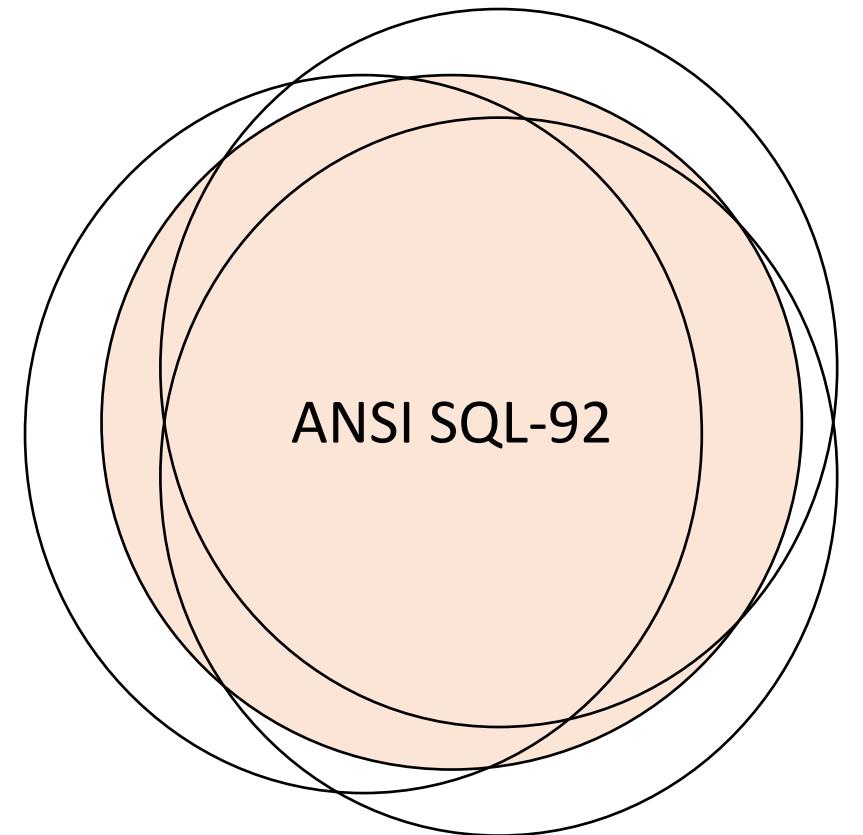
---

- **Atomic:** The transaction is indivisible—either all the statements in the transaction are applied to the database or none are.
- **Consistent:** The database remains in a consistent state before and after transaction execution.
- **Isolated:** While multiple transactions can be executed by one or more users simultaneously, one transaction should not see the effects of other in-progress transactions.
- **Durable:** Once a transaction is saved to the database (in SQL databases via the COMMIT command), its changes are expected to persist even if there is a failure of operating system or hardware.



# “But everyone knows SQL!”

Year	Name	Alias	Comments
1986	SQL-86	SQL-87	First formalized by ANSI.
1989	SQL-89	FIPS 127-1	Minor revision that added integrity constraints, adopted as FIPS 127-1.
1992	SQL-92	SQL2, FIPS 127-2	Major revision (ISO 9075), <i>Entry Level</i> SQL-92 adopted as FIPS 127-2.
1999	SQL:1999	SQL3	Added regular expression matching, <a href="#">recursive queries</a> (e.g. <a href="#">transitive closure</a> ), <a href="#">triggers</a> , support for procedural and control-of-flow statements, non-scalar types (arrays), and some object-oriented features (e.g. <a href="#">structured types</a> ). Support for embedding SQL in Java ( <a href="#">SQL/OLB</a> ) and vice versa ( <a href="#">SQL/JRT</a> ).
2003	SQL:2003		Introduced <a href="#">XML</a> -related features ( <a href="#">SQL/XML</a> ), <a href="#">window functions</a> , standardized sequences, and columns with auto-generated values (including identity-columns).
2006	SQL:2006		ISO/IEC 9075-14:2006 defines ways that SQL can be used with XML. It defines ways of importing and storing XML data in an SQL database, manipulating it within the database, and publishing both XML and conventional SQL-data in XML form. In addition, it lets applications integrate queries into their SQL code with <a href="#">XQuery</a> , the XML Query Language published by the World Wide Web Consortium ( <a href="#">W3C</a> ), to concurrently access ordinary SQL-data and XML documents. <sup>[34]</sup>
2008	SQL:2008		Legalizes ORDER BY outside cursor definitions. Adds INSTEAD OF triggers, TRUNCATE statement, <sup>[35]</sup> FETCH clause.
2011	SQL:2011		Adds temporal data (PERIOD FOR) <sup>[36]</sup> (more information at: <a href="#">Temporal database#History</a> ). Enhancements for <a href="#">window functions</a> and <a href="#">FETCH clause</a> . <sup>[37]</sup>
2016	SQL:2016		Adds row pattern matching, polymorphic table functions, JSON.



One standard, many dialects



# The beginning of NoSQL in the 2000's

---

Google (2006). [BigTable: A Distributed Storage System for Structured Data](#)  
*Proceedings of the 7<sup>th</sup> USENIX Symposium on Operating Systems Design and implementation: pp 205-218*

Abstract Excerpt: Bigtable is a distributed storage system for managing structured data that is designed to scale to a very large size: petabytes of data across thousands of commodity servers. Many projects at Google store data in Bigtable, including web indexing, Google Earth, and Google Finance.

<https://static.googleusercontent.com/media/research.google.com/en//archive/bigtable-osdi06.pdf>



# The beginning of NoSQL in the 2000's

---

Amazon (2007). Dynamo: Amazon's Highly Available Key-value Store  
*21<sup>st</sup> ACM Symposium on Operating Systems Principles: pp 205-218*

Abstract Excerpt:

This paper presents the design and implementation of Dynamo, a highly available key-value storage system that some of Amazon's core services use to provide an “always-on” experience...To achieve this level of availability, Dynamo sacrifices consistency under certain failure scenarios. It makes extensive use of object versioning and application-assisted conflict resolution in a manner that provides a novel interface for developers to use.

<http://www.allthingsdistributed.com/files/amazon-dynamo-sosp2007.pdf>



Northeastern University

# The origins of “NoSQL” – Johan Oskarsson’s SF Meetup

## Braindump

### NOSQL meetup

#### Hadoop summit

I'm going to attend the Hadoop summit in San Francisco in June, had a great time last year, learned a bunch of stuff and got to meet a lot of people I previously only knew by name.

#### NOSQL

To make the most of the flight money I'm putting together a free meetup about "open source, distributed, non relational databases" or NOSQL for short.

It's taking place on the 11th of June, the day after the Hadoop summit in San Francisco. CBS interactive have been kind enough to provide us with both a venue and free lunch!

If you wish to attend, please [register](#).

#### Preliminary schedule

- 09.45: Doors open
- 10.00: Intro session (Todd Lipcon, Cloudera)
- 10.40: Voldemort (Jay Kreps, LinkedIn)
- 11.20: Short break
- 11.30: Cassandra (Avinash Lakshman, Facebook)
- 12.10: Free lunch (sponsored by CBSi)
- 13.10: Dynomite (Cliff Moon, Powerset)
- 13.50: HBase (Ryan Rawson, StumbleUpon)
- 14.30: Short break
- 14.40: Hypertable (Doug Judd, Zvents)
- 15.20: Panel discussion
- 16.00: End of meetup, relocate to a pub called Kate O'Brien's nearby

[Follow johanoskarsson](#)

[tumblr.](#)

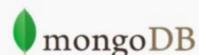
05/12/09

### Things @skr

Search

[view archive](#)

**“open source, distributed, non-relational databases”**



Dynomite



# Characteristics of NoSQL (but there are always exceptions)

---

- Non-relational
- Open-source (but many have commercial versions with add-ons and support)
- Designed for distributed storage / cluster environments
- Schema-less: Fields can be freely added on the fly.
- Non-SQL-like query languages
- Not ACID-compliant. Different options for consistency and distribution of data.
- Databases designed for 21<sup>st</sup> century web estates with “big data” needs.

NoSQL might mean “No SQL” or “Not Only SQL”

NoSQL supports a trend towards **Polyglot persistence**: using different data stores in different circumstances as part of a complete data-engineering architecture.



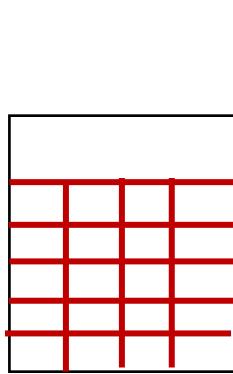
# Schema-less

- NoSQL databases operate without a schema: fields added freely w/o redefinition
- Ideal for non-uniform data and custom fields

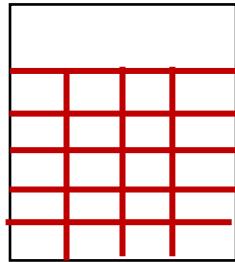
	<b>Relational</b>	<b>NoSQL</b>
Data model(s)	Tables	Key-Value, Documents, Graphs, Columns
Computing model(s)	Centralized, rate limiting	Distributed, scalable
<b>Schema</b>	<b>Pre-declared</b>	<b>Flexible</b>
Datatypes on attributes	Fixed	Not fixed
Hierarchical data types	No	Yes
Cost	Licensing costs become prohibitive as you scale vertically.	Open-source software. Hardware costs scale linearly



# Types of models

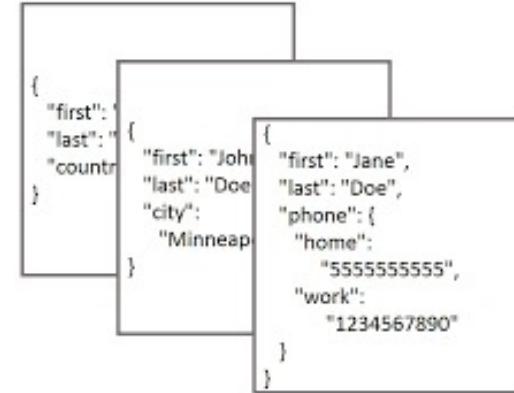


Relational

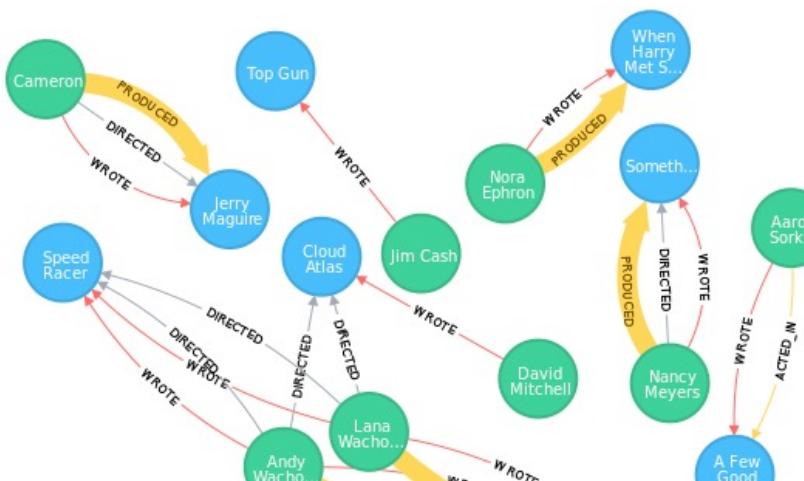


Key	Value
K1	AAA,BBB,CCC
K2	AAA,BBB
K3	AAA,DDD
K4	AAA,2,01/01/2015
K5	3,ZZZ,5623

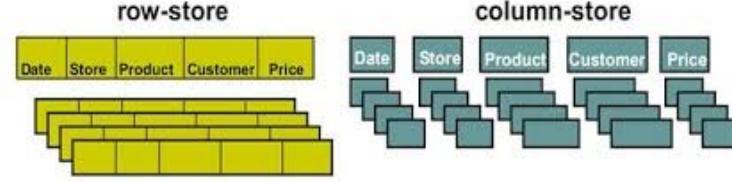
Key-Value



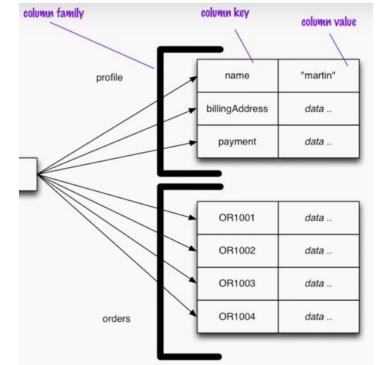
Document



Graphs



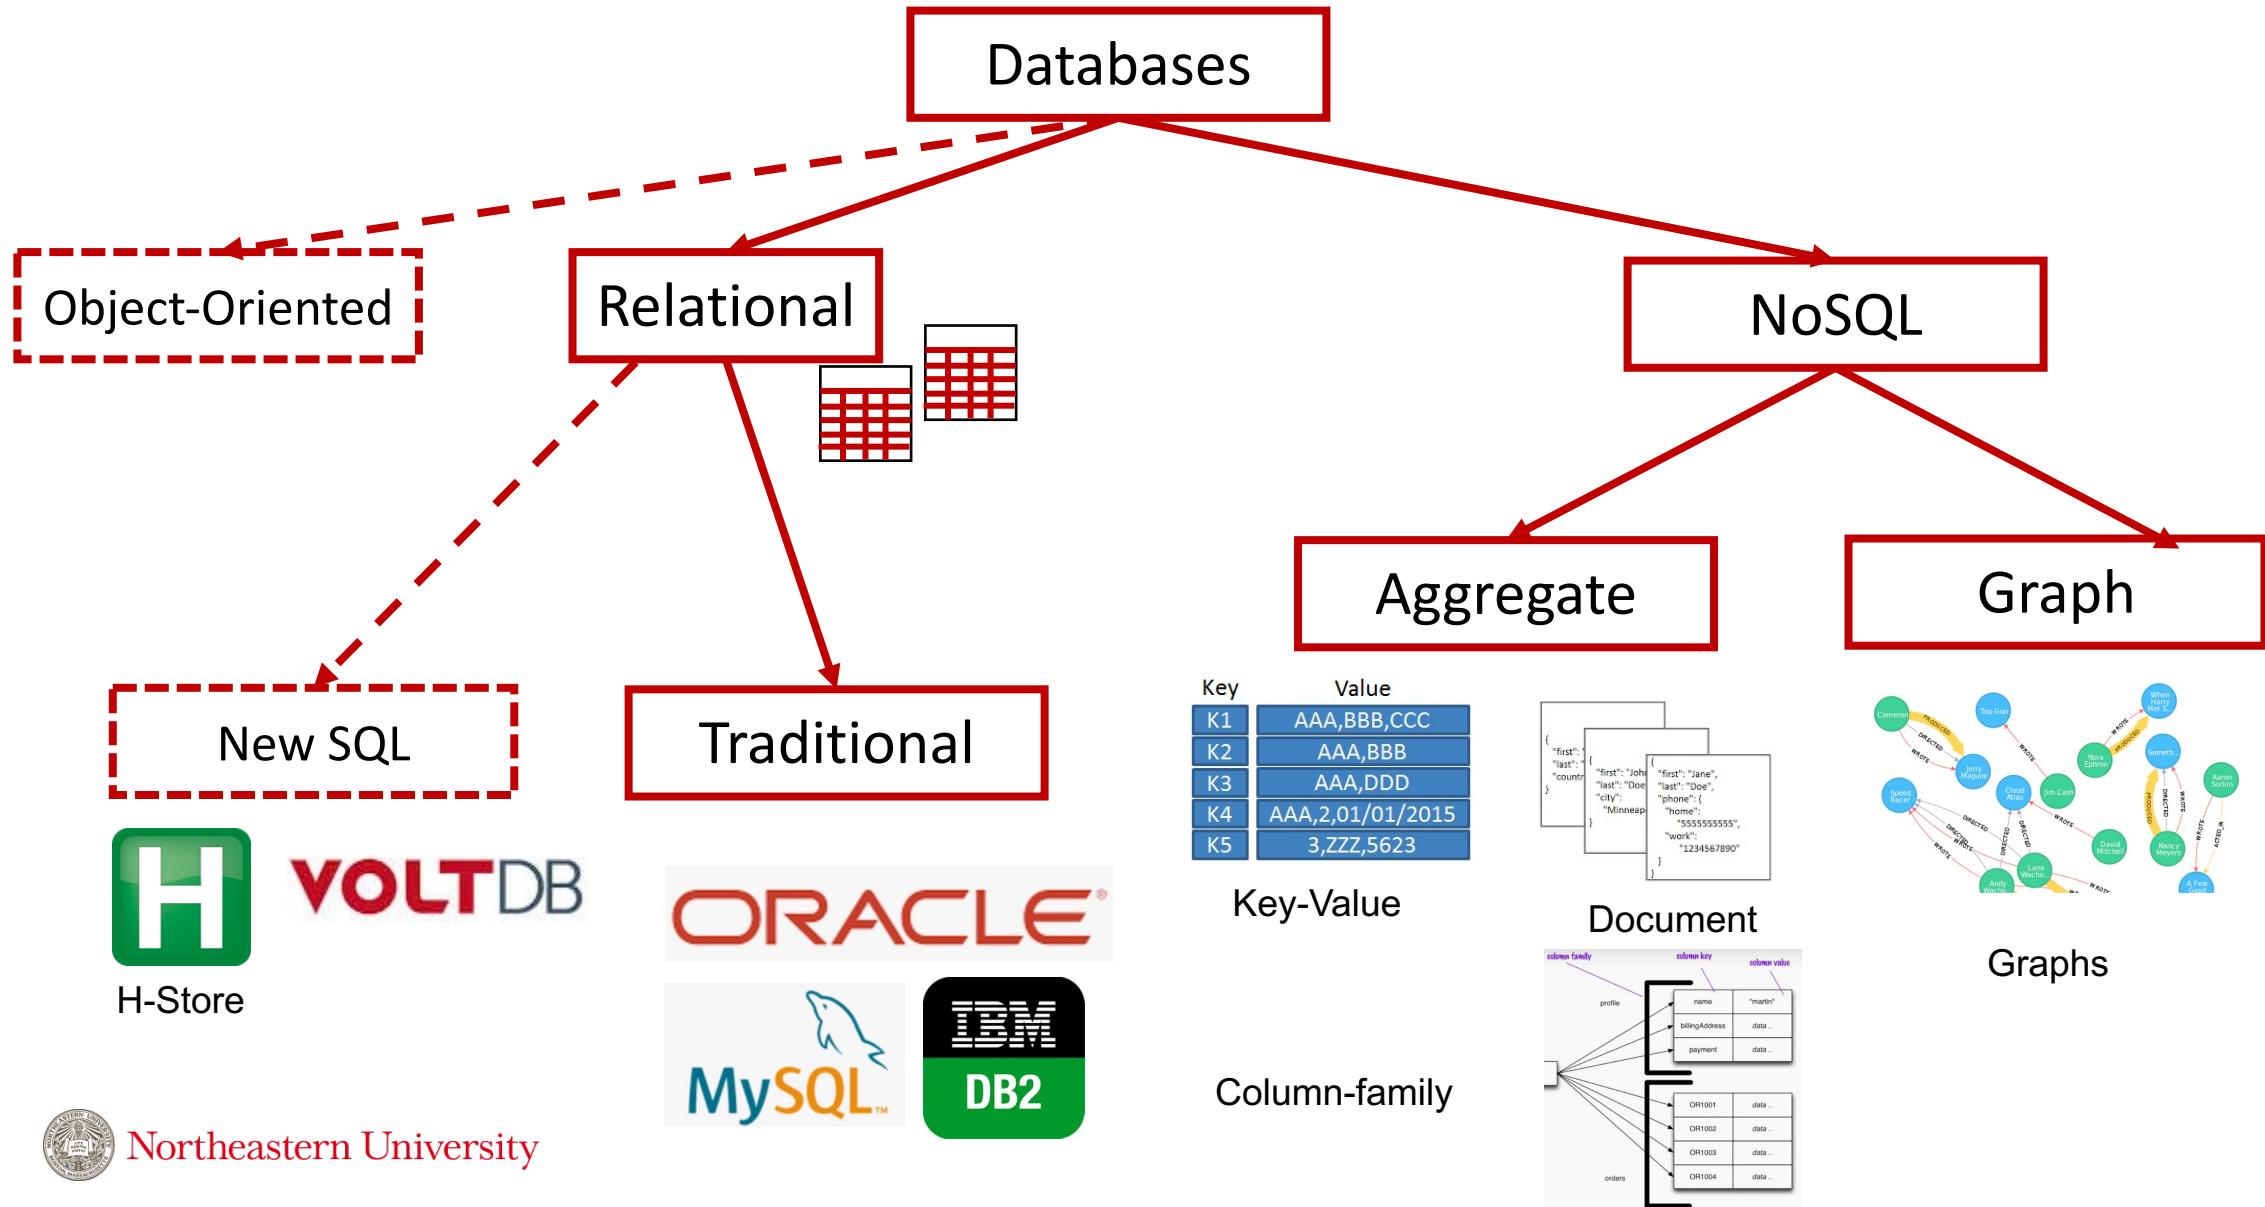
Column-oriented storage  
Column-family storage



# The non-relational explosion

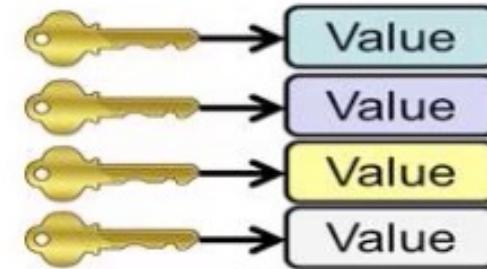


# The hierarchy of NoSQL



# Key Value

## Key-value



- Your access to data is only via a key, but the value could be anything: a simple value, a list, a document, even a database
- Often used when you need very fast but simple operations
- Common use case: managing online shopping carts.



# Document Databases

---

- A **document database** is a non-relational database that stores data as structured documents, usually XML or JSON formats.
- Simple, flexible, scalable

```
// in customers
{
  "id":1,
  "name":"Martin",
  "billingAddress":[{"city":"Chicago"}]
}

// in orders
{
  "id":99,
  "customerId":1,
  "orderItems":[
    {
      "productId":27,
      "price": 32.45,
      "productName": "NoSQL Distilled"
    }
  ],
  "shippingAddress":[{"city":"Chicago"}]
  "orderPayment": [
    {
      "ccinfo":"1000-1000-1000-1000",
      "txnid":"abelif879rft",
      "billingAddress": {"city": "Chicago"}
    }
  ],
}
```



# XML (eXtensible Markup Language): A document format

- Used as an information storage standard
- XML + CSS (Cascading Style Sheets) enables website design that distinguishes data and format.
- The foundation for web-service specifications (SOAP)
- A rich ecosystem of tools and standards that includes:

**Xpath** – Element retrieval

**Xquery** – XML document querying

**XML schema** – meta-specification for validation

- **XSLT** – XML transformations, e.g., HTML

**DOM** – A document navigation API

```
<customer_record>
<customer_id>187693</customer_id>
<name>"Kiera Brown"</name>
<address>
    <street>"1232 Sandy Blvd."</street>
    <city>"Vancouver"</city>
    <state>"Washington"</state>
    <zip>"99121"</zip>
</address>
<first_order>"01/15/2013"</first_order>
<last_order>"06/27/2014"</last_order>
</customer_record>
```



# Relational to NoSQL: Concept Mapping

RDBMS		MongoDB
Database		Database
Table, View		Collection
Row		Document (BSON)
Column		Field
Index		Index
Join		Embedded Document
Foreign Key		Reference
Partition		Shard

- Collection is not strict about what it stores
- Schema-less
- Hierarchy is evident in the design
- Embedded document are used to represent relationships explicitly.

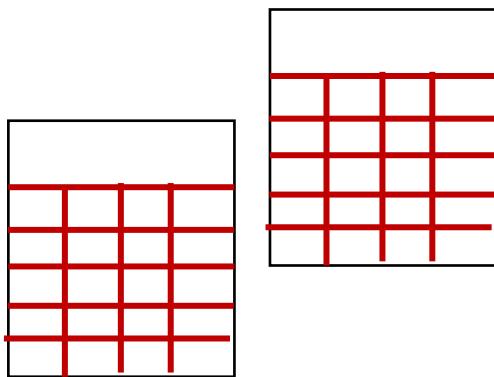


# Document databases (MongoDB): a happy medium

Relational Databases

Pre-defined / rigid schema

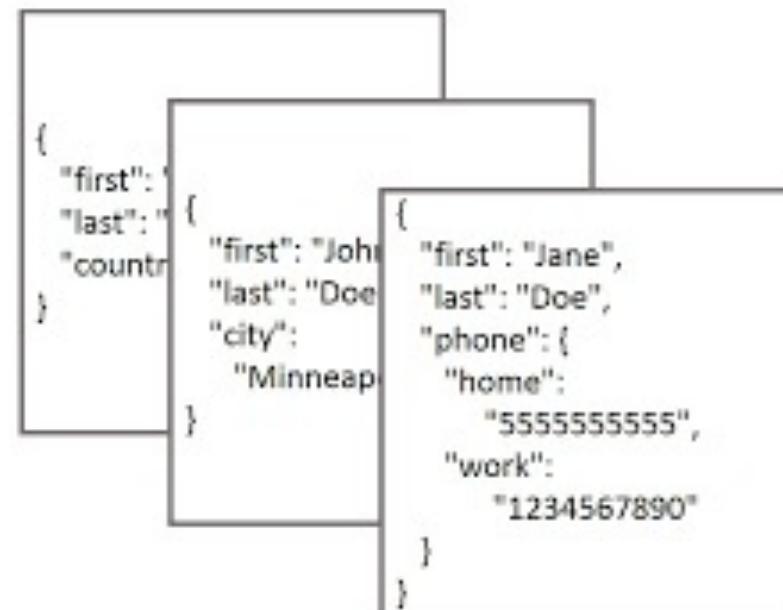
SQL: Structured Query Language



Document Databases

Data is self-describing and flexible

Query by content



Key-Value Stores

Data is simple hash tables

Query by key only (with exceptions)

Key	Value
K1	AAA,BBB,CCC
K2	AAA,BBB
K3	AAA,DDD
K4	AAA,2,01/01/2015
K5	3,ZZZ,5623



## Goal : Blend performance and functionality

Scalability & Performance

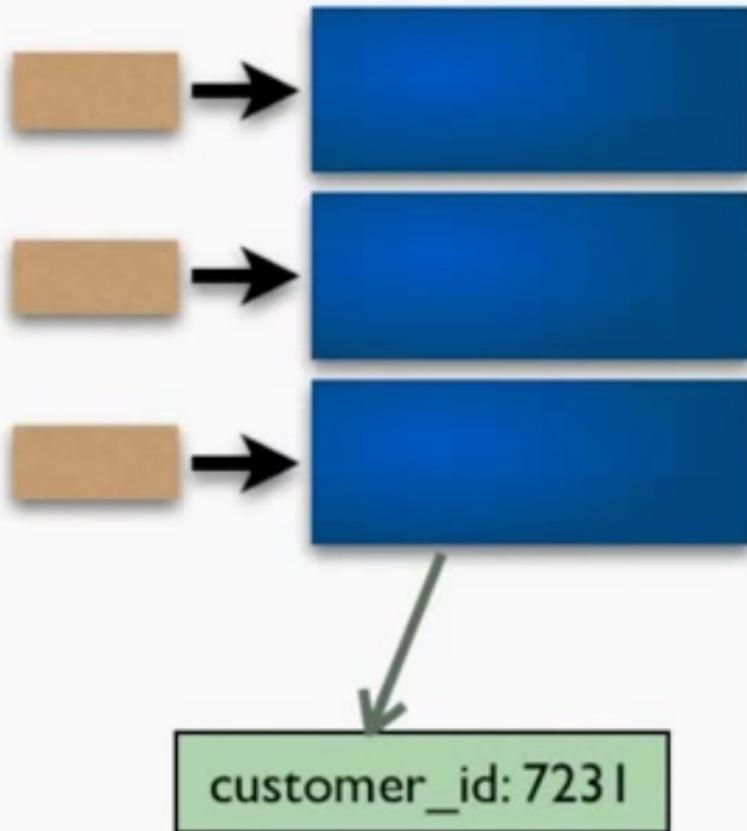
- memcached
- key/value stores
- MongoDB
- RDBMS

Depth of Functionality



The boundary between key-value and document is blurry

## Key-Value

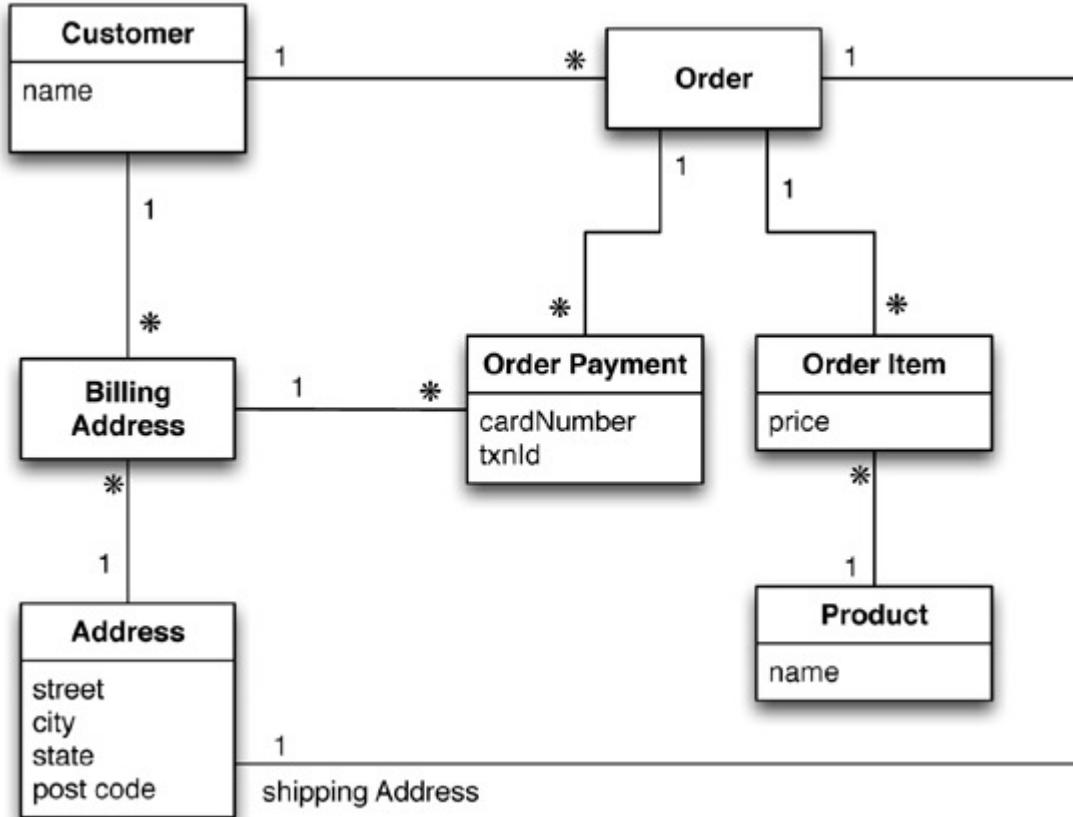


## Document

```
{"id": 1001,  
 {"id": 1002,  
 "customer_id": 7231,  
 "line-items": [  
 {"product_id": 4555, "quantity": 8},  
 {"product_id": 7655, "quantity": 4},  
 {"product_id": 8755, "quantity": 3}]  
 }  
 }
```

metadata

# Aggregate vs. Relational



Customer	
Id	Name
1	Martin

Orders		
Id	CustomerId	ShippingAddressId
99	1	77

Product	
Id	Name
27	NoSQL Distilled

BillingAddress		
Id	CustomerId	AddressId
55	1	77

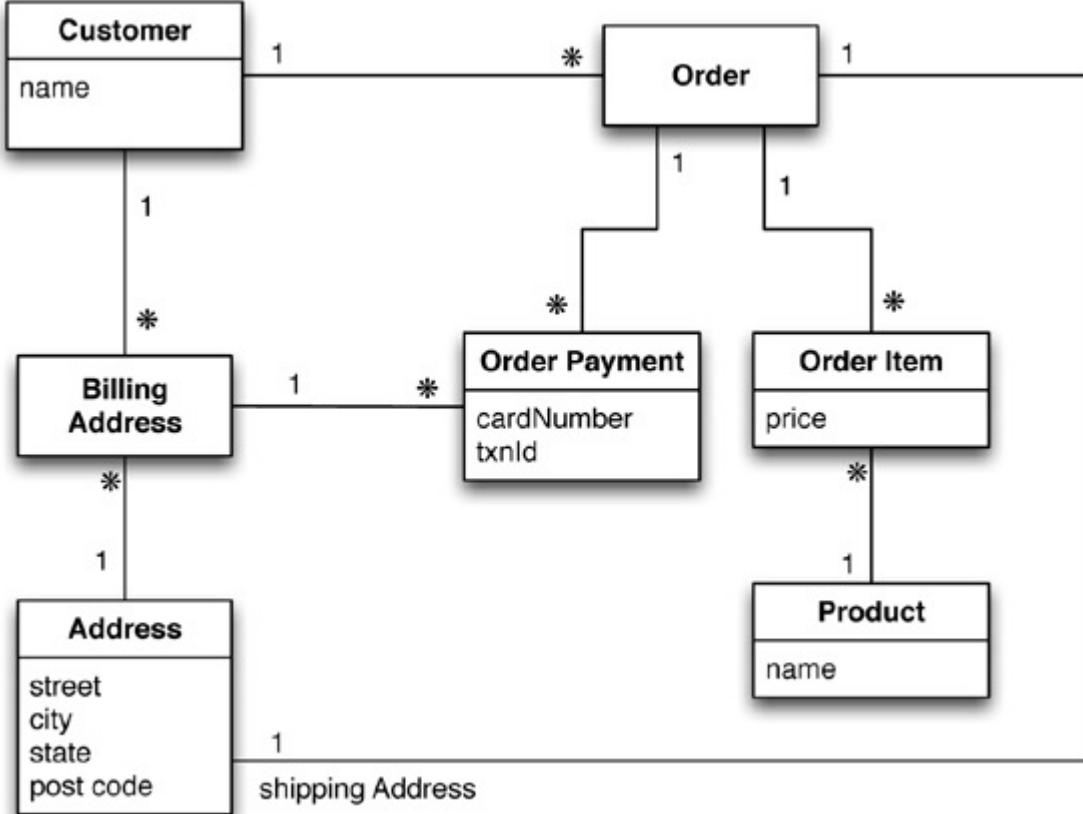
OrderItem			
Id	OrderId	ProductId	Price
100	99	27	32.45

Address	
Id	City
77	Chicago

OrderPayment				
Id	OrderId	CardNumber	BillingAddressId	txnId
33	99	1000-1000	55	abelif879rft



# Aggregate vs. Relational



```
// in customers
{
  "id":1,
  "name":"Martin",
  "billingAddress":[{"city":"Chicago"}]
}

// in orders
{
  "id":99,
  "customerId":1,
  "orderItems":[
    {
      "productId":27,
      "price": 32.45,
      "productName": "NoSQL Distilled"
    }
  ],
  "shippingAddress":[{"city":"Chicago"}]
  "orderPayment":[
    {
      "ccinfo":"1000-1000-1000-1000",
      "txnId":"abelif879rft",
      "billingAddress": {"city": "Chicago"}
    }
  ],
}
```



# Why have document databases flourished since 2008?

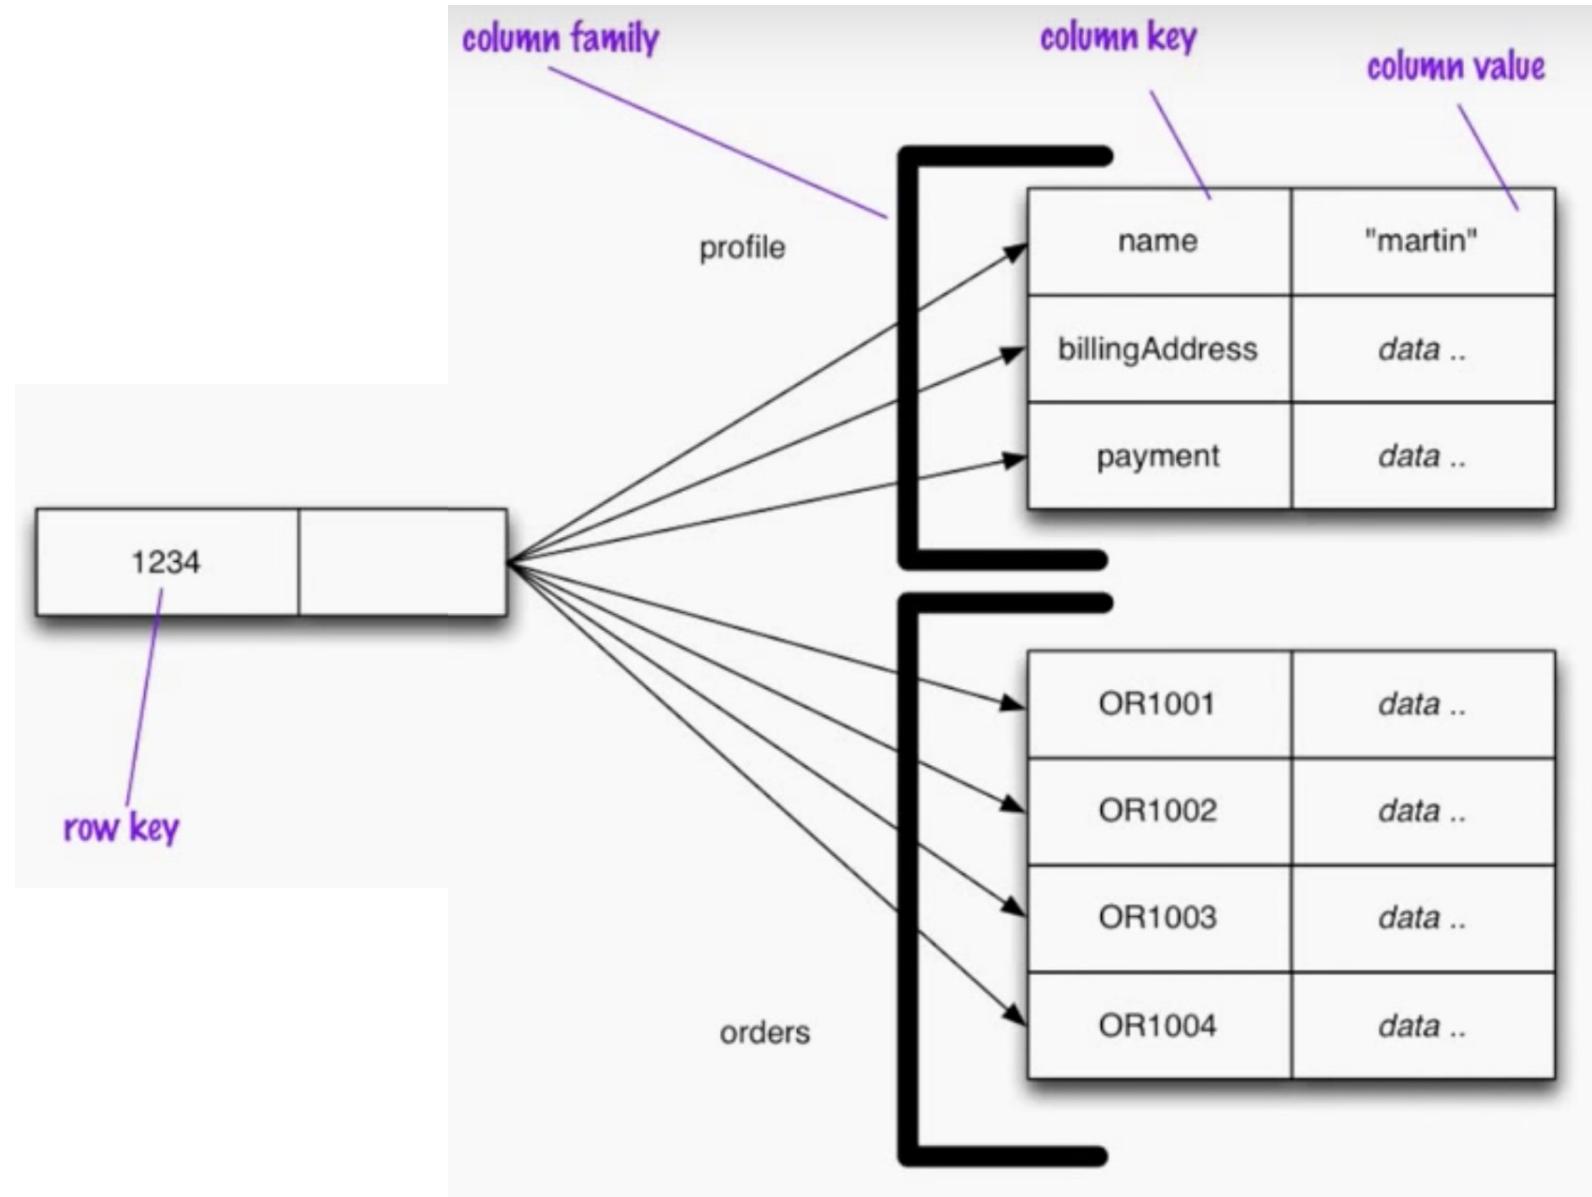
---

- They address the *impedance mismatch* problem associated with object-oriented developers needing to persist their data to relational databases. (Object-Relational Mapping)
- Self-describing formats support *ad hoc* query by value absent in pure key-value stores.
- Aligned with web-service based programming models using JSON / XML as a transport layer.
- Suitable for *horizontal scaling* in the world of big data



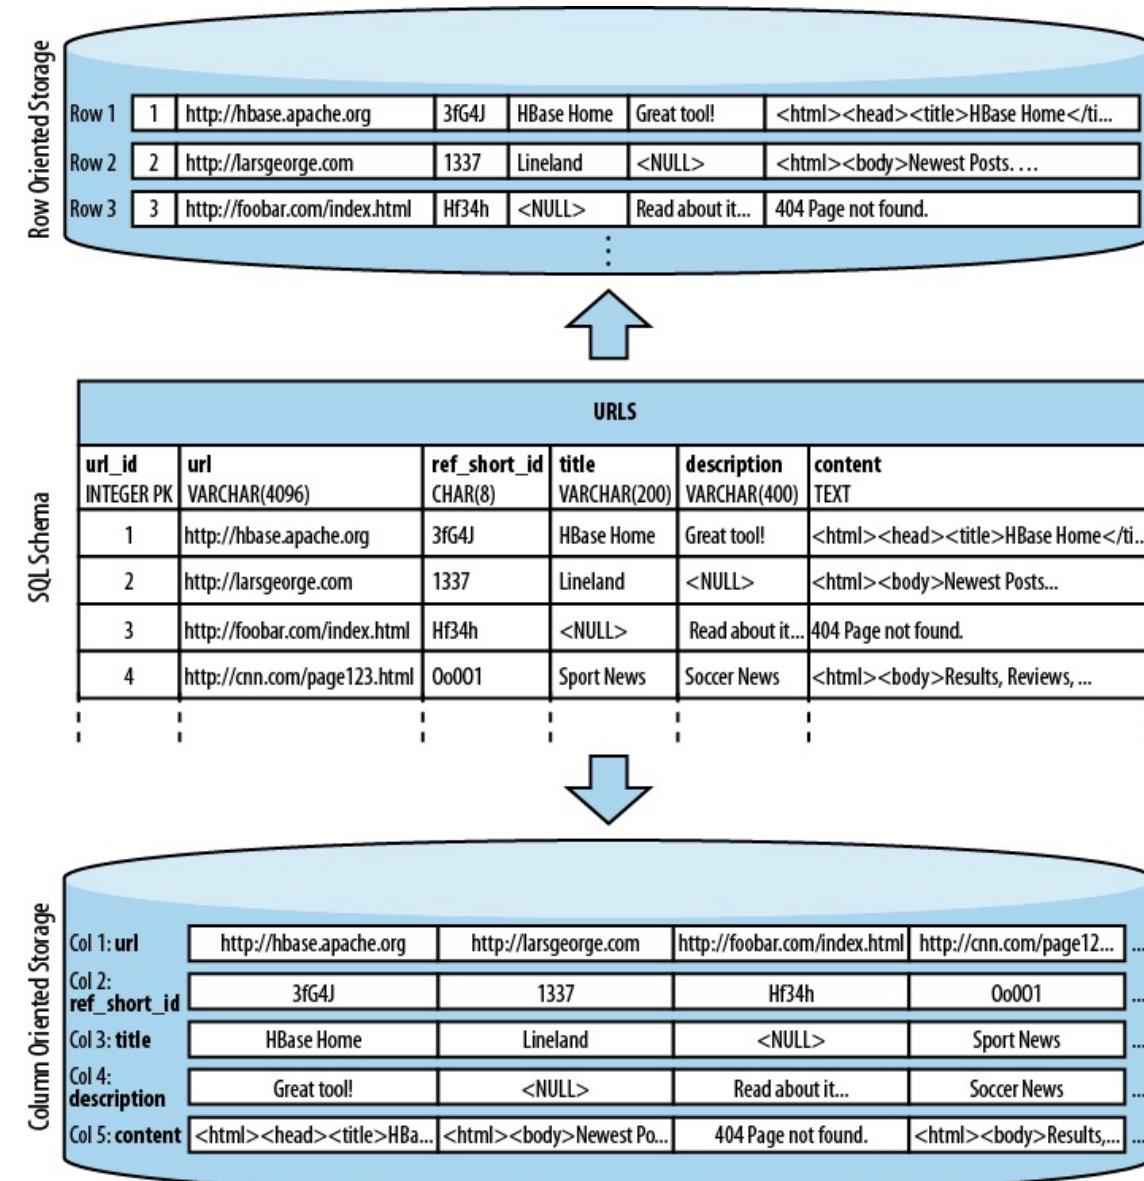
# Column Family Databases

Also “aggregate”

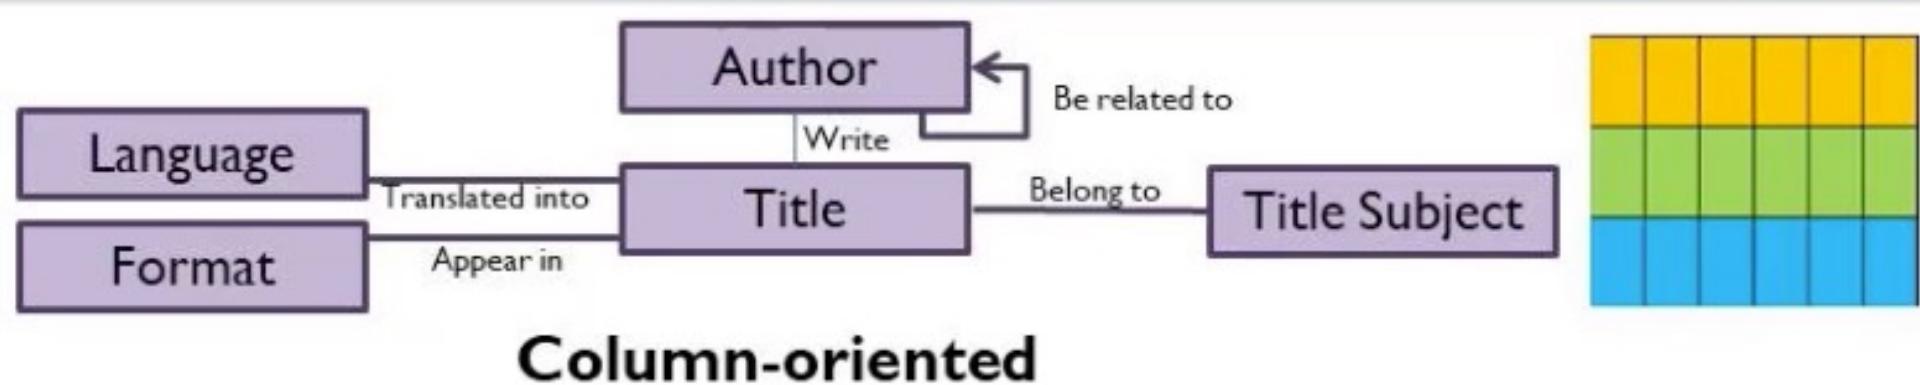


# Column-Oriented Databases

- Column databases store data as columns instead rows
- More efficient for certain types of queries / analytics.
- Reduced I/O
- Better compression
- Poor performance on row-level operations (Why?)



# Each column attribute is its own table.



9780977140060:001, 9781935504535:002, 9781935504511:003,

Data Modeling Made Simple:001, Extreme Scoping:002, UML Database Modeling Workbook:003,

Hoberman:001, Moss: 002, Blaha:003,

Steve: 001, Larissa, 002, Michael:003,

English:001,002,003,

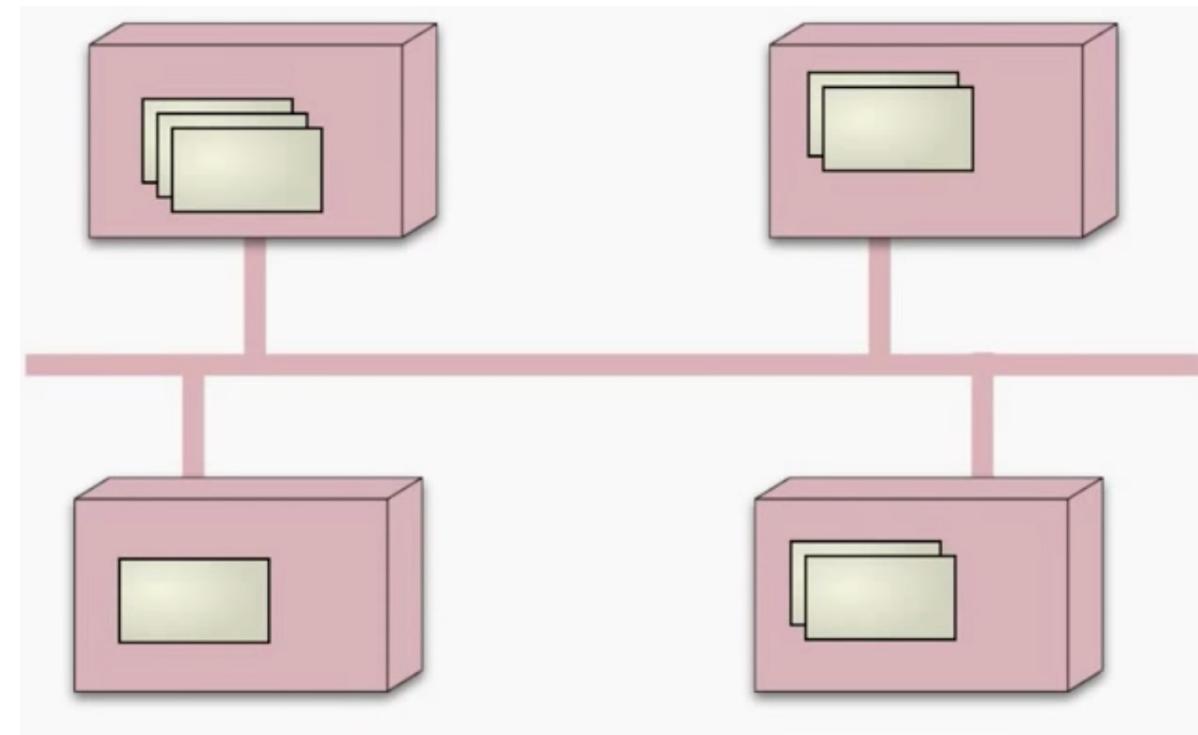
Print Version:001,002,003,

Database Design:001,003, Data Warehousing:001,002, Data Architecture,001,002, ■



# Scalability of Aggregate-Oriented DBs

- A document lives on one node where all the data can be retrieved in one operation.
- Avoid cross-node joins
- So aggregate-oriented DBs work well when you want to store your data across a cluster and your operations center around particular aggregates!



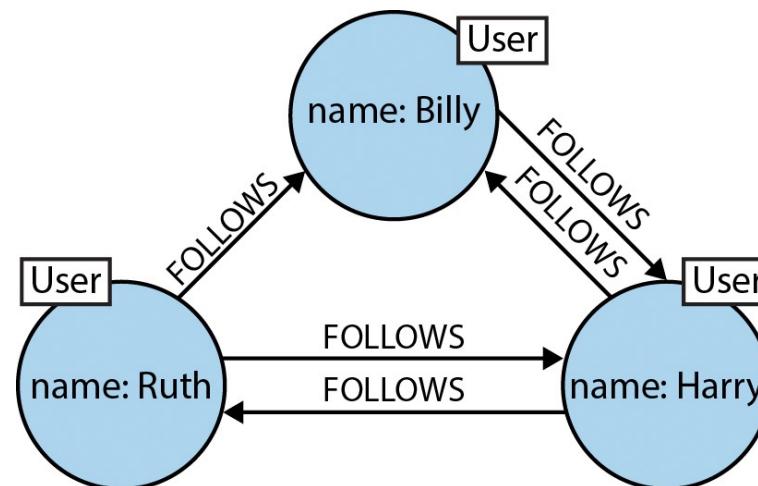
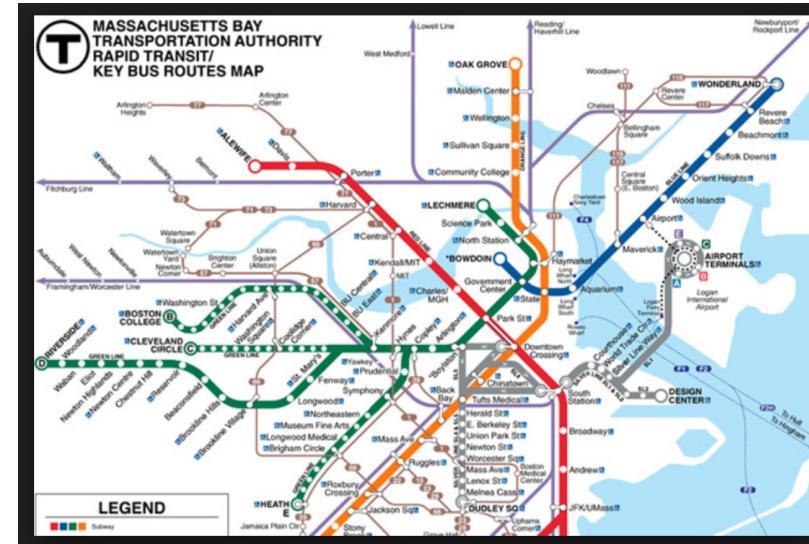
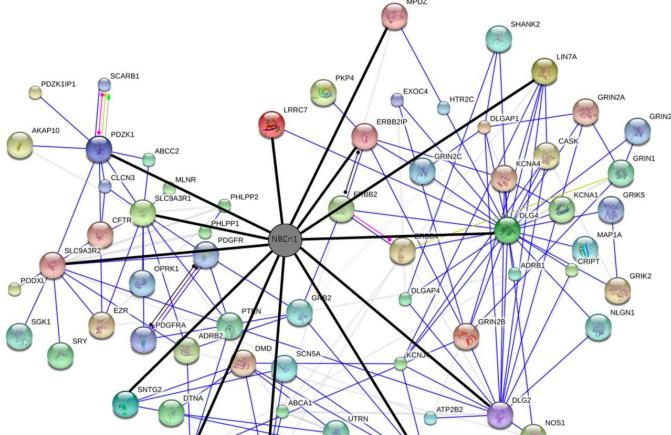
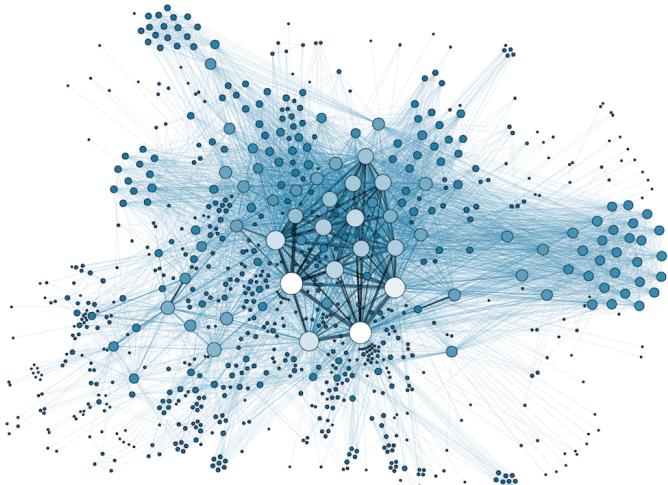
# Disadvantage of aggregates

- If you want to slice and dice your data differently, say by product rather than by order, you have to re-aggregate your data so that product is the root of the aggregate.
- This can be a complex and time-consuming operation!
- With relational, it's simple! Just alter your query.

Product	revenue	prior revenue
321293533	3083	7043
321601912	5032	4782
131495054	2198	3187
...	...	...
...	...	...
...	...	...
...	...	...
...	...	...
...	...	...
...	...	...

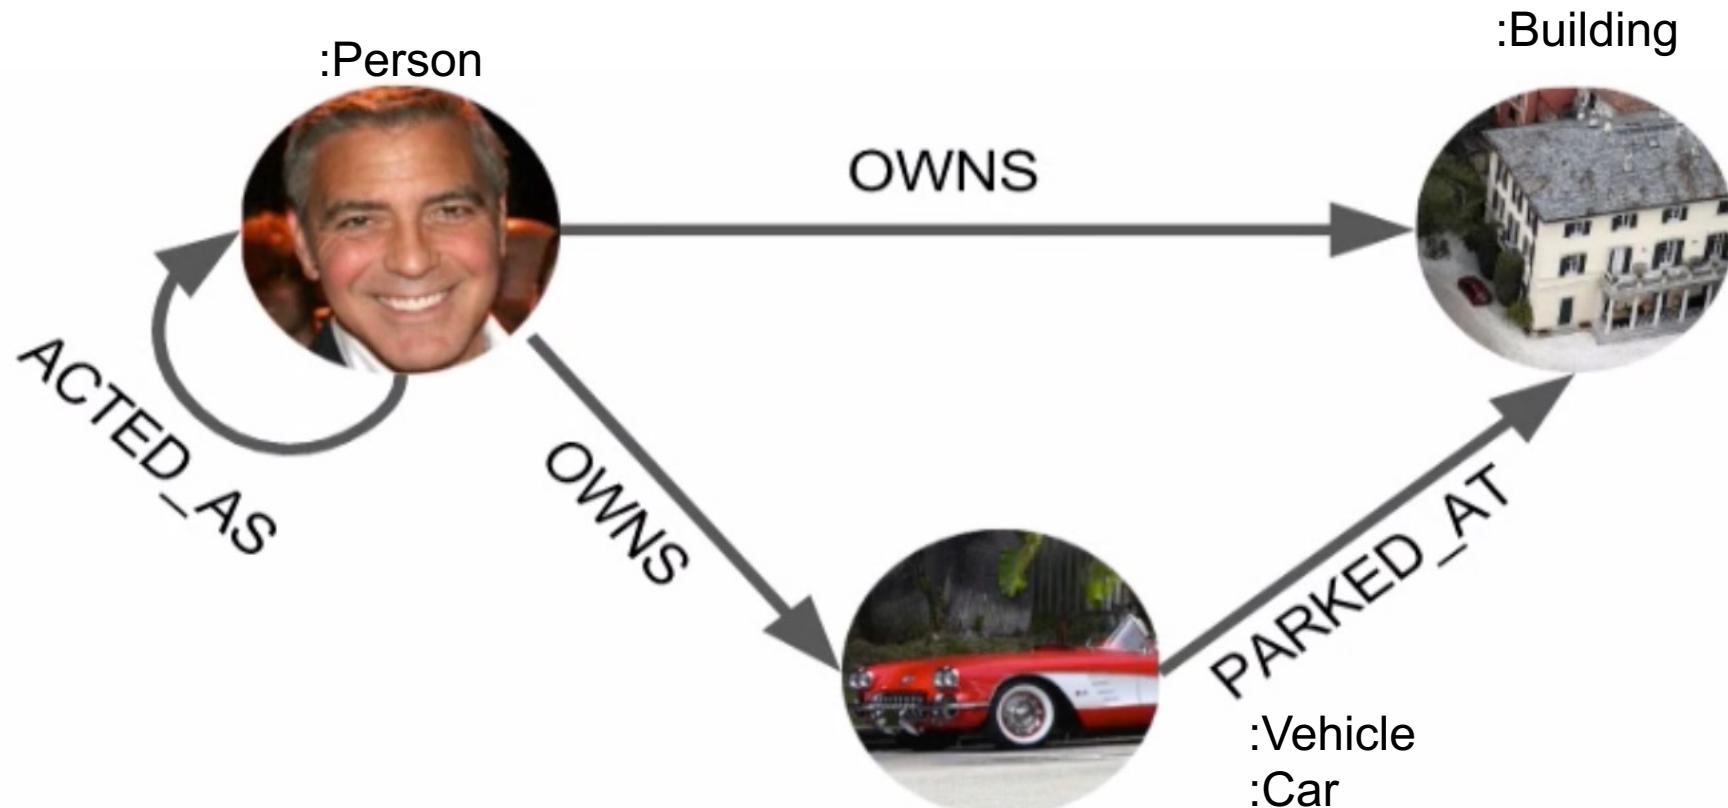


# Graph Databases



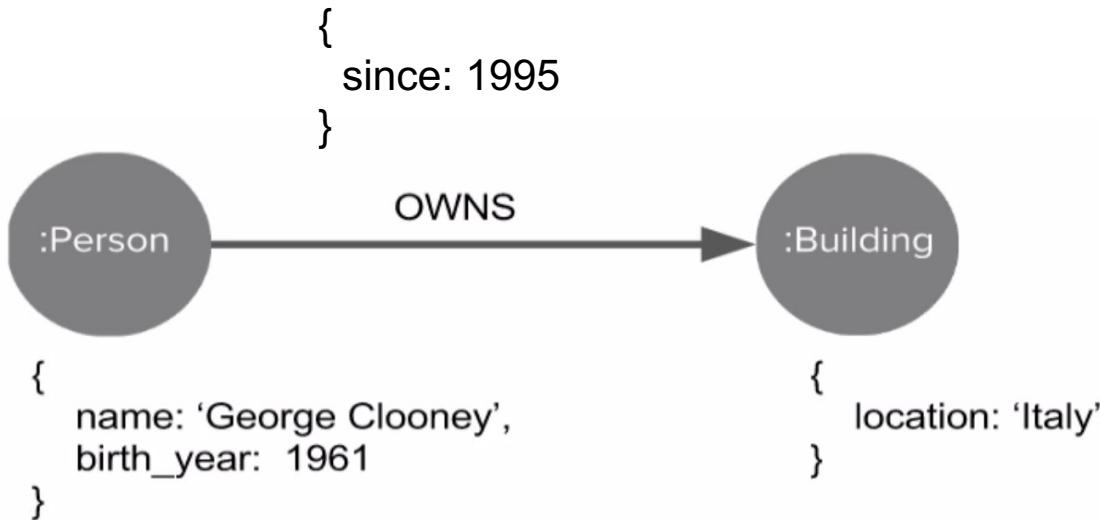
# What is a Graph?

- A graph is a collection of *nodes* and *relationships*.
- Nodes have one or more categorical labels and represent things
- Relationships have a single type and are a connected arc between two nodes



# Graph properties

- Properties are key value pairs that are assigned to a node or a relationship
- Nodes have one or more categorical labels and represent things
- Relationships have a single type and are a connected arc between two nodes

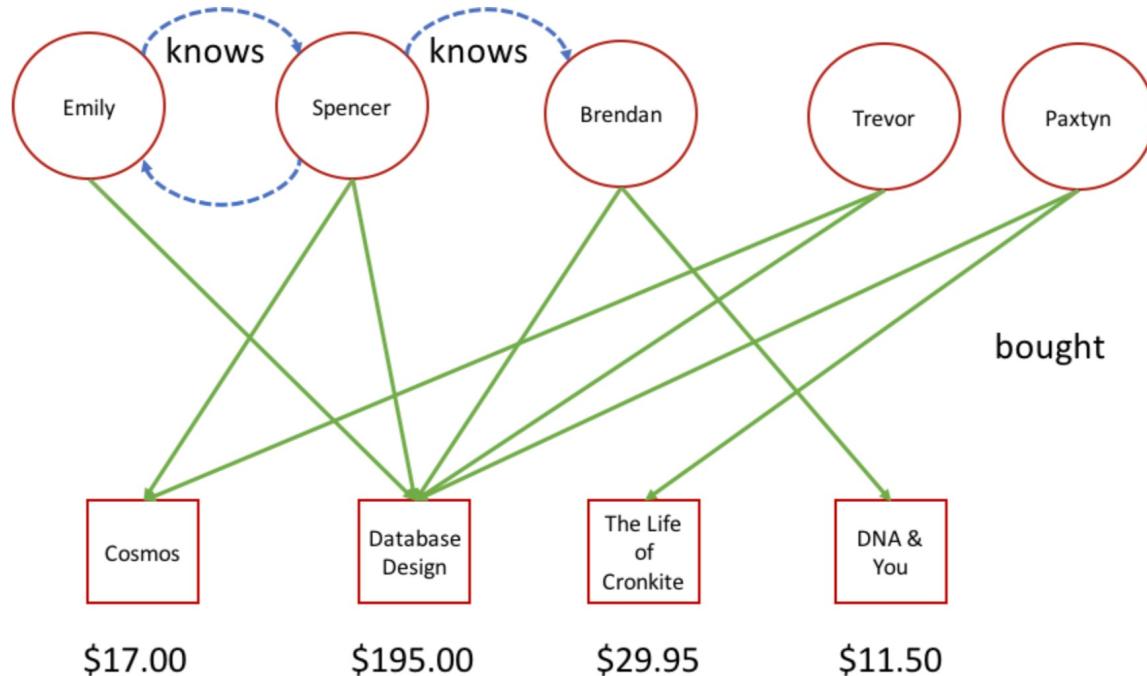


Value Types

Type	Example
Boolean	true, false
Text	"AKA string"
Numbers	123, 56.70
Lists	[ 'must', 'be', 'same', 'type' ]



# Modeling graphs in a relational database is ugly!



SQL:

```
-- What books did spencer buy? Give title and price.  
select npout.string_value title, npout2.num_value price  
from edge e  
join node nin on (e.in_node = nin.node_id)  
join node_props npin on (nin.node_id = npin.node_id)  
join node nout on (e.out_node = nout.node_id)  
join node_props npout on (nout.node_id = npout.node_id)  
join node_props npout2 on (nout.node_id = npout2.node_id)  
where e.type = 'bought'  
and npin.propkey = 'name'  
and npin.string_value = 'Spencer'  
and npout.propkey = 'title'  
and npout2.propkey = 'price'  
order by price desc;
```

Cypher (Neo4J):

```
match (e:Person {name:'Spencer'})-[:Buys]->(x:Book) return x.title, x.price
```



# Reasoning about connected knowledge

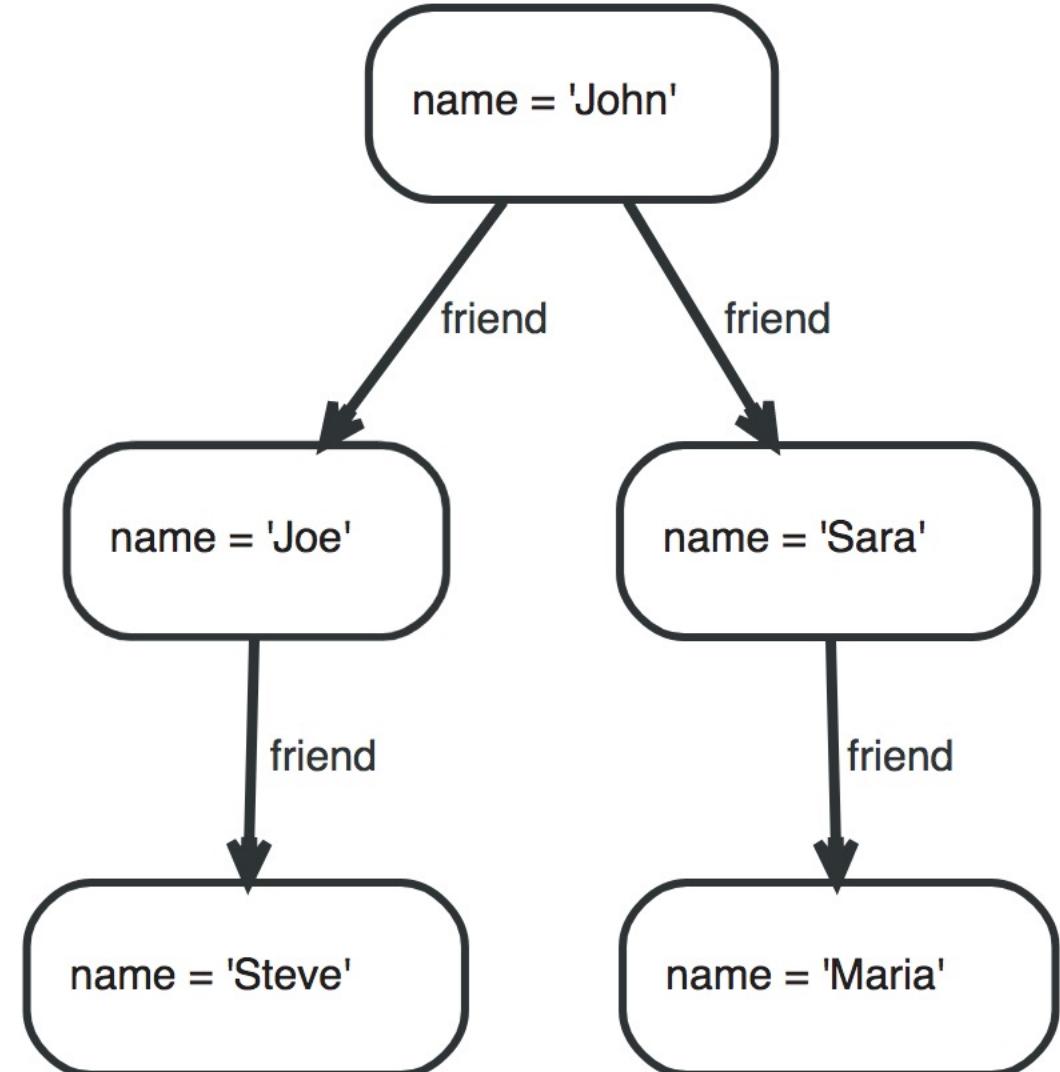
For example, here is a query which finds a user called 'John' and 'John's' friends (though not his direct friends) before returning both 'John' and any friends-of-friends that are found.

```
MATCH (john {name: 'John'})-[:friend]->()-[:friend]->(fof)  
RETURN john.name, fof.name
```

Resulting in:

john.name	fof.name
"John"	"Maria"
"John"	"Steve"

2 rows

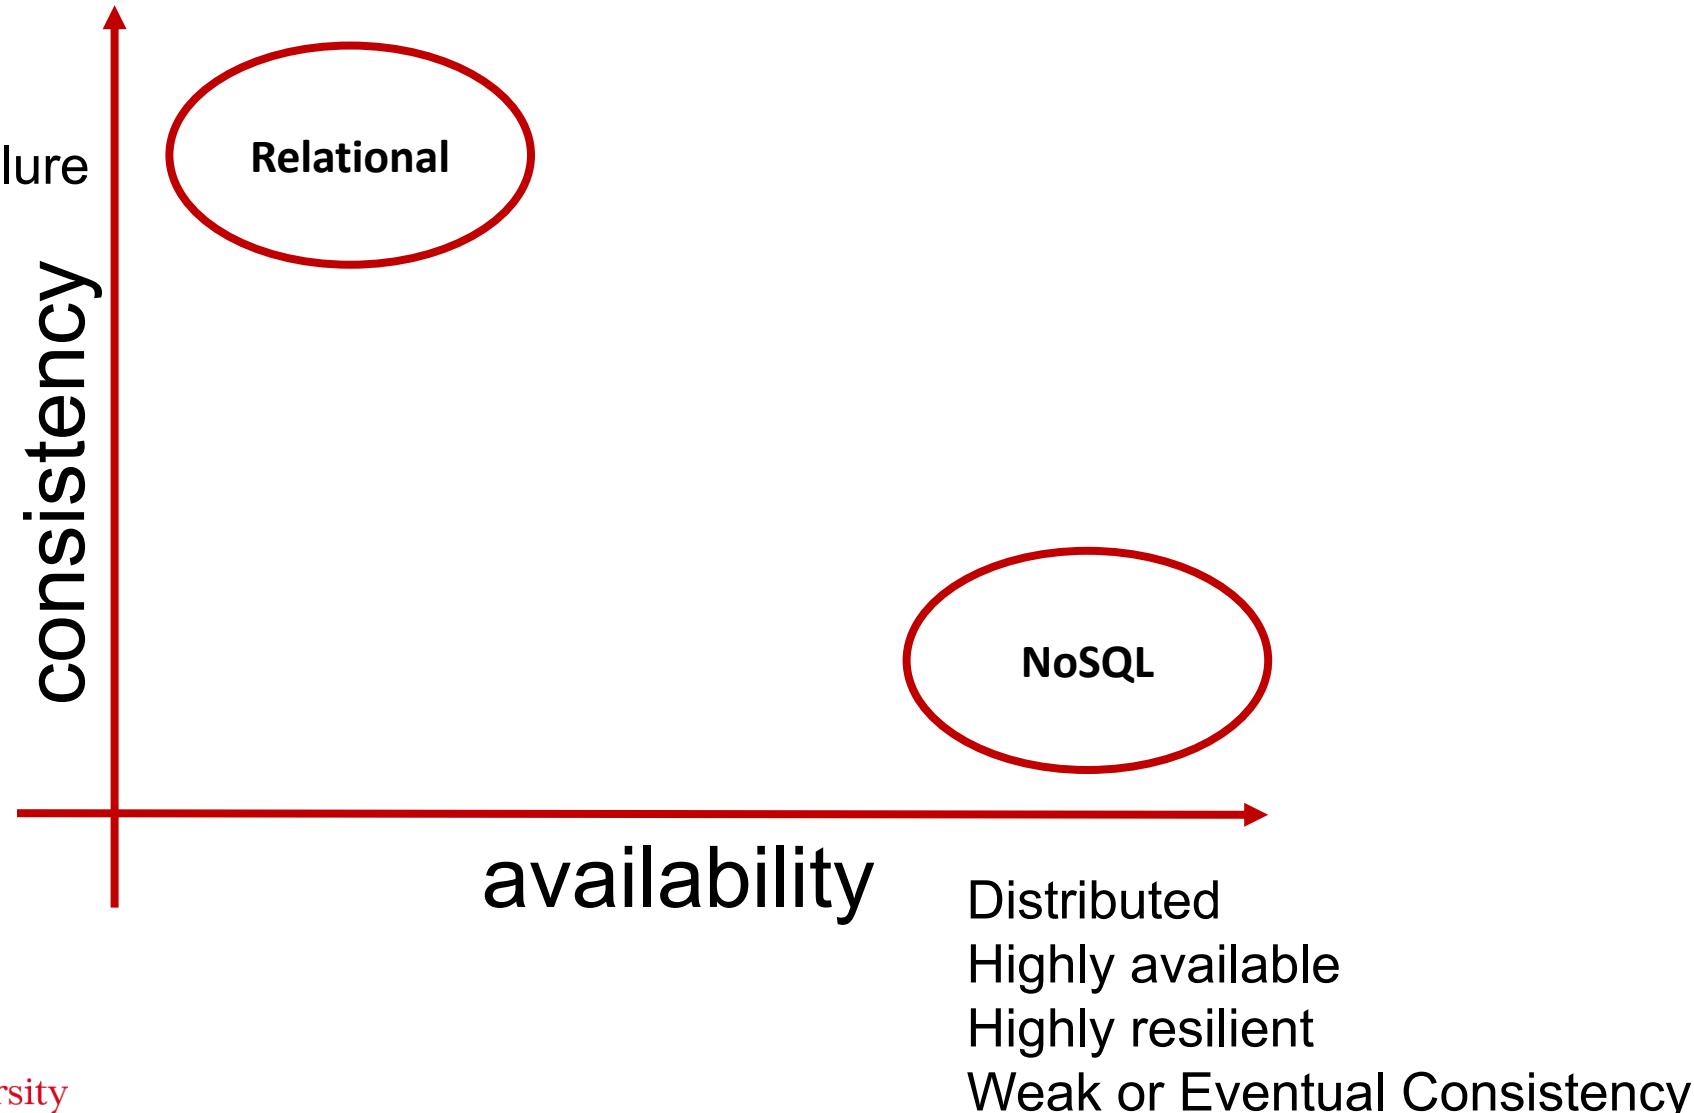


Downside: Difficult to scale



# Relational vs. NoSQL

Centralized  
Strict consistency  
Transactional  
Single Point of Failure



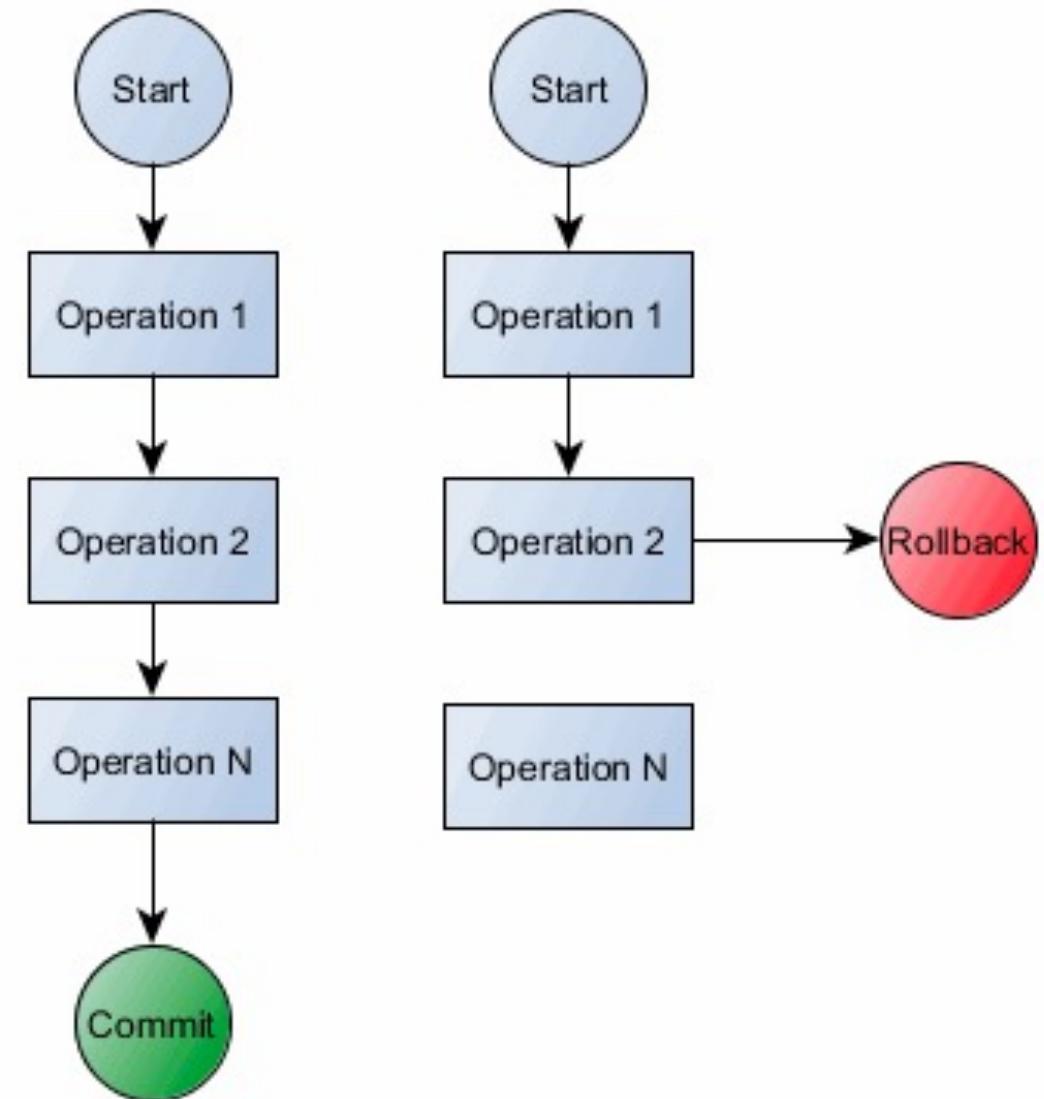
# ACID Properties

**Atomicity:** Transactions are *all or nothing*: They either succeeds or fails as a whole.

**Consistency:** Transactions on a database must transform the database from one consistent state to another.

**Isolation:** Partial effects of incomplete transactions should not be visible to other transactions.

**Durability:** Effects of a committed transaction are permanent and must not be lost because of later failure.



# Transfer \$50 from account A to account B

---

Read (A)

$A = A - 50$

Write (A)

Read (B)

$B = B + 50$

Write (B)

**Atomicity:** Don't take money from A without giving to B

**Consistency:** No money should be lost or gained

**Isolation:** Other queries shouldn't see A or B change until the entire transaction is completed and both values are updated.

**Durability:** The money does not go back to A



# BASE (An admittedly contrived acronym!)

---

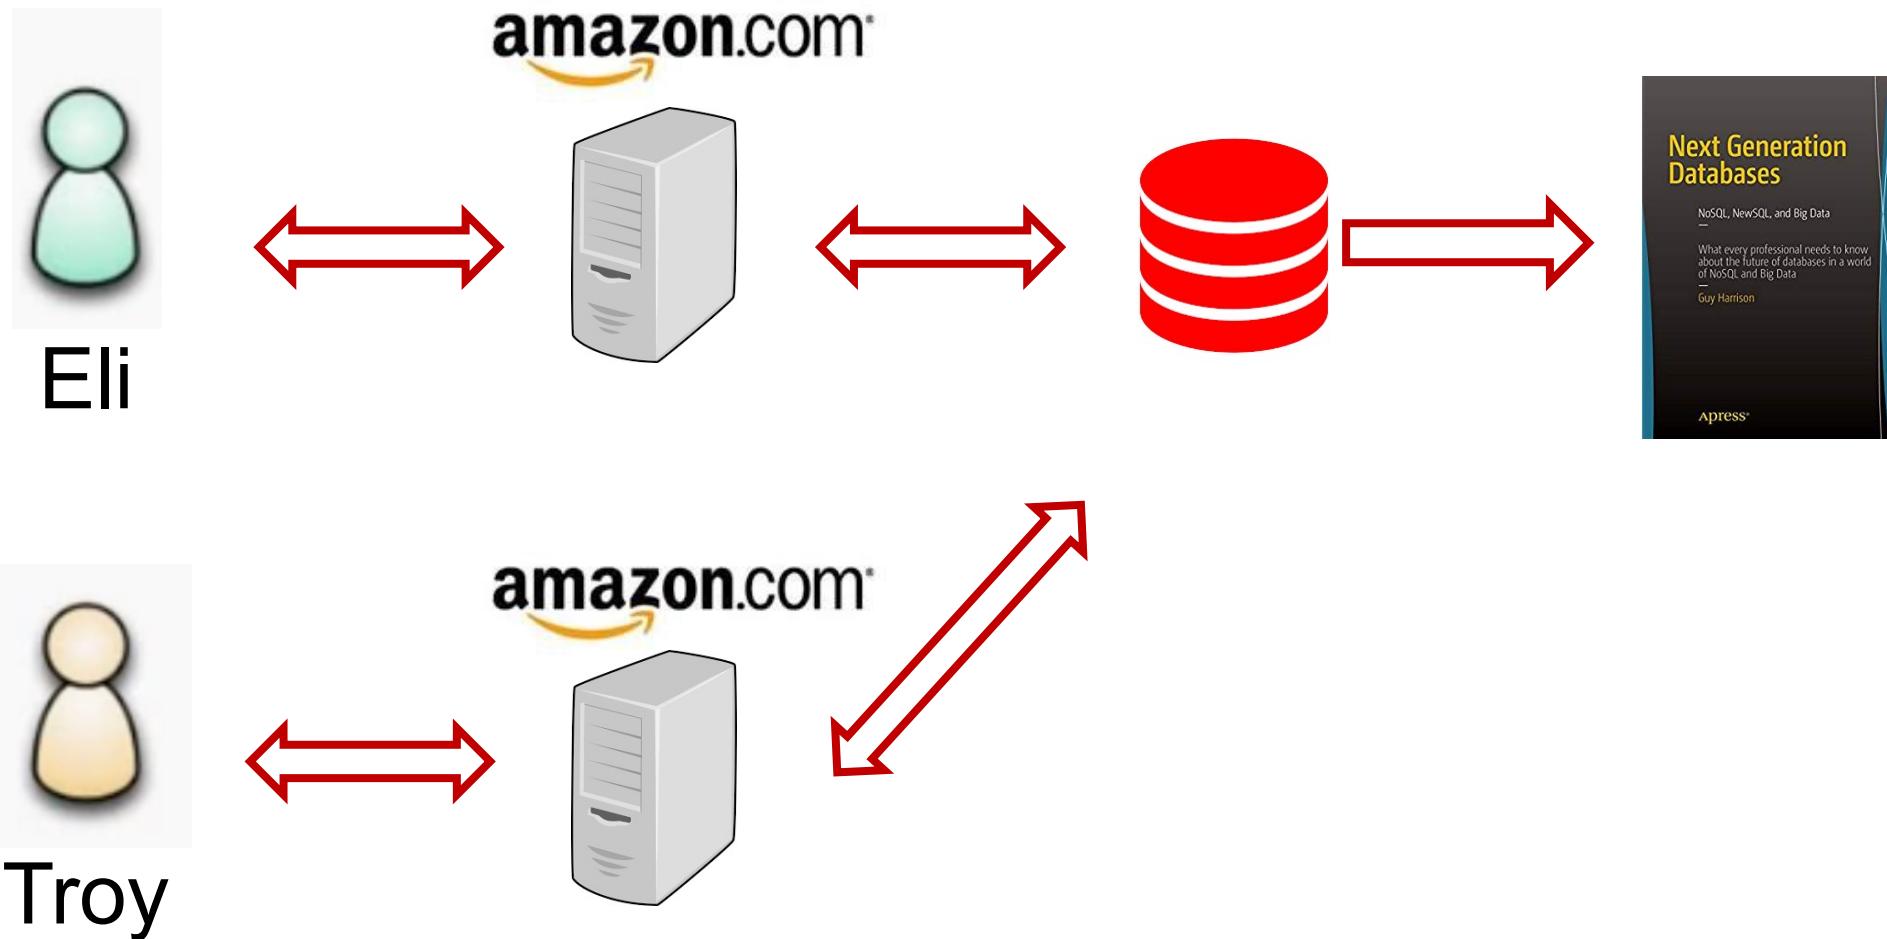
RDBMS → “ACID”

NoSQL → “BASE”

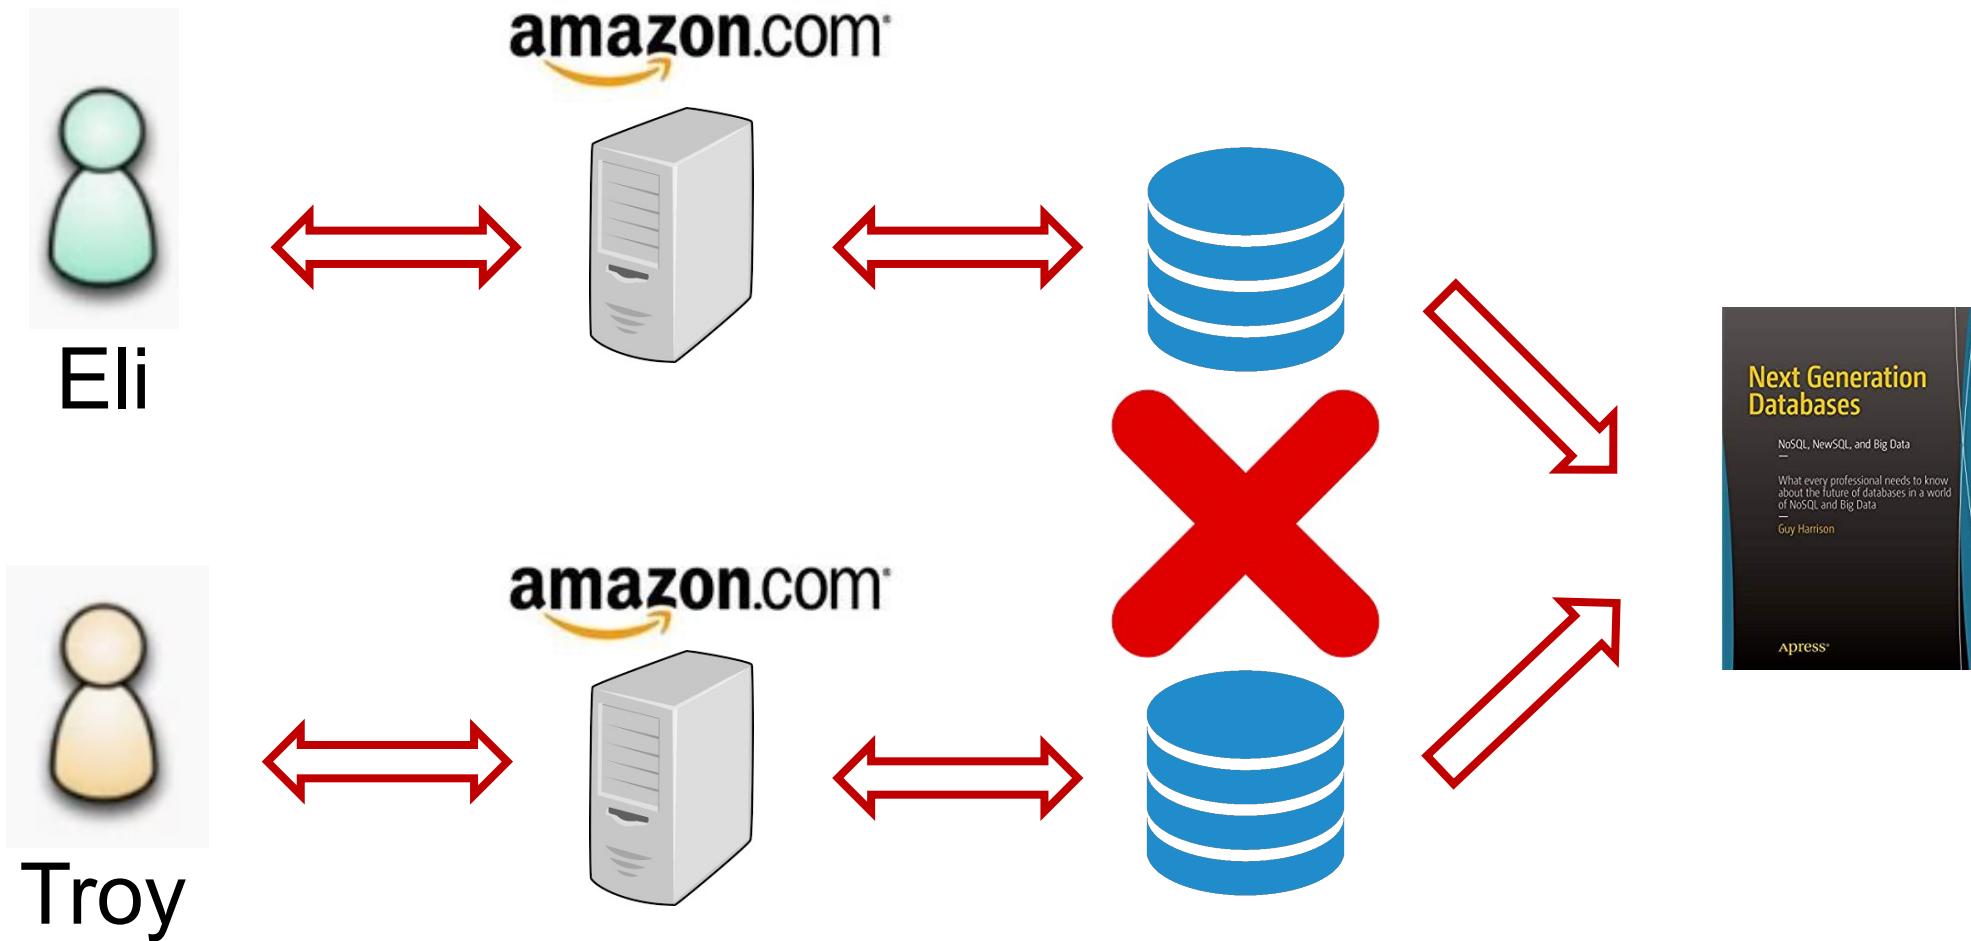
- Basically Available: A partial failure in a distributed system (e.g., one node) doesn't kill the whole system. If each bit of data is replicated to multiple nodes, system can continue unabated.
- Soft state: Data may be over-written with more recent data
- Eventually consistent. There may be times when the database is in an inconsistent state, particularly when data is replicated to multiple nodes for better availability.



# Consistency



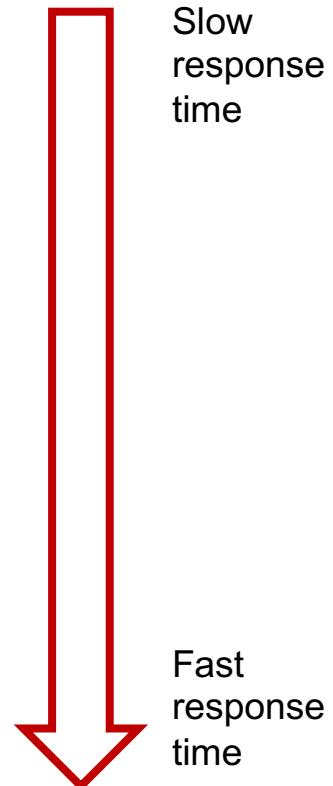
# Consistency vs. Availability



# Levels of Consistency

**Database consistency** is concerned with guaranteeing that a database always appears truthful to the user. Every operation must take the database from one consistent state to another.

- **Strict:** Changes to data are atomic and appear instantaneous
- **Sequential:** Clients see changes in the same order they were applied
- **Causal:** All causally-related changes are observed in the same order by all clients
- **Eventual:** All updates will eventually propagate through the system and all replicas will be consistent
- **Weak:** No propagation guarantees. Changes may appear out of order for some clients



# Drivers for the adoption of NoSQL

- The rise of Big Data and Cluster Computing coupled with the inherent limits of the relational model
- Object-Relational mapping challenges (The *Impedance Mismatch* problem)
- Simplified storage and analytics of semi-structured data

Exhibit 5  
Hadoop and NoSQL accounted for only ~3% of total market, but are expected to grow rapidly going forward.



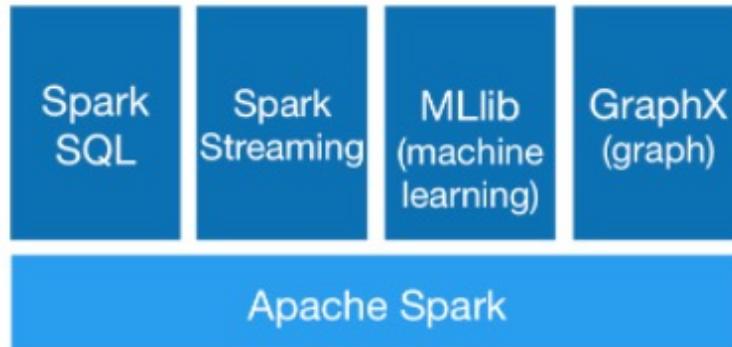
Source: IDC, Bernstein analysis

(But relational isn't going away!)

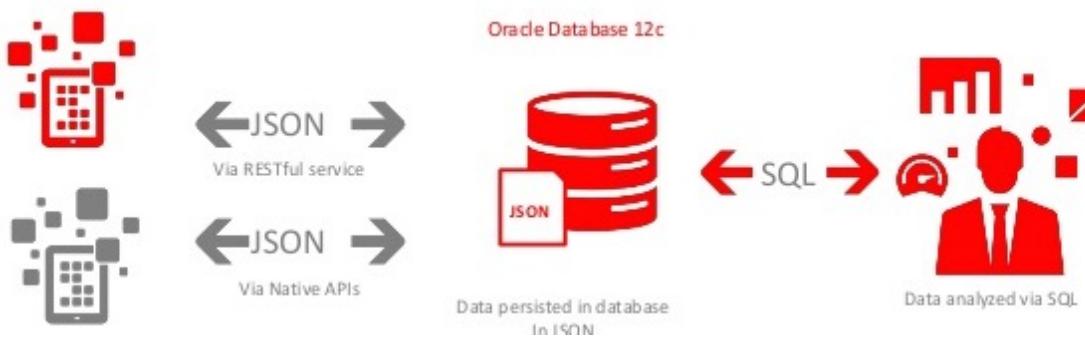


# Overlapping capabilities: Some examples

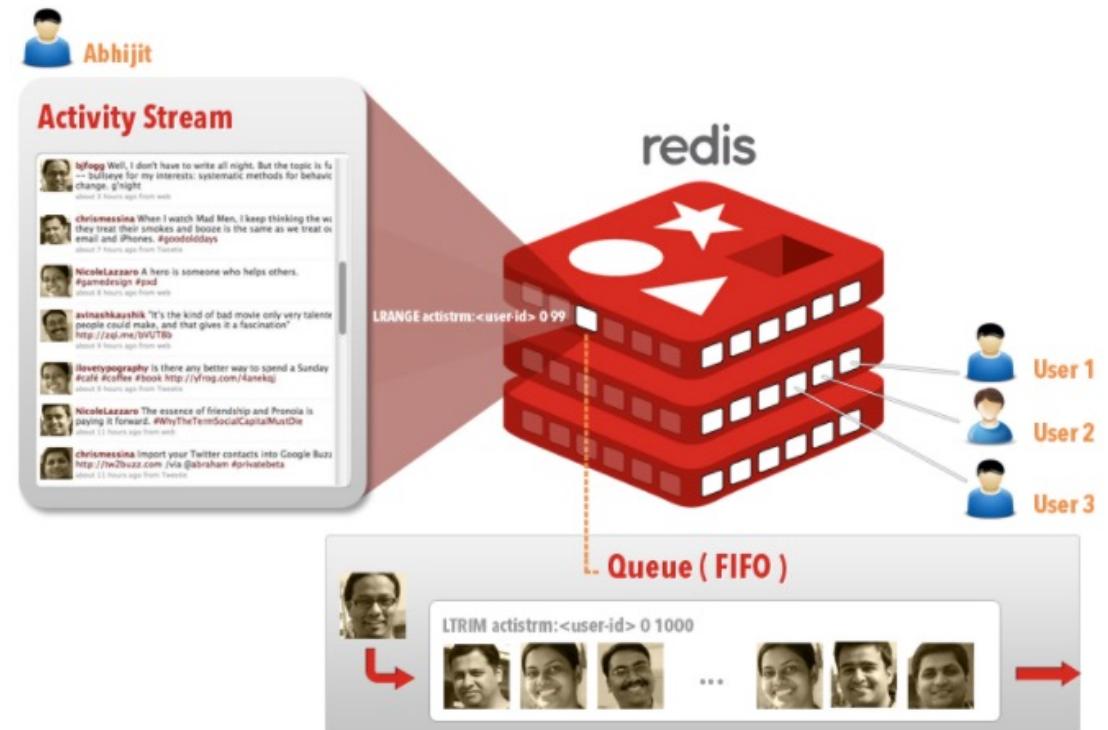
**Apache Spark:** A batch-processing engine with libraries for SQL-like operations, streaming, ML, and graph processing



**Oracle NoSQL:** Document storage (JSON-format) but with SQL-like querying and analytics.

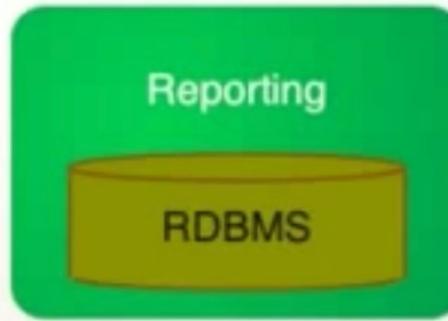
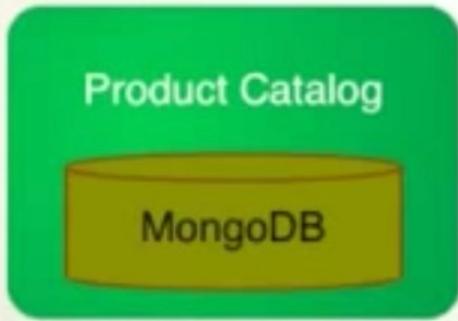
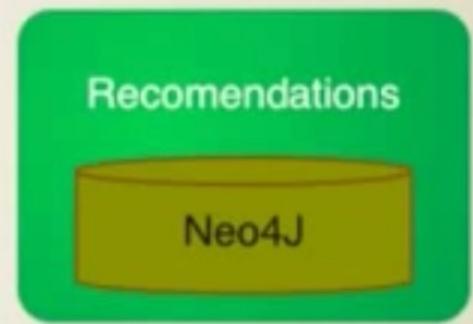
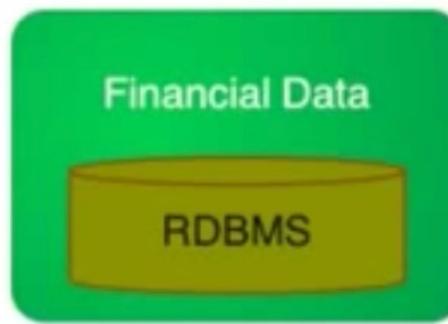
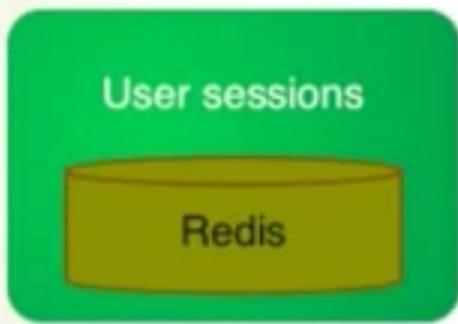


**Redis:** A key-value store but which can be used for stream-processing / event-processing and message-passing.



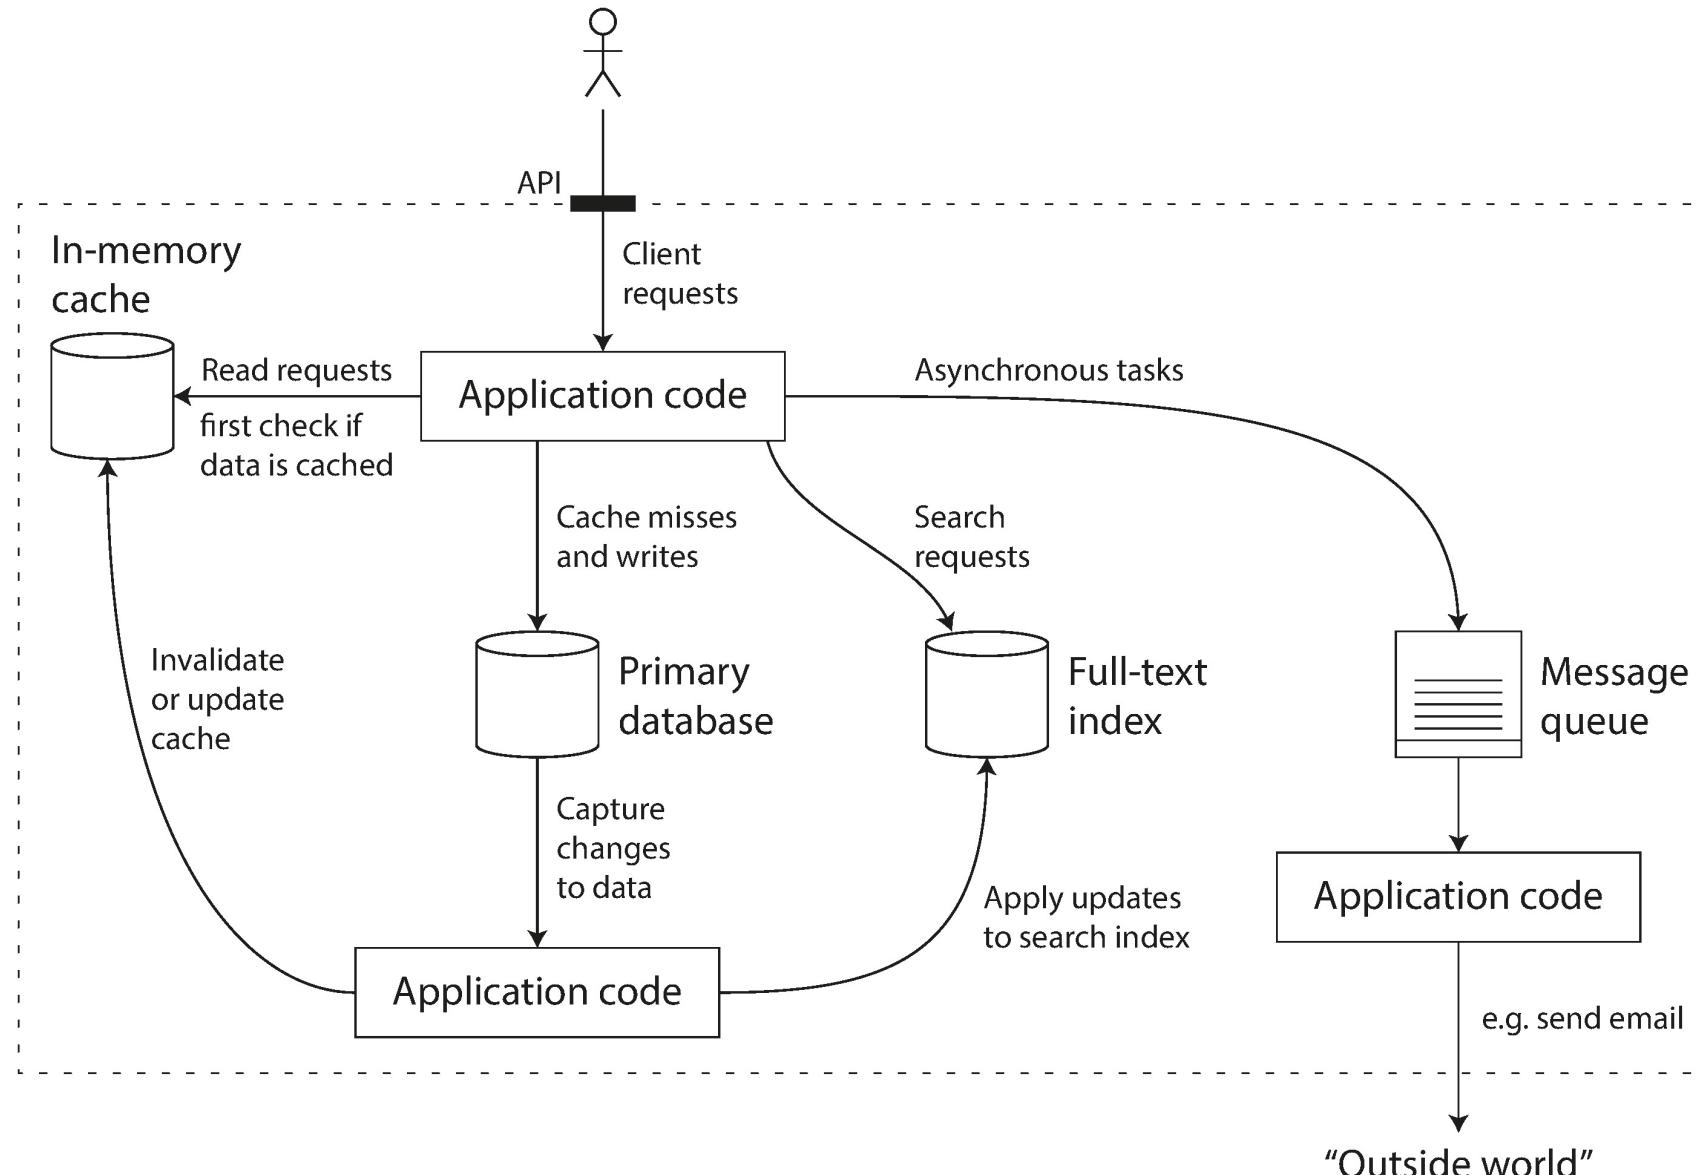
# Polyglot Persistence

## Speculative Retailers Web Application



# Polyglot Persistence

- Many tools optimized for different use-cases. No single tool accommodates all situations. *“One size does not fit all.”*
- Different components are stitched together by application software.

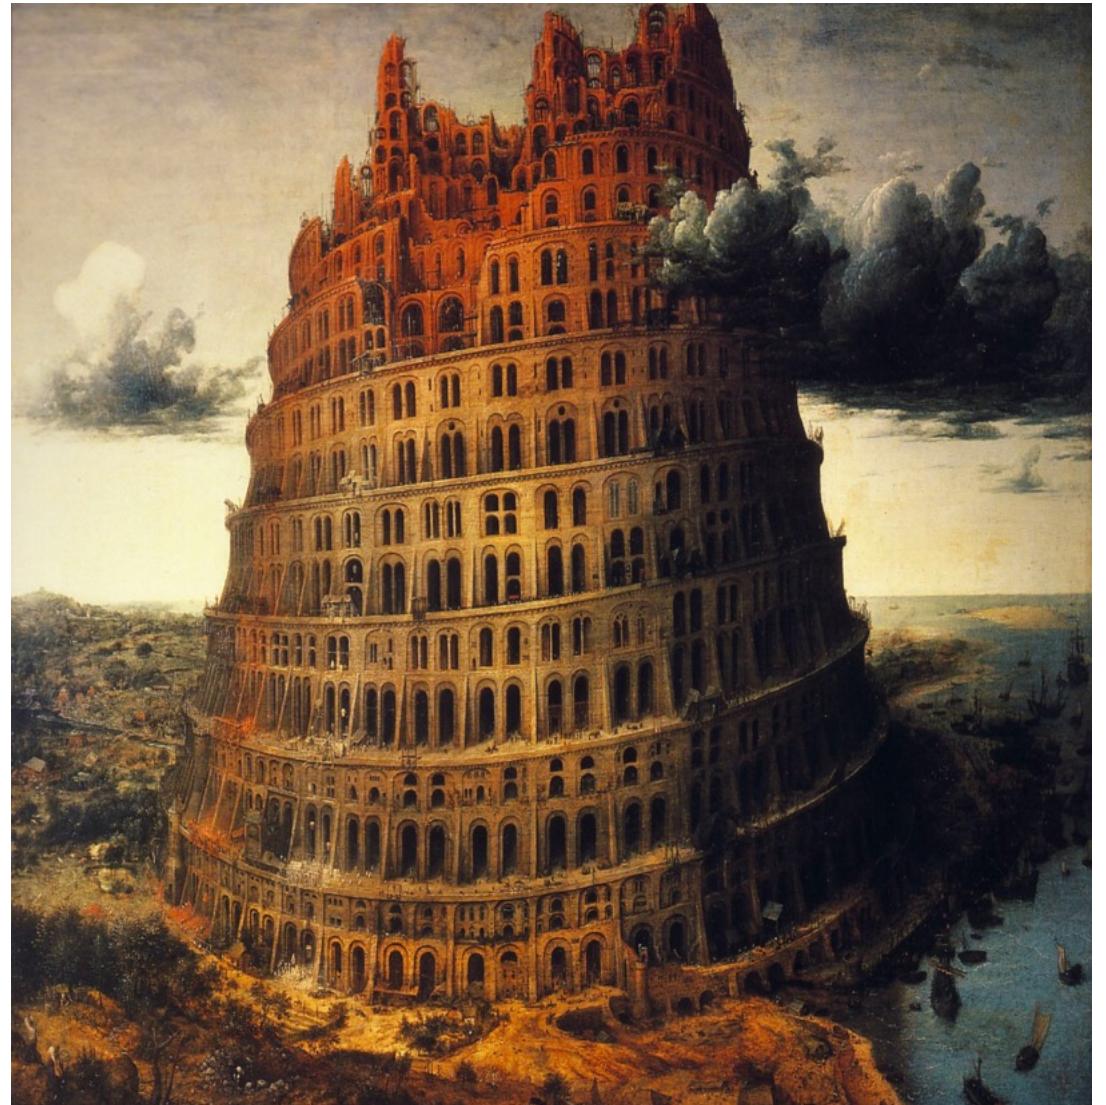


# The database landscape

---

Resembles a Tower of Babel with many systems (large and small) interoperating together.

- Relational DBMS
- NoSQL for special use-cases
- Caching with Key-Value stores
- Messaging brokers for asynchronous communication & coordination / Streaming
- Batch processing with Hadoop or Spark



# The rise of data engineering

- Many applications are *data-intensive* rather (merely) than *compute-intensive*.
- CPUs are not the only limiting factor, but also the size (volume), flow rate (velocity), and complexity (variety) of the data.
- This has given rise to the field of *data engineering* (and demand for data-engineers rather than just software engineers.)

DATA ENGINEER  
“SOFTWARE ENGINEERS BY TRADE”

## Role

*Develops, constructs, tests and maintains architectures (such as databases and large-scale processing systems)*



## Mindset

*All-purpose everyman*

HIRED BY



## Languages

*SQL, Hive, Pig, R, Matlab, SAS, SPSS, Python, Java, Ruby, C++, Perl*

## Skills & Talents

- ✓ *Database systems (SQL & NO SQL based)*
- ✓ *Data modeling & ETL tools*
- ✓ *Data APIs*
- ✓ *Data warehousing solutions*



and  
fitbit  
too!



# Elements of a data-centric application

---

- **Databases:** data storage and retrieval (Relational and NoSQL)
- **Caching:** Remembering the result of an expensive operation to speed up reads
- **Indexing:** Allow users to search data by keyword or filter it in various ways
- **Continuous (Stream) Processing:** Asynchronous data flow programming for real-time data
- **Batch:** Periodically crunching large amounts of accumulated data



# Data Science vs. Data Engineering

**Data Scientist:** Performs descriptive statistics to develop insights, build models, solve a business need.

**Data Engineer:** Architect large scale data processing systems typically using relational, NoSQL, messaging, and distributed computing platforms.

## Data Engineering

**Develop, construct, test, and maintain architectures** (such as databases and large-scale processing systems)



**Ensure architecture** will support the requirements of the business



Discover opportunities for **data acquisition**



**Develop data set processes** for data modeling, mining and production



Employ a variety of languages and tools (e.g. scripting languages) to **marry systems together**



Recommend ways to **improve data** reliability, efficiency and quality



## Data Science

Conduct research to **answer industry and business questions**



**Leverage large volumes of data** from internal and external sources to answer that business



Employ sophisticated analytics programs, machine learning and statistical methods to **prepare data for use in predictive and prescriptive modeling**



Explore and examine data to **find hidden patterns**



**Automate work** through the use of predictive and prescriptive analytics



**Tell stories to key stakeholders** based on their analysis



# Diverse technologies with blurred boundaries

